

AI Coding Agent Buying Guide for Claude Code, Codex, Cursor and GitHub Copilot

2 September 2026 - ID G00856492 - 20 min read

By: Bill Ryan, Hannah Decker, Philip Walsh, Mike Tucciarone, Zach Nagle, Jo Liversidge, Luke Parker, Sara Gravesande, Melanie Alexander

Initiatives: Technologies and Markets; Enabling AI Software Engineering Foundations; Financial

Software engineering leaders are under pressure to decipher AI coding agent offerings from Anthropic, OpenAI, Cursor, and GitHub. Use these insights to optimize investments by negotiating pricing and terms that prevent costs from spiraling out of control.

Insights at a Glance

The market for **AI coding agents** — including Claude Code, Codex, Cursor, and GitHub Copilot — has advanced rapidly from basic code completion into sophisticated, semi-autonomous agents capable of complex reasoning.

Commercially, AI coding agent vendors have shifted from flat per-user pricing to variable, consumption-based models based on metered tokens and shared API credit pools. Consequently, buyers must adopt a nuanced approach to doing business with these vendors focusing on the following:

- **Claude Code:** Pursue a sales-assisted Enterprise deal for large deployments to unlock pricing flexibility, as direct negotiation can yield PUPM fee waivers and usage discounts not available via self-serve Enterprise, Team, or API-only deals.
- **Cursor:** Negotiate a waiver for the \$40 PUPM Enterprise fee, secure precommitted and on-demand token discounts, and lock in the additional \$0.25 per million Cursor token fee for predictable budgeting. Establish clear renewal pricing and terms to delay repackaging given the recent SpaceX acquisition.
- **OpenAI Codex:** Strive for PUPM fee reductions and prepaid credit discounts based on spend commitment. Closely monitor Codex token-to-credit conversion and usage patterns, since changes in the rate card or consumption can lead to unexpected cost increases without notice.

- **GitHub Copilot:** Leverage promotional credit uplift during negotiations, utilize GitHub AI Credit prepurchase plans to minimize monthly overages, and educate developers on efficient usage and Auto mode savings.

To achieve AI-driven productivity gains, software engineering leaders must actively manage negotiation, budgeting, use case selection, and internal spend governance to ensure cost predictability, operational continuity, and sustained value from these vital enterprise tools.

Strategic Planning Assumption

Through 2028, 90% of enterprise AI coding agent contracts will be renegotiated early as vendors revise pricing and customers exceed initial usage commitments.

Issue

The market for AI coding tools — such as Claude Code, Codex, Cursor, and GitHub Copilot — is evolving at a breakneck pace. Vendors have abandoned predictable, flat-rate, per-user pricing with generous entitlements in favor of variable commercial models that rely on metered token billing, API usage pools, and premium credit requirements.

Frequent vendor pricing and model changes combined with unpredictable usage and growth are making costs increasingly difficult to forecast and budget. Without proper controls, heavy developer usage can rapidly deplete shared enterprise credit pools or cause massive budget overages, potentially interrupting development work.

According to the Gartner Peer Community, 55% of leaders reported more than \$200 per developer per month in AI coding agents' token costs in July 2026, generated by tools such as Claude Code, Cursor and OpenAI Codex, with 11% of organizations reporting more than \$1,000 in monthly token costs per developer. ¹ To prevent costs from spiraling out of control, software engineering leaders must proactively navigate changing feature sets and commercial approaches by negotiating optimal hybrid pricing, pooled usage, and API token usage discounts while implementing rigorous spend controls.

Impact

Productivity: According to the 2026 Gartner AI in Software Engineering Survey, 90% of engineering leaders now report productivity gains attributable to AI, but with most reporting gains under 25%. ² This shows organizations are better at integrating AI coding agents, though further productivity gains remain. For more on using AI to increase software development productivity, see [AI's Impact on Developer Productivity Across the SDLC](#).

Financial volatility: The vendor shift to **consumption-based pricing** increases development cost variability, creating severe budgeting challenges. AI is also reshaping engineering budgets, with AI spend rising from ~15% in 2025 to ~22% in 2026 as delivery shifts toward AI-augmented models. ³

Procurement complexity: Software engineering leaders must evaluate changing feature sets and pricing across multiple plans while adopting new strategies for budgeting and negotiating usage-based subscriptions (see [How to Evaluate Enterprise AI Coding Agent Pricing](#)). Otherwise, cost increases could outstrip productivity gains.

More Detail

The Evolution of AI Coding Agents Is Reshaping Pricing and Packaging

AI coding tools have evolved from assistants that suggest code and answer questions in chat into agents that can execute multistep workflows with greater autonomy and operate across more of the software development life cycle. Leading offerings include Claude Code, Codex, Cursor and GitHub Copilot. For a full assessment of this market, see [Magic Quadrant for Enterprise AI Coding Agents](#) and [Critical Capabilities for Enterprise AI Coding Agents](#).

This evolution changes how vendors package and sell products, shifting from per-user subscriptions with bundled entitlements to spend commitments and usage-based charges. These models increase cost variability, risk budget overruns, and reduce predictability. Consequently, software engineering leaders must evaluate pricing alongside usage allowances, overage terms, entitlements, and consumption controls. This guide examines how leading vendors go to market and the commercial implications buyers should consider for AI coding agents.

Claude Code, Codex, Cursor and GitHub Copilot are sold through multiple license options and buying channels. This Explainer Brief focuses primarily on offerings that are sold directly to enterprise customers rather than those tailored for individuals and smaller-scale deployments.

Table 1: AI Coding Agent Commercial Buying Guide Summary

(Enlarged table in Appendix)

Item	Claude Code (Anthropic)	Codex (OpenAI)	Cursor (SpaceXAI)	GitHub Copilot (Microsoft)
Offering	Claude Enterprise	ChatGPT Enterprise	Cursor Enterprise	GitHub Copilot Enterprise
Seat pricing	\$20 PUPM list price. Does not include a usage entitlement.	ChatGPT: Negotiable based on volume, often \$25-\$40 PUPM. Codex-only: \$0 PUPM (purely usage-based).	\$40 PUPM list price. Does not include a usage entitlement.	\$39 PUPM list price. Includes 3,900 AI Credits per user. AI Credits are pooled across the user base on a monthly basis.
Usage pricing	Billed at standard API rates, pooled across the user base.	\$0.04 per credit, pooled. Credits are used to consume Codex tokens.	First-party pool (Composer 2.5, Grok 4.6) and API pool (3P models). Each billed on a pooled usage basis.	Additional usage beyond the included AI Credit pool is charged at \$0.01 per AI credit.
Third-party model access and pricing	N/A (Anthropic models only).	N/A (OpenAI models only). Note: Codex is also available as an open-source stand-alone harness.	3P usage billed at standard API rates plus \$0.25 per million tokens surcharge. OpenAI shutoff 12 November 2026.	No surcharge for accessing 3P models. Auto Mode provides a 10% discount.
Primary negotiation priorities	Demand \$20 PUPM fee waiver and API discounts in exchange for spend commitment.	Seek lower PUPM and credit pricing in exchange for user and spend commitment. If applicable, attempt to retain for one year to plan for transition.	Demand \$40 PUPM waiver, discounts on precommitted and on-demand usage, and partial rollover. Seek clear renewal pricing and terms to mitigate repackaging risk associated with recent acquisition by SpaceX.	Negotiate using promotional uplift as an anchor, leverage GitHub AI Credit prepurchase plans for annual credit pools, educate developers on cost-efficient usage, and align access with business priorities.

Source: Gartner 2026

Claude Code: Seek PUPM Fee Waiver and API Rate Card Discounts

Claude Code is available with Pro, Max, Team, and Enterprise subscriptions, and through an Anthropic Console account billed on API usage. The current single Claude Enterprise seat packaging includes Claude Code.

Claude Enterprise is tailored for large-scale operations and includes capabilities for governance, security, and customization. This includes role-based access control, SCIM for identity management, audit logging, a compliance API for observability, and other features. ⁴

Claude Enterprise seat pricing does not contain a usage entitlement.

- Claude Enterprise has a \$20 PUPM list price, which is purely an access fee to Claude on web, desktop, and mobile, plus Claude Code and Claude Cowork. The \$20 seat fee is billed annually, with midterm seat additions prorated. Organizations enabling U.S. only inference are billed at 1.1x standard API rates for Claude Opus 4.6, Claude Sonnet 4.6, and later models.
- Usage is billed at API rates. ⁵

Two direct purchasing routes for Claude Enterprise:

- **Self-serve (20 seat minimum):** Customers purchase online without sales involvement through credit card or ACH with usage drawn down from the prepaid account. Consumption stops when the balance is exhausted. Customers must agree to Anthropic list pricing and standard terms of service.
- **Sales-assisted (50 seat minimum):** Involves direct interaction with the Anthropic sales team for pricing and terms negotiation. Customers receive standard invoicing and are billed monthly in arrears for actual usage. PUPM fee waivers, usage discounts, and tailored MSA terms and conditions may be offered in exchange for a significant annual spend commitment. In that situation, customers who fail to reach the spend commitment will receive a true-up invoice at the end of the term to fulfill their commitment.

Claude Enterprise is also available indirectly through AWS Marketplace. Marketplace purchases are billed monthly in arrears on actual consumption.

Making contact with Anthropic sales has proved challenging for some prospects

- Anthropic has struggled to scale its account teams quickly enough to meet demand.
- Some prospects struggle to establish a relationship with Anthropic, give up doing so, and execute a self-serve enterprise deal at list price.

Recommendations for Software Engineering Leaders Procuring Claude Code

- When considering licensing for a large number of users and a six-figure annual spend commitment, seek a sales-assisted Enterprise deal.
- Demand a waiver of the \$20 PUPM monthly fee for a sales-assisted Enterprise deal and tie costs strictly to usage. Clarify in the order form that additional users will also receive the PUPM fee waiver.
- Leverage an annual or multiyear spend commitment to secure discounts off standard API rates. *Expect discount offers in the 5% to 15% range in exchange for a significant commitment. Recognize that the situation is fluid, multiyear deals are rare, and results could vary widely based on leverage.*
- Include detailed model pricing (the rate card table) in the Claude Enterprise order form to protect against mid-contract price increases.
- Attempt to negotiate terms to prevent negative unilateral changes by Anthropic while preserving the ability to benefit from cost-lowering changes (e.g., price decreases). Gartner has little evidence this is currently achievable, but Anthropic prospects should still strive for it.

OpenAI Codex: Quantify Token Usage and Convert to Credits

OpenAI's Codex is an AI coding agent available to ChatGPT Plus, Pro, Edu, Business and Enterprise users. OpenAI provides \$0 PUPM usage-based Codex seats for ChatGPT Enterprise customers and grandfathered ChatGPT Business plans active before 24 June 2026. As of that date, new Business plans or workspaces without prior Codex seats can no longer add them.

At the time of publication, OpenAI is offering a new ChatGPT Enterprise pricing variant that shifts financial risk toward the customer. Previously functioning on a more predictable per-user, per-month (PUPM) fee with unlimited instant model usage and an allocation of credits for advanced features, the new usage-based approach increases budget risk by unbundling the credit entitlement from the user entitlement and pricing it separately, at an assumed rate of \$0.04 per credit. For Codex, customers must use credits to consume tokens based on the Codex rate card. ⁶

ChatGPT Enterprise offers Codex access to Codex via a hybrid pricing model.

- PUPM fees provide access to ChatGPT and Codex, but Codex *usage* requires credit consumption.

- Until mid-2026, OpenAI included credits as a per-user entitlement in ChatGPT Enterprise deals. In this deal structure, access to credits is pooled across the user base, with credits expiring *monthly*.
- More recently, Gartner has also observed OpenAI proposing credits as a separate line item in the ChatGPT Enterprise order form. In this deal structure, credits are priced at \$0.04 each, with credits expiring *annually*, which provides improved allocation flexibility as compared to the monthly “use-it-or-lose-it” approach.

Codex usage is priced based on API token usage, calculated as credits per million tokens.

- Credit usage depends on the mix of input, cached input, and output tokens in each task.
- The Codex rate card describes how credits convert to tokens for Codex consumption. ⁷ For example, GPT 5.6 Sol costs 125 credits per million input tokens, and 750 credits per million output tokens. At \$0.04 per credit, this makes Codex pricing equivalent to \$5 per million input tokens and \$30 per million output tokens. Other OpenAI models follow the same logic.

Recommendations for Software Engineering Leaders Procuring Codex

- Negotiate PUPM and credit fees, pushing for lowered pricing based on user volume.
- Seek to utilize the previous user-plus model that includes pooled credits on a PUPM basis for at least one year, especially if your organization has already bought under this model and is coming up for renewal. This will provide your organization with time to prepare for the operational and financial impacts of the new structure that prices credits separately from user fees.
- If buying under the new pricing structure that splits out credits on a separate line item at \$0.04 each, secure a discounted allocation of prepaid credits and negotiate the standard overage rate down if leverage allows (e.g., from \$0.04 to \$0.035 or lower).
- Calculate and monitor how Codex token usage converts to credits. Changes to the Codex rate card can increase credit requirements without changing the pricing model, creating stealth cost increases.

- Investigate billing dashboards to help determine usage and make forecasting consumption easier, even downloading .csv files if necessary, to drill down to model-level specifics and align analytics with renewal planning.
- Set appropriate monthly usage caps by persona — such as developers relying heavily on Codex — and proactively manage unpredictable consumption to avoid massive budget blowouts (see [How to Optimize Token Consumption for AI Coding Agents](#)).

Cursor: Negotiate Token Rates and Enterprise PUPM Fee Waiver

Cursor is an enterprise AI coding agent platform that offers access to third-party models (including models from Anthropic and Google) as well as its own first-party models. On 14 August 2026 it was announced that Cursor was officially acquired by SpaceX (see [First Take: SpaceX's Cursor Acquisition Gives Its Enterprise AI Ambitions a Launch Vehicle](#)).⁸ The acquisition has enabled Cursor to access the compute needed to enter the frontier model race, work with SpaceX to create Grok 4.6, which is now its flagship first-party model. Following the acquisition, OpenAI stated on 28 August 2026 that it plans to terminate its model provision agreement with SpaceX (Cursor), targeting a shutoff date of 12 November 2026 (see [First Take: OpenAI's Cursor Withdrawal Signals a Shift Toward Vertically Integrated AI Coding Agents](#)).

Cursor offers two distinct usage pools per plan: the first-party models pool and the API pool.⁹

- The first-party models pool includes Cursor's first-party models (Composer 2.5, Grok 4.6). Usage of these models is metered by token type (input, output, cache read/write).
- The API pool covers usage of third-party models, billed at each model's standard API rate *plus an additional \$0.25 per million token fee*. Model selection directly impacts how quickly this allowance is consumed. When users exceed their included usage, they can either pay on-demand at the same per-token rates or upgrade to a higher plan tier.

Customers should not assume that access to third-party frontier models will remain unchanged.

Auto Mode Streamlines Model Selection and Can Reduce Costs. Cursor's Auto mode uses the recently released Cursor Router to automatically choose the best AI model for each coding task, optimizing performance and cost based on the user's selection of Cost, Balance or Intelligence modes. It simplifies user experience and typically routes requests to included, lower-cost proprietary models. To use Auto mode, Cursor requires that Grok 4.5 be enabled.

Team and Enterprise plans introduce additional features like centralized billing, team administration, and enhanced agent limits. Team pricing is per user and includes usage entitlements each month, with Premium seats offering higher usage limits. Enterprise plans offer a single seat type priced at \$40 PUPM with usage charged on a pooled basis.

For Teams and Enterprise, a Cursor Token Rate (\$0.25 per million tokens) applies to third-party model requests, on top of third-party API pricing, while first-party models are exempt. Large enterprise deals involve precommitted API usage blocks and annual payment terms. Actual spend depends on model selection, agent usage, and whether users leverage Cursor's Auto/Composer models or manually select third-party models. Contracts are typically 12 months with automatic renewal.

Stronger spend controls as a selling point. Cursor added usage visibility and spend alerts in 2026. Procurement-friendly spend controls could become a differentiator against GitHub Copilot, Claude Code, and OpenAI Codex. For more on Cursor's approach to delivering customer-facing controls for AI usage and model management, see [Emerging AI Pricing and Customer Cost Controls Accelerate Adoption](#).

Recommendations for Software Engineering Leaders Procuring Cursor

- Request that Cursor waive the \$40 PUPM monthly fee when negotiating an enterprise deal and tie costs strictly to usage. Clarify in the order form that additional users will also receive the PUPM fee waiver.
- Negotiate discounts on token pricing for precommitted and on-demand usage volume, and additional discounts if you pursue a multiyear commitment. Strive for on-demand discounts to match precommitted discounts, which are often lower.
- Negotiate renewal protections to maintain the \$0.25 per million Cursor Token Rate for the next contract term.

- Cursor includes a one-year autorenewal clause in its order forms but it requires mutual agreement. Revise the renewal clause to ensure the decision to renew for at least one additional year is yours, not Cursor's. This is critical following SpaceX's acquisition, which heightens the risk of Cursor changing its pricing and packaging.
- Model access guarantees: Auto mode uses Cursor's proprietary routing. Require that third-party frontier models remain accessible at current credit costs for the term of the agreement, or that you receive a price credit if access degrades.

GitHub Copilot: Usage-Based Pricing Shift Creates Buying Implications

GitHub Copilot is an AI coding agent accessible through IDE integrations, CLI, cloud, web and mobile experiences integrated with the GitHub platform. GitHub Copilot offers access to 1P and 3P models, and unlike Cursor, does not add a surcharge for 3P model usage. It also offers Auto mode (auto model selection) to pick the most efficient and cost-effective AI model based on task complexity and other requirements.

GitHub Copilot's go-to-market strategy has evolved significantly in response to shifts in AI technology economics and enterprise demand. Initially, GitHub Copilot was offered under a flat, per-user seat pricing model, which provided predictable, subsidized costs and unlimited access to AI-powered coding assistance. This approach facilitated rapid adoption, as organizations could easily forecast expenses and scale usage without concern for overages.

GitHub Copilot's transition to usage-based pricing introduces greater cost volatility.

- Effective 1 June 2026, GitHub Copilot transitioned from the more predictable seat-based licensing to a variable, usage-based pricing model (see [First Take: 1 June Marks the End of Subsidized AI Coding as GitHub Copilot Switches to Usage-Based Pricing](#)).¹⁰
- Under this new structure, GitHub officially replaced Premium Request Units (PRUs) that were allocated on a per-seat basis with pooled AI credits. This transition effectively turns AI coding tokens into a metered resource and shifts the financial risk of high consumption from the vendor directly to the enterprise customer. *Client preview bills have indicated that GitHub Copilot spend could increase by three to 10 times under this new model* (see [How to Respond Proactively to the GitHub Copilot Pricing Change](#)).

Notable characteristics of GitHub Copilot's new pricing model include:

- The \$39 PUPM GitHub Copilot Enterprise fee provides 3,900 AI Credits (\$0.01 per credit), with overage credits also priced at \$0.01. This introduces variable spend potential and requires organizations to actively monitor usage rather than budgeting solely on license counts.
- Included credits are aggregated into a shared enterprise pool that all licensed users consume from. Unused credits are forfeited at the end of each month, while exceeding the credit pool will incur incremental costs.
- Once pooled credits are exhausted, GitHub no longer offers an automatic fallback to a free model. Instead, usage continuation is dictated by the enterprise's configured budget controls, and AI budget exhaustion can actively interrupt live development work.
- Developer behavior now has a measurable cost impact. Auto mode reduces premium request consumption by 10%, creating opportunities to lower costs.
- A temporary, promotional credit uplift applies between 1 June and 1 September 2026, providing a buffer period before steady-state pricing takes effect. GitHub Copilot customers receive 7,000 AI credits PUPM during this period instead of 3,900.¹⁰

Recommendations for Software Engineering Leaders Procuring GitHub Copilot

- Leverage the promotional uplift as part of your negotiation strategy. GitHub Copilot is on record offering 7,000 credits per user, per month from 1 June to 1 September 2026, making the effective price per credit \$0.006 during the promotional period.⁸ Couple that data point with competitive pressure and adoption uncertainty to drive PUPM and credit price reductions.
- Utilize GitHub AI Credit prepurchase plans which provide an annual allotment of credits. Azure customers who need 2M GitHub AI credits or more can benefit from tiered discounts starting at 5%. Further, as GitHub AI credits are billed through an Azure subscription, MACC customers can use this spend to draw down annual commitment targets. Note, that the Azure Commitment Discount (ACD) does not stack on top of the prepurchase plan discount, making PAYG more cost-effective when the ACD is higher than 5%.¹¹
- Educate developers on the financial impact of usage. Make developers aware of the shift to credit-based pricing and how their behavior directly affects credit consumption and costs. Train developers on cost-efficient prompt engineering and encourage the use of Auto mode for its 10% discount on premium requests.

- Prioritize GitHub Copilot access for teams and projects that generate measurable productivity gains rather than broadly deploying licenses across the entire developer population. Establish clear governance policies for usage, such as leveraging spending alerts and requiring explicit overage approvals from budget owners to exceed spend commitments.

Evidence

¹ **Gartner Peer Community Poll.** This poll was conducted in July and August 2026 with 46 respondents. Gartner Peer Community is a platform for technology leaders to engage in discussions with peers and share knowledge in real time. The results are not representative of the market as a whole or of Gartner's position on the topic.

² **Gartner Software Engineering Survey for 2026.** This survey was conducted to provide a comprehensive understanding of the current landscape in software engineering, as well as to determine the priorities and strategic challenges of software engineering leaders. It also aimed to identify the demand for various roles and skills within software engineering organizations, and assess their budget expectations, team structures and organizational outcomes. Finally, it explored the integration of AI in software engineering workflows and its impact on engineering organizations. The survey was conducted online from August through November 2025 among 482 respondents from the U.S. (n = 360) and U.K. (n = 122). Qualifying organizations operated in multiple industries and reported enterprisewide revenue for fiscal year 2024 of at least \$250 million or equivalent. Qualified participants were highly involved in managing software engineering/application development teams and the activities they perform. Disclaimer: The results of this survey do not represent global findings or the market as a whole, but reflect the sentiments of the respondents and companies surveyed.

³ **2026 Gartner AI in Software Engineering Survey.** This survey was conducted to explore how software engineering functions were approaching AI, including budget allocations, the current state of implementation and impact, AI strategy, and use cases. The survey focused on two key areas: the use of AI tools (such as **AI code assistants** and AI code agents) across the software development life cycle (SDLC) and the development of AI-powered solutions (such as, **AI agents**) within the software engineering function. The research was conducted online from April through May 2026 among 300 respondents across various industries and regions, including North America (n = 150), Europe (n = 100), and Asia/Pacific (n = 50). Qualifying organizations reported enterprisewide annual revenue of at least \$250 million (or equivalent) in fiscal year 2025 and had experience using AI tools within the SDLC, as well as developing AI-powered solutions within their software engineering function. Participants were senior engineering leaders and were involved in decision making related to the use of AI within the engineering function. Disclaimer: The results of this survey do not represent global findings or the market as a whole, but reflect the sentiments of the respondents and companies surveyed.

⁴ [Plans & Pricing](#), Claude.

⁵ [What is the Enterprise plan?](#), Claude Support.

⁶ [ChatGPT Rate Card \(Business, Enterprise/Edu, credit-based pricing\)](#), OpenAI.

⁷ [Pricing](#), OpenAI Business.

⁸ [Cursor is now a part of SpaceX](#), Cursor.

⁹ [Models & Pricing](#), Cursor.

¹⁰ [GitHub Copilot is moving to usage-based billing](#), GitHub Blog.

¹¹ [GitHub Enterprise pricing](#), Microsoft Azure.

Recommended by the Authors

Some documents may not be available as part of your current Gartner subscription.

[Avoid AI Budget Blowouts: Master Enterprise Pricing for ChatGPT, Claude and Gemini](#)

[4 Strategies to Manage the Escalating Cost of AI Coding Assistants and Agents](#)

[How to Manage Token Costs for Custom-Built AI Agents](#)

[How to Evaluate Enterprise AI Coding Agent Pricing](#)

[How to Keep Up With the AI Coding Agent Market](#)

[Emerging AI Pricing and Customer Cost Controls Accelerate Adoption](#)

[Magic Quadrant for Enterprise AI Coding Agents](#)

[First Take: SpaceX's Cursor Acquisition Gives Its Enterprise AI Ambitions a Launch Vehicle](#)

[AI's Impact on Developer Productivity Across the SDLC](#)

© 2026 Gartner, Inc. and/or its affiliates. All rights reserved. Gartner is a registered trademark of Gartner, Inc. and its affiliates. This publication may not be reproduced or distributed in any form without Gartner's prior written permission. It consists of the opinions of Gartner's Business and Technology Insights Organization, which should not be construed as statements of fact. While the information contained in this publication has been obtained from sources believed to be reliable, Gartner disclaims all warranties as to the accuracy, completeness or adequacy of such information. Although Gartner insights may address legal and financial issues, Gartner does not provide legal or investment advice and its insights should not be construed or used as such. Your access and use of this publication are governed by [Gartner's Usage Policy](#). Gartner prides itself on its reputation for independence and objectivity. Its insights is produced independently by its Business and Technology Insights Organization without input or influence from any third party. For further information, see "[Guiding Principles on Independence and Objectivity](#)." Gartner insights may not be used as input into or for the training or development of generative artificial intelligence, machine learning, algorithms, software, or related technologies.

Table 1: AI Coding Agent Commercial Buying Guide Summary

Item	Claude Code (Anthropic)	Codex (OpenAI)	Cursor (SpaceXAI)	GitHub Copilot (Microsoft)
Offering	Claude Enterprise	ChatGPT Enterprise	Cursor Enterprise	GitHub Copilot Enterprise
Seat pricing	\$20 PUPM list price. Does not include a usage entitlement.	ChatGPT: Negotiable based on volume, often \$25-\$40 PUPM. Codex-only: \$0 PUPM (purely usage-based).	\$40 PUPM list price. Does not include a usage entitlement.	\$39 PUPM list price. Includes 3,900 AI Credits per user. AI Credits are pooled across the user base on a monthly basis.
Usage pricing	Billed at standard API rates, pooled across the user base.	\$0.04 per credit, pooled. Credits are used to consume Codex tokens.	First-party pool (Composer 2.5, Grok 4.6) and API pool (3P models). Each billed on a pooled usage basis.	Additional usage beyond the included AI Credit pool is charged at \$0.01 per AI credit.
Third-party model access and pricing	N/A (Anthropic models only).	N/A (OpenAI models only). Note: Codex is also available as an open-source stand-alone harness.	3P usage billed at standard API rates plus \$0.25 per million tokens surcharge. OpenAI shutoff 12 November 2026.	No surcharge for accessing 3P models. Auto Mode provides a 10% discount.
Primary negotiation priorities	Demand \$20 PUPM fee waiver and API discounts in exchange for spend commitment.	Seek lower PUPM and credit pricing in exchange for user and spend commitment. If applicable, attempt to retain for one year to plan for transition.	Demand \$40 PUPM waiver, discounts on precommitted and on-demand usage, and partial rollover. Seek clear renewal pricing and terms to mitigate repackaging risk	Negotiate using promotional uplift as an anchor, leverage GitHub AI Credit prepurchase plans for annual credit pools, educate developers on cost-efficient usage, and align

associated with recent
acquisition by SpaceX.

access with business
priorities.

Source: Gartner 2026