

Avoid AI Budget Blowouts: Master Enterprise Pricing for ChatGPT, Claude and Gemini

15 May 2026 - ID G00846413 - 15 min read

By: Bill Ryan, Melanie Alexander, Christopher Dixon

Initiatives: Financial; Pioneer New Approaches to Optimize Costs and Realize Greater Value; Master AI Contracting Complexities and Orchestrate AI Buying Across the Enterprise; Alleviate Tech Buyers' Regret to Stop Value Leakage Across the Enterprise; Technologies and Markets

IT SPVM leaders are under pressure to decipher and negotiate Anthropic Claude, Google Gemini, and OpenAI ChatGPT enterprise offerings. Use these insights to optimize hybrid pricing, maximize pooled usage, and secure contract provisions that prevent costs from spiraling out of control.

Insights at a Glance

Demand for rapid GenAI adoption is surging, with Gartner client interactions about Anthropic Claude, Google Gemini, and OpenAI ChatGPT skyrocketing 90% in Q1 2026 compared to Q1 2025. Despite this strong interest, organizations face significant financial risks — such as budget overruns and acquiring unused “shelfware”— if they procure these consumption-based technologies without rigorous commercial diligence.

Given divergent pricing models and rigid vendor negotiation stances, IT sourcing, procurement and vendor management (SPVM) leaders must develop a targeted commercial strategy for each vendor's unique offering to navigate opaque licensing and avoid a margin-killing deployment.

Key insights and actions:

- All three vendors employ hybrid pricing models — per-user, per-month (PUPM) fees plus usage-based charges — with negotiable variables that are unique to each and constantly evolving.

- For Anthropic Claude Enterprise, aim to negotiate a waiver of the \$20 per-user, per-month (PUPM) fee and obtain discounts on API rates when making a significant spend commitment. Crucially, the waiver must accommodate user growth, and the agreed-upon API pricing should be explicitly documented on the order form.
- For Google Gemini Enterprise, maximize value and avoid overage fees by aligning users to the appropriate edition. Negotiate discounted per-user, per-month (PUPM) fees and secure price protection for the duration of the contract to optimize costs.
- For OpenAI ChatGPT Enterprise, negotiate lower PUPM fees based on user volume and per-user credit allocation to consume advanced features. Extend monthly credit expiration to align with the agreement term, and extract discounts on overage rates and API spend. Protect against unilateral rate card changes during the term that can significantly increase costs.

This Explainer Brief focuses specifically on the enterprise versions of ChatGPT, Claude, and Gemini. Please note that Anthropic, Google, and OpenAI each provide additional offerings tailored for individual users, teams, and smaller-scale deployments, which are not covered in this analysis.

Strategic Planning Assumption

By 2028, 80% of enterprises adopting usage-based AI will be forced to impose strict internal cost controls to prevent budget overruns, up from less than 10% in 2026.

Issue

IT sourcing, procurement, and vendor management (SPVM) leaders face major challenges procuring enterprise-grade AI capabilities. The 2026 Gartner Heads of IT SPVM Signature Survey shows 61% struggle to forecast AI/GenAI costs, and 60% find securing cost protections (caps, rate guarantees, credit rollovers) difficult to negotiate. ¹

In this environment, IT SPVM leaders must rapidly deliver services from providers like Anthropic, Google Cloud, and OpenAI despite a complex and changing landscape of licensing and deployment options that include:

- **Divergent pricing approaches:** Anthropic, Google, and OpenAI all use hybrid pricing models that share similar principles. They all include both user- and usage-based components, and API pricing to target developers focused on custom applications and integration. However, the unique characteristics of each vendor's offering make direct comparisons exceptionally challenging.

- **High risk of unbudgeted costs:** Usage-based models with uncapped access risk significant budget overruns if not rigorously monitored. While flexible and consumption-aligned, failing to implement spending caps can lead to “runaway” costs driven by viral adoption, inefficient coding practices that drive excessive token consumption, or a lack of chargeback mechanisms to encourage more judicious use.
- **Vendor resistance to negotiation:** Anthropic, Google, and OpenAI have all taken a more rigid approach to negotiation, even with very large prospects. Negotiating pricing and usage provisions is challenging due to the combination of high customer demand for these services and vendor challenges scaling their own sales organizations to satisfy the demand.

Impact

Entering talks to procure Anthropic Claude Enterprise, Google Gemini Enterprise, or OpenAI ChatGPT Enterprise without adequate commercial due diligence can lead to significant financial and operational damage.

Opaque pricing and hybrid procurement models act as hidden margin killers, risking severe budget overruns or forcing enterprises to pay for unused “shelfware.” Without strict internal spend controls or negotiated price protections, variable credit and token usage costs can easily spiral.

In the current environment, where change is constant and established commercial standards are lacking, IT SPVM leaders must act decisively to master the nuances of Anthropic, Google, and OpenAI enterprise offerings, work with business stakeholders to accurately forecast usage and negotiate concessions to protect the organization from a cost and risk perspective.

More Detail

Shared Pricing Principles With Unique Variables

Procuring generative AI at scale is complex. Anthropic, Google, and OpenAI have repeatedly evolved their business models and shifted toward hybrid, consumption-based pricing structures. As a result, IT SPVM leaders face high risk of budget overruns, unused shelfware, and restrictive usage caps if contracts are poorly negotiated.

Recognize common pricing model elements. Anthropic, Google, and OpenAI all sell enterprise offerings with hybrid pricing models that include a mix of user- and usage-based components. For example, all three vendors:

- Lead with a per-user, per-month fee.
- Account for usage on top of the per-user, per-month price in the form of credits, tokens or quotas.
- Have online price lists that specify credit and/or token pricing for usage.
- Charge additional fees when customers exceed the entitlement or spend commitment.

Identify unique pricing model characteristics and their impact on costs. Anthropic, Google, and OpenAI enterprise offerings each have unique commercial elements relative to their competitors. For example:

- Anthropic's Claude Enterprise charges a per-user, per-month fee, but the user license does not include a usage entitlement. ² Instead, usage is billed on a pooled basis according to the Claude API rate card, which may or may not be discounted based on a spend commitment. ³
- Google's Gemini Enterprise offers the choice to license Gemini Enterprise Business, Gemini Enterprise Standard, or Gemini Enterprise Plus, with enhanced user rights at a higher price point for the Plus edition. Usage is pooled across licensed users, with usage limits (described as "quotas") that are measured on a per-day basis. ⁴
- OpenAI's ChatGPT Enterprise provides a user license with unlimited access to certain models and includes a monthly credit allocation that is pooled across licensed users and can be used to consume advanced features such as thinking, deep research and image generation. ⁵

IT SPVM leaders that invest sufficient time and effort to understand each vendor's pricing model nuances will be better positioned to quantify usage, budget accurately, and negotiate concessions (see Table 1).

Table 1: Decode ChatGPT, Claude, and Gemini Enterprise Offerings

(Enlarged table in Appendix)

Variable	Claude Enterprise	Gemini Enterprise	ChatGPT Enterprise
Core positioning	AI agents, safe AI, long-context reasoning	AI agents, unified agentic platform, enterprise search	AI agent platform, powered by business context and frontier models
Pricing model	PUPM user fee and pay-as-you-go (PAYG) usage consumption at API pricing	PUPM user fee and pooled quota limits on features and usage	PUPM user fee with PUPM credit allocation for consuming advanced features
PUPM fee provides:	Access to features, but no usage	Access to features, up to daily quota limits	Access to GPT-5.5 Instant, included enterprise features, and credits to access certain advanced capabilities*
Pooled or per-user usage	Pooled	Pooled	Pooled
Credit/quota limits reset	N/A	Daily	Typically monthly
Sales approach	Self-service or Sales-assisted	Self-service for Business, sales-led for Standard and Plus	Sales-led
Negotiation points specific to each vendor	■ PUPM fee (remove) for both current and additional users ■ Spend commitment ■ API rate card discounts	■ PUPM fee (reduce), overage fees ■ Charge overages instead of imposing restrictions when quotas are reached	■ PUPM fee (reduce) ■ Credit allocation \$0.07 overage rate monthly credit expiration
Negotiation points applicable to all three vendors	■ Prevent vendor ability to unilaterally degrade rates and increase multipliers during term ■ Clarify renewal term length and pricing ■ Customize MSA terms and conditions		
* At the time of publication, subject to change.			

Source: Gartner (May 2026)

Anthropic Claude Enterprise: Seek PUPM Fee Waiver and API Discount

The Claude Enterprise plan is tailored for large-scale operations and includes capabilities for governance, security, and customization. This includes role-based access control, SCIM for identity management, audit logging, and a compliance API for observability.

Anthropic is on its third iteration of Claude Enterprise pricing in the past year. It was originally introduced as a seat-based model with Standard and Premium seat types and per-seat limits on usage. It then shifted to a two-seat model differentiated by functionality (Claude Chat and Claude Chat + Claude Code) with usage charged on a consumption basis.

In February 2026, Anthropic transitioned once again, this time to a single seat type and PUPM fee, with additional fees based on actual consumption. It includes a \$20 PUPM list price, with usage pooled across its user base and billed at API rates. A minimum of 20 seats and an annual commitment are required.

Two notable characteristics of Claude Enterprise include:

- **Seat pricing that does not contain a usage entitlement.** Unlike ChatGPT Enterprise and Gemini Enterprise, Anthropic's Claude Enterprise PUPM fee does not include a usage-based entitlement. It is simply the price of admission to use Claude's capabilities and only provides each user with access to Claude on web, desktop, and mobile, plus Claude Code and Claude Cowork.
- **Two flavors of its Enterprise offering.** ⁶
 - The **"Self-serve"** version allows customers to purchase online without sales involvement. It requires prepaid user and API licensing via credit card or ACH with usage drawn down from the prepaid account. In addition, customers must agree to Anthropic list pricing and standard terms of service.
 - The **"Sales-assisted"** version is a more typical enterprise deal. It involves direct interaction with the Anthropic sales team for pricing and terms negotiation. Customers receive standard invoicing and are billed monthly in arrears for actual usage. PUPM fee waivers, usage discounts, and tailored MSA terms and conditions may be offered in exchange for a significant annual spend commitment. In that situation, customers who fail to reach the spend commitment will receive a true-up invoice at the end of the term to fulfill their commitment.
- **Making contact with Anthropic sales can be hard.** One reason for the "self-serve" vs. "sales-assisted" bifurcation of Claude Enterprise is Anthropic's challenges scaling its salesforce quickly enough to meet incoming demand. As a result, many prospects struggle to establish a relationship with Anthropic, give up on doing so, and execute a self-serve Enterprise deal at list price.

Recommendations for IT SPVM Leaders Negotiating Claude Enterprise Deals

- When considering licensing for hundreds of users and a six-figure annual spend commitment, seek a sales-assisted Enterprise deal. Be patient and persistent, and use every formal or informal connection at your disposal to make contact with Anthropic Sales.

- Push Anthropic to waive the \$20 PUPM monthly fee and tie costs strictly to usage. Clarify in the order form that additional users will also receive the PUPM fee waiver.
- Leverage an annual or multiyear spend commitment to secure discounts off standard API rates. Expect discount offers in the 5% to 15% range in exchange for a significant commitment. Recognize that the situation is fluid, multiyear deals are rare, and results could vary widely based on negotiation leverage and your relationship with Anthropic.
- Include detailed model pricing in the Claude Enterprise order form.
- Attempt to negotiate terms to prevent negative unilateral changes by Anthropic while preserving the ability to benefit from cost-lowering changes. Gartner has little evidence this is currently achievable, but Anthropic prospects should still strive for it.

Google Gemini Enterprise: Optimize Edition Selection and Usage Sizing

Google Gemini Enterprise is a per-user enterprise AI subscription sold through Google Cloud that combines Gemini and third-party model access, permission-aware enterprise search, data connectors, prebuilt and custom agents, security and governance controls.

Google has taken a “good, better, best” approach to Gemini Enterprise packaging. It is priced per user and scaled via enterprise editions. For example:

- Business edition is positioned as the lighter, lower-cost option for up to 300 users.
- Standard and Plus are for larger organizations needing higher quotas, stronger security/compliance controls, and support for custom/full-code agents and Gemini Code Assist. ⁴
- The Plus edition is most frequently considered by Gartner clients when evaluating Gemini Enterprise.

Usage quotas and overages are key considerations when evaluating Standard and Plus editions. All Gemini Enterprise editions provide access to chat, search, and 3P connectors. Standard and Plus provide higher quotas than Business, including video/image generation and Deep Research. All quotas are pooled across the organization, which gives customers more flexibility than when usage is limited on a per-user basis. However, Google uses daily quotas for many usage measures, which introduces the risk of waste, overage fees, or the need to buy additional licenses to expand daily pooled usage capacity. ⁷

Recommendations for IT SPVM Leaders Negotiating Gemini Enterprise Deals

- Align licensed users to the appropriate edition based on expected workloads and use cases.
- Benchmark user-based PUPM fees and negotiate lower pricing based on user commitment.
- If reaching daily quota limits is a risk, determine whether to purchase additional licenses to increase the size of the feature pool available, as amending daily quota limits via negotiation to maximize value and optimize usage is unlikely.
- Request deployment credits: In addition to pushing for direct per-seat discounts, ask for Google Cloud credits or professional services vouchers. During the late stages of a negotiation, Google may be more willing to concede on providing credits than further lowering PUPM fees and deal size.

OpenAI ChatGPT Enterprise: Negotiate PUPM Fees and Credit Flexibility

OpenAI's ChatGPT Enterprise offering provides what OpenAI describes as "enterprise-grade AI, security, and support at scale." This includes advanced security (SOC 2 compliance), single sign-on (SSO), analytics dashboards, dedicated support, and data privacy assurances — by default, customer content from ChatGPT Enterprise is not used to train OpenAI models. ⁸

ChatGPT Enterprise pricing has evolved since the option was first introduced. Originally, ChatGPT Enterprise was offered as a user-based package that allowed unlimited usage for a \$60 per-user, per-month (PUPM) fee. OpenAI then lowered the PUPM price from \$60 to \$40 per month in late 2024. In 2025, OpenAI transitioned to a "user+" model, where the PUPM fee provides unlimited access to certain features such as ChatGPT 5.3 Instant, and a monthly credit allocation that is provided to consume advanced features (see [OpenAI Procurement Options for ChatGPT Enterprise or APIs](#)).

Enterprise users share a credit pool at the contract level. All Enterprise users draw from this pool when using advanced features. There are no per-seat usage caps by default, but administrators can use role-based access controls (RBAC) to set spend controls by group. ⁹ Typically, the pricing for ChatGPT Enterprise ranges from \$25 to \$40 per user per month, with credit allocations averaging in a range from 200 to 400 credits per user per month, pooled at the tenant level.

ChatGPT Enterprise pricing and terms are negotiable. Customers and prospects can negotiate the PUPM fee based on the number of users and the size of the credit allocation. The larger the user count, the greater the ability to negotiate lower PUPM pricing. However, the larger the credit allocation, the more challenging it will be to negotiate lower pricing. OpenAI's API models can also be purchased based on token-based pricing and API spend commitments; discounted pricing can sometimes be incorporated into Enterprise deals.

Recommendations for IT SPVM Leaders Negotiating ChatGPT Enterprise Deals

- Benchmark user-based PUPM fees and negotiate lower pricing based on user commitment and credit allocation. Recognize that more users and fewer credits can be used to improve the PUPM price, but higher credit allocations can make price negotiation more challenging.
- Negotiate an extension of OpenAI's standard monthly credit expiration to maximize value and minimize waste and overages. Aim for annual or multiyear expiration, aligning with your overall agreement period.
- Negotiate the standard \$0.07 overage price per credit to \$0.065 or lower if leverage allows.
- Include the ChatGPT rate card in the order form rather than relying on the hyperlinked pricing page.
- Attempt to negotiate terms to prevent unilateral changes by OpenAI that increase costs, while preserving the ability to benefit from cost-lowering changes. Although Gartner has not yet seen this achieved, OpenAI prospects should still strive for it.
- Negotiate MSA terms to clarify data ownership, use and access, secure indemnification protection for IP infringement, and improve transparency regarding OpenAI's controls related to accuracy and responsible AI.

Evidence

¹ [2026 Gartner Heads of IT Sourcing, Procurement and Vendor Management Signature Survey](#). This survey was conducted to understand the key goals and challenges of heads of IT sourcing, procurement, and vendor management (SPVM). It also sought to understand how they are currently navigating the increased pressure felt across their organizations related to a growing set of technology and geopolitical demands compounded by expectations to support intensifying cost optimization pressures and an ever-broadening set of vendor relationships while being pushed beyond traditional metrics toward strategic value enablement and risk management. The research was conducted from 3 November through 12 December 2025. It had 203 respondents from North America (n = 103), Western Europe (n = 50), and Asia/Pacific (n = 50). Qualified organizations operated in multiple industries and reported enterprisewide full-time equivalents (FTEs) of at least 500 employees. Qualified respondents included CIOs, CTOs, and senior leaders of IT SPVM. All respondents held sole or shared decision-making authority for IT supplier selection and vendor management. Disclaimer: The results of this study do not represent global findings or the market as a whole, but reflect the sentiments of the respondents and companies surveyed.

² [Plans and Pricing: Claude by Anthropic](#), Claude.

³ [Pricing: Claude API Docs](#), Claude.

⁴ [Compare editions of Gemini Enterprise](#)

⁵ [ChatGPT Rate Card](#), ChatGPT.

⁶ [What is the Claude Enterprise Plan?](#), Claude.

⁷ [Gemini Enterprise Usage Quotas for Standard and Plus Editions](#).

⁸ [ChatGPT Business and Enterprise Pricing](#), ChatGPT.

⁹ [Role-Based Access Controls in your ChatGPT workspace](#)

Recommended by the Authors

Some documents may not be available as part of your current Gartner subscription.

[How to Control Costs in Credit and Usage-Based SaaS Negotiations](#)

[OpenAI Procurement Options for ChatGPT Enterprise or APIs](#)

[Protect Your Enterprise From AI Pricing and Contract Risks](#)

[Negotiate AI and Generative AI Pricing to Avoid Skyrocketing Costs](#)

© 2026 Gartner, Inc. and/or its affiliates. All rights reserved. Gartner is a registered trademark of Gartner, Inc. and its affiliates. This publication may not be reproduced or distributed in any form without Gartner's prior written permission. It consists of the opinions of Gartner's Business and Technology Insights Organization, which should not be construed as statements of fact. While the information contained in this publication has been obtained from sources believed to be reliable, Gartner disclaims all warranties as to the accuracy, completeness or adequacy of such information. Although Gartner insights may address legal and financial issues, Gartner does not provide legal or investment advice and its insights should not be construed or used as such. Your access and use of this publication are governed by [Gartner's Usage Policy](#). Gartner prides itself on its reputation for independence and objectivity. Its insights is produced independently by its Business and Technology Insights Organization without input or influence from any third party. For further information, see "[Guiding Principles on Independence and Objectivity](#)." Gartner insights may not be used as input into or for the training or development of generative artificial intelligence, machine learning, algorithms, software, or related technologies.

Table 1: Decode ChatGPT, Claude, and Gemini Enterprise Offerings

Variable	Claude Enterprise	Gemini Enterprise	ChatGPT Enterprise
Core positioning	AI agents, safe AI, long-context reasoning	AI agents, unified agentic platform, enterprise search	AI agent platform, powered by business context and frontier models
Pricing model	PUPM user fee and pay-as-you-go (PAYG) usage consumption at API pricing	PUPM user fee and pooled quota limits on features and usage	PUPM user fee with PUPM credit allocation for consuming advanced features
PUPM fee provides:	Access to features, but no usage limits	Access to features, up to daily quota limits	Access to GPT-5.5 Instant, included enterprise features, and credits to access certain advanced capabilities*
Pooled or per-user usage	Pooled	Pooled	Pooled
Credit/quota limits reset	N/A	Daily	Typically monthly
Sales approach	Self-service or Sales-assisted	Self-service for Business, sales-led for Standard and Plus	Sales-led
Negotiation points specific to each vendor	■ PUPM fee (remove) for both current and additional users ■ Spend commitment ■ API rate card discounts	■ PUPM fee (reduce), overage fees ■ Charge overages instead of imposing restrictions when quotas are reached	■ PUPM fee (reduce) ■ Credit allocation \$0.07 overage rate monthly credit expiration

Negotiation points applicable to all three vendors

- Prevent vendor ability to unilaterally degrade rates and increase multipliers during term
- Clarify renewal term length and pricing
- Customize MSA terms and conditions

* At the time of publication, subject to change.

Source: Gartner (May 2026)