

TMT: Security Software

**Junaid Siddiqui**

212-326-6141  
Junaid.Siddiqui@truist.com

**Miller Jump**

212-401-4427  
Miller.Jump@truist.com

**Jonathan Ngo**

212-419-4482  
Jonathan.Ngo@truist.com


26 Page Document

**Reasons for this report**

✓ Sector Update

## Security Software - Rise of the Models: Cybersecurity in the Age of Mythos and Daybreak

The rollout of advanced AI-native security platforms such as Claude Mythos and Daybreak is redefining the cybersecurity landscape by enabling faster threat detection, autonomous defense, and AI-driven vulnerability discovery. AI is accelerating, not displacing, cybersecurity demand as frontier models compress attack timelines, expand the attack surface, and make identity, data, and runtime protection more mission-critical. We continue to favor platform leaders CRWD, NET, PANW and ZS, alongside data/identity beneficiaries CVLT, RBRK, OKTA and SAIL, as we see platformization, telemetry scale, and growing agentic security needs driving durable share gains in cyber.



Source: Dall-E (OpenAI)

## Cybersecurity in the Frontier Age

- The rollout of advanced AI-native security platforms such as **Claude Mythos (Anthropic)** [Private] and **Daybreak (OpenAI)** [Private] is redefining the cybersecurity landscape by enabling faster threat detection, autonomous defense, and AI-driven vulnerability discovery. While this strengthens overall defensive capabilities, it is also creating heightened anxiety among incumbent cybersecurity vendors, as traditional solutions risk obsolescence if they fail to keep pace with AI-driven architectures. For investors, this shift introduces a mix of fear and opportunity—fueling volatility in cybersecurity stocks, expanding valuation dispersion between AI-forward leaders and legacy players, and increasing concern around margin compression, competitive disruption, and required reinvestment cycles across the sector.
- **Claude Mythos** and **Daybreak** signal a new reality: AI models can identify vulnerabilities in near real time. What once took days or weeks is now compressed into minutes, drastically shrinking the window between discovery and exploitation. At the same time, rapidly expanding attack surfaces are increasing exposure faster than most teams can manage.
- Emerging initiatives such as Anthropic's **Project Glasswing**, CrowdStrike's **Quiltworks**, and OpenAI's **Trusted Access for Cyber** highlight a shift toward deeply integrated, AI-first security ecosystems—where access control, threat detection, and response are increasingly automated, context-aware, and model-driven. This signals a future cybersecurity paradigm built on continuous identity validation, autonomous defense agents, and tightly governed AI access layers, accelerating the convergence of identity, endpoint, and AI security while reshaping competitive dynamics across the sector.
- **Partnerships validate incumbents' relevance:** The formation of ecosystems and coalitions among Anthropic, CrowdStrike, OpenAI, and established cybersecurity vendors signals that AI model providers are complementary rather than disintermediating. These collaborations reinforce that enterprise security will remain a layered stack—where LLM providers supply foundational capabilities, while cybersecurity vendors retain critical roles in orchestration, integration, compliance, and frontline defense—alleviating investor concerns about outright displacement while still pressuring vendors to evolve.
- **Platformization + Data Flywheel Advantage:** AI-native security platforms will win via data scale + integration flywheels, embedding models across end-to-end workflows (detect → respond) and compounding advantage through proprietary telemetry, ecosystem connectivity, and vendor consolidation, while point solutions without platform gravity are marginalized as replaceable signals or absorbed into broader stacks.
- **Stay Selective on the Security Space:** Platforms are where we believe value will accrue to. We favor platform vendors such as CrowdStrike (CRWD, Buy), Cloudflare (NET, Buy), Palo Alto Networks (PANW, Buy), and Zscaler (ZS, Buy), that are building an AI defensible moat ready for the agentic future. Additionally, Data & Identity Security are becoming core focus areas in age of agentic security, benefiting names such as Commvault (CVLT, Buy), Rubrik (RBRK, Buy), Okta (OKTA, Buy), and SailPoint (SAIL, Buy). As agentic AI evolves, the need to secure those autonomous agents becomes increasingly mission-critical. In an agent-driven environment, data and identity security are not just protective measures, they are the backbone of trust.

## Key Themes

### Security remains the top priority of CIO's and cybersecurity budgets continue to grow.

- Security is among the most defensible, and highest-priority verticals of IT software: Overall security spend is expected to be \$240 billion in 2026, according to Gartner, up around 12.5% YoY from \$213 billion in 2025, and expected to increase to \$300Bn+ by 2030.
- Cybersecurity is a **non-discretionary budget item** creating a large market opportunity as security growth exceeds other areas of IT.
- AI advancements are a double-edged sword. Attackers can use the same models to boost their abilities, while defenders gain continuous monitoring, faster threat identification, and automated patching at a scale no human team could match. Gen AI will dramatically grow the attacker population as it puts advanced cybercrime tools in the hands of attackers of all skill levels while at the same time helping organizations scale their cybersecurity posture. Gartner predicts that by 2027, 17% of all cyberattacks will involve generative AI.
- Data is the “gold” of digital security data has become a critical task as regulatory compliance and corporate governance objectives. Ensuring the security, availability, and integrity of data has become a critical task as regulatory compliance and corporate governance objectives affecting many organizations mandate the creation of multiple copies of data with longer and more complex retention requirements.

### The attack surface continues to expand in the agentic era while security transformation has lagged

- AI-generated attacks are rapidly escalating. According to the 2026 CrowdStrike Global Threat Report, AI-enabled adversaries increased attacks by 89% year-over-year, while according to the 2025 Armis Cyberwarfare Report, 66% of organizations surveyed agree that organizations continue to underestimate the resources required to defend against AI-powered threats, revealing a widening gap between perceived preparedness and actual resilience.
- The rise of (non-human) identities is expanding the threat landscape. As the number of machine identities surges, we are observing an evolution in the threat landscape. According to IDC, on average enterprises have >100 machine identities for every one human identity in the ecosystem.
- Government and Regulation Governments continue to expand regulation, and we expect enforcement to increase. These regulations have elevated cybersecurity to a board-level mandate for organizations. The EU AI Act, set to become effective on August 2, 2026, is the world's first comprehensive, binding law regulating artificial intelligence, adopted by the European Union in 2024. Non-compliance with the EU AI Act's requirements for high-risk systems can result in fines of up to 3% of a company's global annual revenue (or €15 million, whichever is higher), underscoring the material financial consequences of inadequate AI governance.
- White House on June 2<sup>nd</sup> signed a scaled-back executive order requiring AI firms to provide models to the government 30 days pre-release for oversight and cybersecurity coordination, balancing national-security concerns with efforts to avoid slowing innovation.

### Platformization

- As companies look at their security needs for 2026 and beyond, beneficiaries will be AI-native platforms with data scale + ecosystem integration as opposed to just tools that address distinct security needs, in our opinion.

### Headwinds and Risk

- LLM providers internalize core security capabilities—such as threat detection, identity validation, and autonomous response—thereby compressing vendor differentiation and shifting value toward AI platform owners.

# Actionable Ideas

## Platforms

**CrowdStrike** - CrowdStrike has evolved from an endpoint vendor into the AI-native cybersecurity platform of choice for the agentic era, with Falcon serving as the consolidation point as enterprises rationalize fragmented security stacks. Through tools like Charlotte AI, AgentWorks, its Agentic Security Workforce, and initiatives like Project QuiltWorks, the company is positioning itself as the orchestration layer for agent-driven cybersecurity, enabling automated defense, vulnerability remediation, and a shift toward fully autonomous SOC operations. Being selected as one of only two cybersecurity vendors among the 12 launch partners in Anthropic's Project Glasswing signals that CrowdStrike is viewed as mission-critical infrastructure for securing frontier AI systems, not just another security tool.

**Cloudflare** - Cloudflare's competitive moat, scalable connect-and-protect architecture, and large market opportunity make this an attractive long-term investment. The developer platform (Workers, R2, AI Inference) is becoming a meaningful growth pillar as agentic AI workloads drive a new wave of edge compute and inference demand. Cloudflare is positioning itself as a foundational edge-based infrastructure layer for the agent-driven internet, offering an integrated AI security and compute platform that enables organizations to deploy, govern, and secure autonomous AI agents at scale across multi-cloud environments.

**Palo Alto Networks** - We see Palo Alto Networks as a key beneficiary of vendor consolidation given platform breadth, with the AI-native Cortex strategy positioning the company well for agentic SOC adoption as enterprises automate Tier 1 and Tier 2 workflows. By embedding agentic AI across its platform and consolidating fragmented tools, the company aims to deliver an integrated, identity- and cloud-centric architecture capable of securing increasingly complex, AI-driven attack surfaces while reducing manual security operations. The selection of Palo Alto Networks as a core participant in Anthropic's Project Glasswing highlights its role as a foundational security platform for the AI era, reinforcing the view that large-scale, integrated cyber platforms are central to governing, securing, and operationalizing frontier models like Claude Mythos as enterprises adopt increasingly autonomous AI systems.

**Zscaler** - Zscaler is evolving its Zero Trust Exchange into a core security and governance layer for agent-driven AI environments, providing visibility, policy enforcement, and data protection for autonomous systems operating at machine speed. Through its AI Security Suite, Secure Access to AI, and role in initiatives like Project Glasswing, Zscaler aims to enable safe enterprise AI adoption by combining real-time control, deep observability, and autonomous threat detection across both human and machine interactions.

## Actionable Ideas (cont'd)

---

### Data Security:

**Commvault** - Commvault has repositioned into a cyber resilience platform that is directly benefiting from AI tailwinds, as enterprises prioritize faster threat detection, automated response, and reduced recovery times in an increasingly complex ransomware environment. Its integration of AI—including emerging agentic capabilities to autonomously orchestrate detection and recovery workflows—along with broad multi-cloud and SaaS coverage, is driving demand for its platform as organizations seek scalable, intelligent resilience solutions that minimize downtime and operational burden.

**Rubrik** - Rubrik is well-positioned to benefit from accelerating AI tailwinds as enterprises prioritize data security and cyber resilience amid rising AI-enabled threats and ransomware activity. The company's ability to manage and secure data across on-prem, cloud, and SaaS environments—combined with expanding use cases in DSPM, AI data governance, and support for agentic workloads—positions it as a critical platform for protecting and operationalizing data in AI-driven environments. High recurring revenue, strong retention, and continued module expansion underpin a durable growth profile as organizations consolidate data protection, governance, and recovery onto a unified, intelligent platform.

### Identity Security:


**Okta** – We see Okta as a leading identity platform, with workforce identity providing a stable revenue base while customer identity (CIAM via Auth0) and newer Identity Governance and Privileged Access offerings drive the next leg of growth. As enterprises adopt agentic AI and manage a surge in nonhuman identities, Okta's Identity Security Posture Management, Auth for GenAI, Okta for AI Agents, and broader machine identity capabilities position the company to secure and orchestrate identity across increasingly autonomous, AI-driven environments, enabling it to benefit directly from this structural shift.

**SailPoint**- We see SailPoint as a leader in identity governance, with its core identity security platform providing a strong foundation as enterprises shift toward more dynamic, AI-driven environments. As agentic AI adoption accelerates and nonhuman identities proliferate, SailPoint's strengths in identity lifecycle management, access governance, and machine identity control position it to benefit from the need to securely manage, provision, and audit autonomous agents at scale—driving increased demand as identity becomes a central control plane for AI-enabled enterprises.

## Cybersecurity TAM is Big and Growing

### A Large and Growing TAM: \$300Bn+ by 2030

2025: ~\$213Bn, 5-YR CAGR ~7%



Source: Truist Securities Research

# LLM Model Families

| Developer                                                                           | Model Family       | Associated Versions                                                                                                                                                                                                                                 |
|-------------------------------------------------------------------------------------|--------------------|-----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
|    | GPT-5 / o-Series   | <ul style="list-style-type: none"> <li>GPT-5.5 Pro</li> <li>GPT-5.5 Reasoning</li> <li>GPT-5.4 Thinking</li> <li>GPT-5.4 Mini/Nano</li> <li>GPT-5.3 Instant</li> <li>O3/O1</li> <li><b>GPT-5.5 Cyber (Daybreak) Experimental Preview</b></li> </ul> |
| <b>ANTHROPIC</b>                                                                    | Claude 4           | <ul style="list-style-type: none"> <li>Claude 4.8 Opus</li> <li>Claude 4.7 Opus</li> <li>Claude 4.6 Sonnet/Opus</li> <li>Claude 4.5 Haiku</li> <li><b>Mythos (Experimental Preview)</b></li> </ul>                                                  |
|    | Gemini             | <ul style="list-style-type: none"> <li>Gemini 3.1 Pro</li> <li>Gemini 3.1 Flash/Flash Lite</li> <li>Gemini 3 Pro/Flash</li> <li>Gemini 2.5 Pro</li> </ul>                                                                                           |
|    | Grok               | <ul style="list-style-type: none"> <li>Grok-4.4 (Upcoming/1T)</li> <li>Grok-4.3 (Current)</li> <li>Grok-4.20</li> <li>Grok-3</li> <li>Grok 5 (In Training/10T)</li> </ul>                                                                           |
|   | Muse / Llama       | <ul style="list-style-type: none"> <li>Muse Spark (Llama Replacement)</li> <li>Llama 4 Scout (17B)</li> <li>Llama 4 Maverick (17B)</li> <li>Llama 3.1 / 3.0</li> </ul>                                                                              |
|  | Mistral / Devstral | <ul style="list-style-type: none"> <li>Mistral Large 3</li> <li>Mistral 3 (14B, 8B, 3B)</li> <li>Mistral Small 4</li> <li>Devstral Small 2 / Devstral 2</li> </ul>                                                                                  |
|  | DeepSeek V         | <ul style="list-style-type: none"> <li>DeepSeek V4</li> <li>DeepSeek V3.2</li> <li>DeepSeek R1 / R1-0528 (Reasoning)</li> </ul>                                                                                                                     |

Source: Truist Securities Research



# Mythos and Daybreak: Driving a Paradigm Shift in Cybersecurity

---

- **Frontier AI capability inflection:** Mythos (launched in April 2026) and Daybreak (launched in May 2026) represent a step-change in autonomous cybersecurity, demonstrating end-to-end AI-driven vulnerability discovery and exploit generation at scale—capabilities previously confined to elite human researchers.
- **Industrialization of Zero-Day discovery:** The model's ability to uncover thousands of high-severity, previously unknown vulnerabilities shifts the industry from scarcity to abundance, materially increasing the volume of exploitable flaws entering the ecosystem.
- **Compression of attack timelines:** Autonomous exploit development compresses the timeline from discovery to weaponization from weeks to hours, significantly elevating zero-day risk and rendering traditional patch cycles insufficient.
- **Shift to system-level risk (attack paths vs. single CVEs):** Mythos and Daybreak highlight a move toward chaining vulnerabilities into full attack paths, reducing the relevance of point-in-time scoring frameworks and increasing the importance of holistic attack surface management.
- **AI-driven arms race (offense vs. defense):** Dual-use capabilities necessitate rapid adoption of AI-enabled defensive tools (e.g., automated patching, continuous monitoring), accelerating the transition to machine-speed cyber defense architectures.
- **Governance and regulatory overhang:** Restricted deployment (e.g., Project Glasswing, Trusted Access for Cyber) underscores national security concerns, with implications for tighter access controls, regulatory scrutiny, and rising importance of AI governance frameworks—particularly in critical sectors like financial services.



## AI Creating and Discovering More Vulnerabilities Than Ever Before

More Exposures are being *Created* → and *Discovered*

- *This is accelerating software development + deployment*
- *New system = New entry point*
- *Attack surface expanding faster than teams can track*

- *Frontier Models finding vulnerabilities in minutes, not months*
- *Attackers gaining same ability*
- *Discovery that was impossible yesterday is trivial today*

*These dynamics are unfolding simultaneously across an increasingly expansive and interconnected attack surface.*

## Time To Exploit Has Compressed Significantly

**771 → 1.6**  
*Days to exploit*

From 2021 to today, the average time from disclosure to exploitation has fallen to 1.6 days

**500+**  
*Zero-days found*

These were discovered autonomously in open-source code by Opus 4.6

**Thousands**  
*of vulnerabilities found by Mythos*


*Discovered in weeks following the April 2026 launch of Mythos Preview*

**99%**  
*Remain unpatched*

*Vendor disclosure processes were not built for AI scale discovery volumes*

# AI Frontier Benchmark Comparisons

Frontier models continue to show rapid, broad-based gains across domains, with steady upward progress on security (exploit generation in CyberGym), software engineering (SWE-bench variants), system execution (Terminal-Bench), and general intelligence (HLE)—reinforcing that capability scaling is no longer siloed but compounding across coding, reasoning, and real-world task execution simultaneously, widening the gap vs. prior-generation models and smaller competitors.



Source: Truist Securities Research

# LLM Efficiency and Economics

OpenAI and Anthropic are improving their ability to detect software flaws, but it comes with a price.

| Company | AI Model        | Steps Completed (per 100M tokens) | Unit Cost (per 100M tokens) |
|---------|-----------------|-----------------------------------|-----------------------------|
|         | Mythos preview  | 25                                | \$12,500                    |
|         | Claude Opus 4.7 | 11                                | \$2,500                     |
|         | Claude Opus 4.6 | 16                                | \$2,500                     |
|         | GPT-5.5-Cyber   | 23                                | \$3,000                     |
|         | GPT-5.4         | 14                                | \$1,500                     |
|         | Codex-5.3       | 11                                | \$1,400                     |

Source: Truist Securities Research

# AI Frontier Models Not Displacing Cybersecurity Vendors

- **Pre-production code scanning ≠ post-production runtime security:** While LLMs are increasingly effective at static analysis of publicly available code (e.g., identifying vulnerabilities, insecure patterns, or misconfigurations before deployment), this is fundamentally different from securing live infrastructure in production environments. Runtime security requires continuous monitoring of dynamic workloads, identity behavior, network traffic, and real-time threat detection/response, often under highly variable conditions (cloud, containers, ephemeral compute). It also demands deterministic enforcement (blocking, isolation, remediation), low-latency decisioning, and integration with production systems (SIEM, EDR, identity, cloud control planes)—capabilities that go well beyond pattern recognition in code. As a result, even as frontier models improve pre-deploy scanning and developer productivity, they do not substitute for the always-on, operational, and enforcement-heavy layer that defines modern cybersecurity platforms.
- **Identity is the #1 threat vector:** 80%-90% of cyber attacks stem from compromised credentials and privilege abuse—not novel exploits—allowing attackers to operate as legitimate users. Unlike static code flaws, identity risk is dynamic (who/what/where/when) and requires continuous monitoring of behavior, lateral movement, and escalation. While frontier models can assist with analysis, they lack the deterministic, real-time enforcement (auth controls, session revocation, least-privilege) needed in production—reinforcing the central role of specialized identity and security platforms.
- **Integration complexity favors incumbents:** Security stacks (SIEM, EDR, identity, network) are deeply embedded; vendors own data pipelines, telemetry, and workflows, making them the natural layer to incorporate AI rather than be displaced by it.
- **Trust, liability, and compliance barriers:** Organizations require auditable, deterministic, and regulator-approved controls, which are delivered through established vendors with SLAs—not standalone frontier models, which are probabilistic and non-deterministic, making them difficult to validate, reproduce, and certify in regulated environments. By contrast, enterprises require auditable, deterministic controls with clear accountability, logging, and SLAs, especially for incident response, access control, and policy enforcement. This gap means organizations continue to rely on established cybersecurity vendors that can wrap AI inside governed, testable, and compliant systems, ensuring explainability, repeatability, and regulatory approval (e.g., SOC 2, FedRAMP), rather than trusting raw model outputs in isolation.
- **Adversarial dynamics increase demand:** As frontier models also enhance attacker capabilities, enterprises need more—not fewer—defensive tools, reinforcing spend with security vendors.
- **AI becomes a feature, not a replacement:** Cyber vendors are rapidly embedding frontier models into their platforms, capturing the productivity benefits while retaining customer ownership and economic value.
- **Google/Wiz Acquisition Example:** Despite having one of the leading frontier models (Gemini), Google still paid \$32Bn to acquire Wiz (closed March 2026), the leading cloud security vendor, underscoring that LLMs do not replace the security stack—they augment it. The move highlights that enterprises value integrated, operational security platforms (visibility, enforcement, runtime protection) over standalone model capabilities, with LLMs positioned as an input into security workflows rather than a substitute for vendors.

## Partnerships the Way Forward

---

- **Structured access frameworks (Glasswing, TAC) signal “AI-inside-the-stack,” not replacement of it:** Anthropic’s Project Glasswing and OpenAI’s Trusted Access for Cyber (TAC) both deliberately gate frontier model capabilities behind vetted partnerships, identity controls, and auditable environments, rather than exposing them as standalone tools. This reflects a recognition that deploying high-capability cyber models requires integration with existing security infrastructure (SIEM, identity, endpoint, cloud control planes), governance, and human-in-the-loop workflows. In practice, these programs position frontier models as co-pilots embedded within the incumbent cyber stack, reinforcing the role of established vendors as the control, enforcement, and distribution layer, rather than disintermediating them.
- **Project Glasswing (Anthropic)** is a controlled-access initiative (launched April 7, 2026) that grants select, vetted organizations exclusive access to Claude Mythos Preview strictly for defensive use, with no commercial release and no public availability outside the program. It enforces gated distribution, identity controls, and auditable usage, underscoring that even frontier cyber models are deployed through trusted partners and existing security stacks—not as standalone tools—reinforcing a partnership-driven model rather than disintermediation of the cyber ecosystem.
- **Project Glasswing** consortium, originally comprised of ~50 vetted partners (now expanded to ~200), which include 12 founding partners, including CrowdStrike and Palo Alto Networks in cybersecurity. This core group consists of high-tier technology, finance, and security firms that are deeply integrated into the "offensive-for-defense" testing of the Mythos engine. These additional members include government agencies, sovereign entities, and critical infrastructure providers that remain largely unnamed for security reasons. Cloudflare and Zscaler have publicly disclosed their participation in Project Glasswing.
- On May 18<sup>th</sup>, Anthropic announced it is revising its earlier position to allow users of its Mythos cybersecurity model to disclose their involvement in Glasswing and, at their own discretion, share findings, best practices, tools or code developed through the program.

## Partnerships the Way Forward (cont'd)

---

- On June 2<sup>nd</sup>, Trump administration signed an executive order increasing government oversight of models before they are released (30 days before release, down from 90 days earlier contemplated). Additionally, the administration approved Anthropic's expansion of Mythos access from ~50 to ~200 global organizations to help identify and fix software vulnerabilities, reversing earlier security concerns.
- On May 11<sup>th</sup>, OpenAI launched **Daybreak**, OpenAI's answer to Anthropic's Project Mythos. The effort is built on the foundations of three models: GPT-5.5 (which has standard safeguards for general purpose use), GPT-5.5 with **Trusted Access for Cyber** (for verified defensive work in authorized environments), and GPT-5.5-Cyber (a permissive model for red teaming, penetration testing, and controlled validation).
- **Trusted Access for Cyber (TAC)** is OpenAI's controlled deployment framework for frontier AI in cybersecurity, providing vetted, tiered access with identity verification, auditability, and use-case restrictions, ensuring models like Daybreak are safely embedded within existing security platforms rather than used as standalone tools.
- In contrast to Project Glasswing, which is a highly restricted, invite-only consortium focused on securing the world's most critical software, TAC is a scalable, identity-vetted program designed to democratize access for the broader cybersecurity community.
- CrowdStrike's **Project Quiltworks**, announced in April, and subsequently expanded in May, is an industry-wide coalition of ecosystem partners to assess, prioritize, and continuously remediate the wave of vulnerabilities in production code now being discovered by frontier AI models. Project Quiltworks acts as a "Remediation Ecosystem" designed to operationalize AI-driven patching across the global enterprise tech stack. Project QuiltWorks includes the world's leading systems integrators and frontier AI models from OpenAI and Anthropic to deliver assessments, board-level risk reporting, adversary-informed prioritization, and guided remediation.




## Recent AI Cyber-Incidents 2024-2026

| Date      | Company                                                                             | Attack Type                | Impact / Sophistication                                                                                                                             |
|-----------|-------------------------------------------------------------------------------------|----------------------------|-----------------------------------------------------------------------------------------------------------------------------------------------------|
| Feb 2024  |    | Deepfake CFO Impersonation | <b>\$25.6MM</b> stolen via a multi-person deepfake video conference.                                                                                |
| Aug 2025  |    | AI-Powered Ransomware      | First documented ransomware using local LLMs to generate encryption scripts on the fly.                                                             |
| Sept 2025 |    | AI-Driven Extortion        | 95,000 user accounts breached; hackers threatened to feed data into AI training sets.                                                               |
| Sept 2025 |    | Jailbreak Cyber Espionage  | ~30 High-value global organizations, including major technology companies, financial institutions, chemical manufacturers, and government agencies. |
| Jan 2026  |   | AI-Assisted Data Breach    | Potentially <b>7.2MM</b> accounts leaked; linked to Everest group's automated exploitation tools.                                                   |
| Apr 2026  |  | Autonomous AI Worm         | Infected <b>1,100+ repositories</b> in 3 hours by autonomously discovering/exploiting zero-days.                                                    |

Source: Truist Securities Research


# Generative AI is the New Strategic Battleground

**AI is a double-edged sword.** Artificial Intelligence (AI) has been used to enhance cybersecurity tools for the last decade. However, generative AI is a force multiplier for cyber adversaries. The integration of artificial intelligence by attacker campaigns is escalating the complexity of cyber threats. The deployment of Generative AI by threat actors to orchestrate complex attacks against enterprise cloud environments underscores the critical role security needs to play in the cloud and application development environments. At the same time, AI is transforming cybersecurity by enabling a more proactive and comprehensive approach to threat detection and response. Security posture should improve with the increase in automation and enhanced vulnerability scanning for code.



Source: Truist Securities Research

# Top Agentic Threats



Source: Palo Alto Networks, Truist Securities Research

# Agentic Threat Vectors in a Gen AI World

| Attack Vector Category    | Threat Description                                                                                  | Public Leaders                                             | Private Innovators                 |
|---------------------------|-----------------------------------------------------------------------------------------------------|------------------------------------------------------------|------------------------------------|
| Agentic AI & Authority    | Autonomous agents executing unauthorized actions (deleting data, sending funds) at machine speed.   | CrowdStrike; Zscaler                                       | Glean; Zero Networks               |
| Prompt Injection          | Hijacking LLM behavior via malicious inputs or hidden external data to bypass safety controls.      | CrowdStrike; Palo Alto; Google (Wiz); Zscaler; SentinelOne | Prompt.Security                    |
| Data & Model Poisoning    | Corrupting training data or weights to embed backdoors or bias outputs before deployment.           | Palo Alto; Zscaler; CrowdStrike                            | HiddenLayer                        |
| Supply Chain & MCP        | Compromised third-party models and Model Context Protocol (MCP) servers introducing hidden threats. | Zscaler; Cloudflare                                        | MintMCP; Cystore                   |
| Shadow AI                 | Unsanctioned AI tools and agents deployed by employees outside IT visibility.                       | CrowdStrike; SailPoint; Okta                               | Cyera                              |
| AI-Generated Code Flaws   | Production codebases scaling insecure or exploitable AI-authored code without review.               | CrowdStrike; Palo Alto                                     | Snyk; Winfunc                      |
| Sensitive Data Exposure   | LLMs leaking PII or proprietary data through outputs or retrieval (RAG) pipelines.                  | Zscaler, Varonis, Rubrik, Commvault                        | Cyera; Veeam; Sentra               |
| Deepfakes & Synth ID      | High-fidelity impersonation (voice/video) used to bypass traditional identity verification.         | CrowdStrike; Okta                                          | Abnormal AI; SandboxAQ             |
| Machine Identities        | Sprawl of machine/agentic identities operating without human oversight or governed access paths.    | SailPoint; Palo Alto (CyberArk), Okta, CrowdStrike         | Ping; BeyondTrust; Delinea         |
| Session Hijacking         | Bypassing MFA via token theft or cookie hijacking to move laterally between IdPs.                   | Okta; Zscaler                                              | Beyond Identity; Adaptive Security |
| Post-Quantum Cryptography | "Harvest-now-decrypt-later" strategies targeting data to be decrypted once quantum is viable.       | Palo Alto, Google                                          | SandboxAQ; QuSecure                |

Source: Truist Securities Research

# What is Agentic Security?

Agentic AI security is the protection of AI agents that can plan, act, and make decisions autonomously. It focuses on securing the agent's reasoning, memory, tools, actions, and interactions so that autonomy does not create new paths for misuse. It addresses risks that arise only when systems execute multi-step tasks, use external tools, or collaborate with other agents.

AI agents introduce risks through autonomous reasoning, tool execution, memory updates, privilege inheritance, and inter-agent communication. These surfaces can be manipulated, poisoned, or redirected, creating pathways for unintended actions, goal drift, privilege misuse, or coordinated influence across multi-agent systems. By applying controls to the internal workflow: reasoning, planning, tool use, memory, and coordination, one can help secure AI agents.

## How do you secure AI agents?

Controls are applied to the agent's internal workflow, placing constraints and validation on reasoning, planning, tool use, memory, and coordination.

## Security concerns of AI agents


Autonomous behavior can be manipulated, leading to unintended actions, privilege misuse, goal drift, or inter-influence.

## Privacy and data security

Agents accessing, storing, or reusing sensitive information create privacy risks, while their data handling raises unauthorized access concerns.

## Detecting or managing manipulation

Anomalies in an agent's decision-making and interactions may reveal compromise, requiring consistent controls across platforms.



Source: Truist Securities Research

# Agentic Security Stack

| Category            | Description                                                                    | Representative Vendors                                        |
|---------------------|--------------------------------------------------------------------------------|---------------------------------------------------------------|
| AI-IAM / NHI        | Managing agent identities, permissions, and delegated access at runtime.       | Palo Alto Networks (CyberArk), SailPoint, Okta, CrowdStrike   |
| Agentic Guardrails  | Real-time monitoring to block malicious intent or unauthorized tool calls.     | Zscaler (AI Guard), CrowdStrike (Pangea), Cloudflare          |
| AI-SPM              | Discovering shadow AI agents and identifying misconfigured entitlements.       | Wiz, Palo Alto (Prisma AIRS), CrowdStrike, Zscaler            |
| Agentic SOC         | Automating triage and investigation using agents to counter AI-driven attacks. | CrowdStrike (Falcon), Zscaler (Red Canary), Palo Alto (XSIAM) |
| Data Security       | Discovering and securing sensitive training/inference data (DLP and vaulting). | Commvault, Rubrik, Varonis, Zscaler                           |
| Exposure Management | Inventorying models, measuring control drift, and prioritizing remediations.   | CrowdStrike, Tenable, Qualys, Rapid7, Zscaler                 |
| AI Gateway & Edge   | Enforcing Zero Trust and abuse prevention at the global edge for LLM traffic.  | Cloudflare                                                    |

Source: Truist Securities Research

# Valuation up from the Bottom


IGV EV/NTM Sales YTD



Source: Truist Securities Research and FactSet

# Relative Performance Overview

## IGV Annual Performance Relative to S&P 500



Source: Truist Securities Research and FactSet



## Companies Mentioned in This Note

**CrowdStrike Holdings, Inc.** (CRWD, \$671.02, Buy, Junaid Siddiqui)  
**Commvault Systems, Inc.** (CVLT, \$118.68, Buy, Junaid Siddiqui)  
**Alphabet Inc.** (GOOGL, \$368.53, Buy, Youssef Squali)  
**Cloudflare, Inc.** (NET, \$250.11, Buy, Junaid Siddiqui)  
**Okta, Inc.** (OKTA, \$118.72, Buy, Junaid Siddiqui)  
**Palo Alto Networks, Inc.** (PANW, \$272.05, Buy, Junaid Siddiqui)  
**Rubrik, Inc.** (RBRK, \$73.41, Buy, Junaid Siddiqui)  
**SailPoint, Inc.** (SAIL, \$18.24, Buy, Junaid Siddiqui)  
**Zscaler, Inc.** (ZS, \$130.78, Buy, Junaid Siddiqui)  
**Anthropic** (Private)  
**OpenAI** (Private)

## Analyst Certification

I, Junaid Siddiqui, hereby certify that the views expressed in this research report accurately reflect my personal views about the subject company(ies) and its (their) securities. I also certify that I have not been, am not, and will not be receiving direct or indirect compensation in exchange for expressing the specific recommendation(s) in this report.

I, Miller Jump, hereby certify that the views expressed in this research report accurately reflect my personal views about the subject company(ies) and its (their) securities. I also certify that I have not been, am not, and will not be receiving direct or indirect compensation in exchange for expressing the specific recommendation(s) in this report.

## Required Disclosures

Analyst compensation is based upon stock price performance, quality of analysis, communication skills, and the overall revenue and profitability of the firm, including investment banking revenue.

As a matter of policy and practice, the firm prohibits the offering of favorable research, a specific research rating or a specific target price as consideration or inducement for the receipt of business or compensation. In addition, associated persons preparing research reports are prohibited from owning securities in the subject companies.

Charts indicating changes in ratings can be found in recent notes and/or reports at our website or by contacting Truist Securities. Please see our disclosures page for more complete information at <https://truist.bluematrix.com/sellside/Disclosures.action>

## Important Disclosures on Equity Research Dissemination, Ratings, Designations, and Coverage Universe

### Dissemination of Research

Truist Securities, Inc. ("Truist Securities") seeks to make all reasonable efforts to provide research reports simultaneously to all eligible clients. Reports are available as published in the restricted access area of our website to all eligible clients who have requested a password. Institutional investors, corporates, and members of the Press may also receive our research via third party vendors including: Thomson Reuters, Bloomberg, FactSet, and S&P Capital IQ. Additional distribution may be done by sales personnel via email, fax, or other electronic means, or regular mail.

For access to third party vendors or our Research website: <https://truistresearch.bluematrix.com/client/library.jsp>

Please email the Research Department at [EquityResearchDepartment@truist.com](mailto:EquityResearchDepartment@truist.com) or contact your Truist Securities sales representative.

### Truist Securities Rating System for Equity Securities

Truist Securities, Inc. ("Truist Securities") rates individual equities using a three-tiered system. Each stock is rated relative to the broader market (generally the S&P 500) over the next 12-18 months (unless otherwise indicated).

**Buy (B)** – the stock’s total return is expected to outperform the S&P 500 or relevant benchmark over the next 12-18 months (unless otherwise indicated)

**Hold (H)** – the stock’s total return is expected to perform in line with the S&P 500 or relevant benchmark over the next 12-18 months (unless otherwise indicated)

**Sell (S)** – the stock’s total return is expected to underperform the S&P 500 or relevant benchmark over the next 12-18 months (unless otherwise indicated)

**Not Rated (NR)** – Truist Securities does not have an investment rating or opinion on the stock

**Coverage Suspended (CS)** – indicates that Truist Securities’ rating and/or target price have been temporarily suspended due to applicable regulations and/or Truist Securities Management discretion. The previously published rating and target price should not be relied upon.

Truist Securities analysts have a price target on the stocks that they cover, unless otherwise indicated. The price target represents that analyst's expectation of where the stock will trade in the next 12-18 months (unless otherwise indicated). If an analyst believes that there are insufficient valuation drivers and/or investment catalysts to derive a positive or negative investment view, they may elect with the approval of Truist Securities Research Management not to assign a target price; likewise certain stocks that trade under \$5 may exhibit volatility whereby assigning a price target would be unhelpful to making an investment decision. As such, with Research Management's approval, an analyst may refrain from assigning a target to a sub-\$5 stock.

## Legend for Rating and Price Target History Charts:

B = Buy

H = Hold

S = Sell

D = Drop Coverage

CS = Coverage Suspended

NR = Not Rated

I = Initiate Coverage

T = Transfer Coverage

Truist Securities ratings distribution (as of 06/08/2026):

| Coverage Universe |       |         | Investment Banking Clients Past 12 Months |       |         |
|-------------------|-------|---------|-------------------------------------------|-------|---------|
| Rating            | Count | Percent | Rating                                    | Count | Percent |
| Buy               | 483   | 66.44%  | Buy                                       | 112   | 23.19%  |
| Hold              | 241   | 33.15%  | Hold                                      | 44    | 18.26%  |
| Sell              | 3     | 0.41%   | Sell                                      | 1     | 33.33%  |

## Other Disclosures

Certain information contained herein has been derived from third-party sources believed to be reliable but is not guaranteed as to accuracy and does not purport to be a complete analysis of the security, company or industry involved. This report is not to be construed as an offer to sell or a solicitation of an offer to buy any security. Truist Securities, Inc. and/or its officers or employees may have positions in any securities, options, rights or warrants. The firm and/or associated persons may sell to or buy from customers on a principal basis. Investors may be prohibited in certain states from purchasing some over the-counter securities mentioned herein. Opinions expressed are subject to change without notice.

Truist Securities, Inc.'s research is primarily provided to and intended for use by Institutional Accounts as defined in FINRA Rule 4512(c). The term "Institutional Account" shall mean the account of: (1) a bank, savings and loan association, insurance company or registered investment company; (2) an investment adviser registered either with the SEC under Section 203 of the Investment Advisers Act or with a state securities commission (or any agency or office performing like functions); or (3) any other person (whether a natural person, corporation, partnership, trust or otherwise) with total assets of at least \$50 million. In addition, certain affiliates of Truist Securities, Inc., including Truist Investment

Services, Inc. (an SEC registered broker-dealer and a member of FINRA, SIPC) and Truist Advisory Services, Inc. (an investment adviser registered with the SEC), may make Truist Securities, Inc. research available, upon request, to certain of their clients from time to time.

Truist Securities, Inc. is a registered broker-dealer and a member of FINRA and SIPC. It is a service mark of Truist Financial Corporation. Truist Securities, Inc. is owned by Truist Financial Corporation and affiliated with Truist Investment Services, Inc. Despite this affiliation, securities recommended, offered, sold by, or held at Truist Securities, Inc. or Truist Investment Services, Inc. (i) are not insured by the Federal Deposit Insurance Corporation; (ii) are not deposits or other obligations of any insured depository institution (including Truist Bank); and (iii) are subject to investment risks, including the possible loss of the principal amount invested. Truist Bank may have a lending relationship with companies mentioned herein. Certain clients may compensate Truist Securities, Inc. for research via hard dollar payments, and Truist Securities, Inc. may be deemed to be an investment adviser to such clients as a result of such payments.

Please see our Disclosure Database to search by ticker or company name for the current required disclosures, including valuation and risks. Link: <https://truist.bluematrix.com/sellside/Disclosures.action>

Please visit the Truist Securities Equity Research Portal for coverage universe, current reports, and analyst roster with contact information. Link: <https://truist-portal.bluematrix.com>

Truist Securities, Inc., member FINRA and SIPC. Truist, Truist Bank, Truist Securities, Truist Investment Services, and Truist Advisory Services are service marks of Truist Financial Corporation.

If you no longer wish to receive this type of communication, please request removal by sending an email to [EquityResearchDepartment@truist.com](mailto:EquityResearchDepartment@truist.com)

© Truist Securities, Inc. 2026. All rights reserved. Reproduction or quotation in whole or part without permission is forbidden. Without Truist Securities' prior written consent, the contents of this report may not be used, in whole or in part, in connection with the training, development, fine-tuning or operation of a machine learning or artificial intelligence system, tool or application.

ADDITIONAL INFORMATION IS AVAILABLE at our website, [TruistSecurities.com](https://TruistSecurities.com), or by writing to: Truist Securities, Research Department, 740 Battery Avenue SE, 3rd Floor, Atlanta, GA 30339-3107