

China AI

Initiation: big enough for multiple winners

- Capability wins the price premium, cost efficiency wins adoption and scale – we see coexistence rather than price convergence
- ARR keeps overtaking published forecasts; we size China's AI application TAM bottom-up and against a USD26.5tn global pool
- Initiate at Positive; as funding and distribution replace model capability on binding constraints, we prefer Alibaba, Z.AI (Zhipu), and Minimax

Negative

Neutral

Positive


Positive
 (initiation)

John Choi

(852) 2773 8730

john.choi@hk.daiwacm.com

Lily Liu

(852) 2848 4435

lily.liu@hk.daiwacm.com



Investment case: We initiate coverage of the China AI sector with a Positive rating, with a specific focus on foundation models. As model capability improves and the cost of intelligence falls, the industry is shifting from model benchmark competition towards mass model consumption. Commercialisation is already advancing faster than anticipated – OpenAI and Anthropic's combined July 2026 annual recurring revenue (ARR) exceeded MiniMax's late-2025 estimate for the entire 2026 global LLM market by more than two-fold. While global AI TAM is projected to reach **USD26.5tn** by SpaceX, in China, our base case estimates AI coding TAM to rise to CNY264bn in 2030E (**266% CAGR**), and consumer AI TAM to reach CNY116bn (**154% CAGR**), with substantially larger long-term opportunities across enterprise and multimodal applications.

Near-term evidence suggests that both premium capability and cost efficiency can be rewarded. High-value coding, research, and agent workloads continue to pay for stronger intelligence, while cheaper models are gaining token share by making AI economical for a broader group of developers, SMEs, and enterprises. Z.AI represents the capability route: its repeated price increases alongside rising usage demonstrate that frontier-level performance can command a premium. MiniMax represents the efficiency route: competitive intelligence, multimodal breadth and lower serving costs position it for high-volume adoption. With Qwen 3.8-Max catching up on capability, Alibaba is combining both, even with proprietary ASIC, cloud ecosystem, and funding capacity. Benchmark leadership will continue to change hands, but release cadence, inference economics, product distribution, and the conversion of adoption into recurring revenue should prove more durable. We therefore expect a **multi-winner market**, with foundation-model providers capturing different customer groups and value pools as AI penetration deepens.

Catalysts: Near-term: Alibaba's stronger-than-expected AI cloud earnings, Z.AI's new model release, ARR conversion and margin improvement, and Minimax's successful new-model launches driving user and monetisation traction. **Long-term:** broader agent adoption and declining inference costs expanding AI consumption across enterprise and consumer workflows.

Valuation: We initiate coverage of China AI with a Positive rating, with a focus on foundation model providers. We initiate coverage on Z.AI and Minimax with Buy (1) calls, while Alibaba is our top pick for its AI cloud earnings. Our sector pecking order is Alibaba > Z.AI ≥ MiniMax.

Risks: Slower-than-expected AI adoption and monetisation; fiercer price competition; deteriorating margin profile; and a slowdown in CSP capex.

Key stock calls

	New	Prev.
Alibaba Group (BABA US)		
Rating	Buy	Buy
Target	180.00	180.00
Upside	▲ 36%	
Alibaba Group - H (9988 HK)		
Rating	Buy	Buy
Target	175.00	175.00
Upside	▲ 38.4%	
Z.AI (2513 HK)		
Rating	Buy	
Target	1,500.00	
Upside	▲ 25.8%	
Minimax Group (100 HK)		
Rating	Buy	
Target	530.00	
Upside	▲ 61.6%	

Source: Daiwa forecasts

Negative	Neutral	Positive	How do we justify our view?		
★	Positive (initiation)		Growth outlook	Valuation	Earnings revisions
			✓ ✓ ✓ ✓ ✓	✓ ✓ ✓ ✓ ✓	✓ ✓ ✓ ✓ ✓

Growth outlook



We size China's AI opportunity from first principles rather than published forecasts, as published forecasts have been overtaken every year since 2023. On a bottom-up basis, we expect China's AI coding TAM to expand to CNY264bn by 2030E (2025-30E CAGR of 266%), and Consumer AI TAM to CNY116bn by 2030E (2025-20E CAGR of 154%). Above these sit two much larger pools that AI can mediate rather than capture: c.CNY10tn of enterprise digital-economy activity and c.CNY4tn of the cultural new economy. The market is expanding faster than any LLM company can absorb, which is the key driver of our Positive sector call. In turn, we believe the market will be large enough to absorb multiple winners.

China AI: TAM and addressable market forecasts

	2025	2030E	2025-30E CAGR
AI Coding (TAM, CNYbn)	0.4	264	266%
Consumer AI applications (TAM, CNYbn)	1.1	116	154%
Enterprise applications (Addressable economic pool, CNYtn)	10		
Multi-modality applications (Addressable economic pool, CNYtn)	4		

Source: Daiwa estimates and forecasts

Valuation



AI monetisation remains at an early stage, with a rapidly evolving industry landscape and few listed pure-play LLM companies to provide meaningful valuation benchmarks. We believe ARR is the most relevant metric for pure-play LLM providers and thus value Z.AI and MiniMax using discounted EV/ARR multiples of 15x and 5x, respectively, applied to their respective average 2030E and 2029E ARR. These are well below the c.65x implied peer average, reflecting limited benchmark quality, differing business positioning and conservative assumptions on financial visibility. We value Alibaba using the SOTP approach, which implies a c.19x 12-month forward PER. We like Alibaba because it offers AI exposure with limited explicit AI premium despite the rapid growth of its AI-cloud business.

EV/ARR valuation comparison table

Company	Valuation date	Valuation (USDbn)	Implied EV/ARR (x)
OpenAI	Oct-2024	157	42.4x
OpenAI	Mar-2025	300	50.0x
OpenAI	Mar-2026	852	34.1x
OpenAI	Jun-2026	1,000	25.0x
Anthropic	Mar-2025	62	73.8x
Anthropic	Sep-2025	183	43.9x
Anthropic	Feb-2026	380	32.6x
Anthropic	May-2026	965	24.6x
Anthropic	Jun-2026	965	15.6x
Moonshot	May-2026	20	200.0x
Moonshot	Jul-2026	35	116.7x
Moonshot	Aug-2026	50	166.7x
Moonshot	Aug-2026	50	62.5x
DeepSeek	Jun-2026	50	227.3x
DeepSeek	Jul-2026	71	157.8x
Average of the latest			65.2x


Source: Daiwa

Earnings revisions



Bloomberg revenue consensus for both Z.AI and MiniMax have continued to rise, indicating strengthening confidence in AI demand and monetisation. The upward revisions accelerated particularly from April for Z.AI and from June for MiniMax, with their 2026E revenue growth forecasts reaching c.480% and 320%, respectively, YTD to late-July. Meanwhile, earnings forecasts have continued to decline as investment in model R&D, computing capacity, and market expansion remains elevated. However, the magnitude of upward revision on revenue is much higher than downward revision on earnings, implying improving scaling economics. We therefore see a favourable direction, with sustained top-line upgrades and stabilising earnings revisions.

Z.AI & Minimax: revenue forecasts revision since listing




Source: Bloomberg

Sector stocks: key indicators

		Share	Rating		Target price (local curr.)			EPS (local curr.)					
								FY1			FY2		
Company Name	Stock code	Price	New	Prev.	New	Prev.	% chg	New	Prev.	% chg	New	Prev.	% chg
Alibaba Group	BABA US	132.32	Buy	Buy	180.00	180.00	0.0%	45.989	45.989	0.0%	65.009	65.009	0.0%
Alibaba Group - H	9988 HK	126.40	Buy	Buy	175.00	175.00	0.0%	5.861	5.861	0.0%	8.542	8.542	0.0%
Baidu	BIDU US	109.50	Buy	Buy	175.00	175.00	0.0%	45.902	45.902	0.0%	49.099	49.099	0.0%
Baidu - H	9888 HK	104.60	Buy	Buy	170.00	170.00	0.0%	5.738	5.738	0.0%	6.137	6.137	0.0%
GDS Holdings	GDS US	31.41	Buy	Buy	48.00	48.00	0.0%	(1.208)	(1.208)	n.a.	2.235	2.235	0.0%
GDS Holdings - H	9698 HK	30.46	Buy	Buy	47.00	47.00	0.0%	(1.208)	(1.208)	n.a.	2.235	2.235	0.0%
Kingdee International Software Group	268 HK	8.99	Buy	Buy	19.50	19.50	0.0%	0.014	0.014	0.0%	0.107	0.107	0.0%
Kingsoft Cloud	KC US	11.96	Buy	Buy	21.00	21.00	0.0%	(2.690)	(2.690)	n.a.	(2.956)	(2.956)	n.a.
Kingsoft Cloud - H	3896 HK	6.17	Buy	Buy	11.00	11.00	0.0%	(0.179)	(0.179)	n.a.	(0.197)	(0.197)	n.a.
Kuaishou Technology	1024 HK	42.64	Buy	Buy	72.00	72.00	0.0%	3.590	3.590	0.0%	4.470	4.470	0.0%
Meitu	1357 HK	4.76	Buy	Buy	10.00	10.00	0.0%	0.205	0.205	0.0%	0.272	0.272	0.0%
Minimax Group	100 HK	328.00	Buy		530.00			(1.067)			(0.955)		
Phancy Group	6682 HK	29.62	Buy	Buy	47.00	47.00	0.0%	0.063	0.063	0.0%	0.461	0.461	0.0%
Tencent Holdings	700 HK	470.80	Buy	Buy	670.00	670.00	0.0%	29.263	29.263	0.0%	32.095	32.095	0.0%
Xiaomi	1810 HK	26.66	Buy	Buy	32.00	32.00	0.0%	0.999	0.999	0.0%	1.516	1.516	0.0%
Z.AI	2513 HK	1,192.00	Buy		1,500.00			(13.572)			(17.026)		


Source: Bloomberg, Daiwa forecasts

Historical ARR trends: OpenAI and Anthropic (USDm)




Source: The Information, Reuters, CNBC, TickerTrends, company disclosures; Daiwa estimates
 Note: ARR is an annualised run-rate, not recognised revenue

Historical ARR trends: Z.AI and MiniMax (USDm)




Source: Company, Daiwa estimates

China model market share by total usage volume on OpenRouter



Source: OpenRouter, compiled by Daiwa

Artificial Analysis Intelligence Index – highest scores achieved by each AI lab, highlighting Qwen, Z.AI and Minimax (as of 10 August 2026)



Source: Artificial Analysis, compiled by Daiwa

When a report covers six or more subject companies please access important disclosures for Daiwa Capital Markets Hong Kong Limited at http://www.hk.daiwacm.com/research_disclaimer.html or contact your investment representative or Daiwa Capital Markets Hong Kong Limited at Level 26, One Pacific Place, 88 Queensway, Hong Kong

Table of contents

Premium vs. Efficiency: both co-exist and are relevant.....	5
Capability will continue to win the price premium	5
Cost efficiency will be the key to adoption and scalability.....	6
ARR trends: the AI market continues to outgrow its own forecasts.....	10
US: still compounding, but the slope weakened after June	10
China: faster growth from a smaller base	11
Volume trends: token volumes surge, but leadership remains fluid	13
Usage leadership changes within weeks	13
Durable advantages may form in ecosystem development	14
Pricing trends: frontier model premium vs. falling price of intelligence	15
LLM pricing is both increasing and decreasing	15
Deep dive on capability vs. efficiency: discussion on scaling impact and activation ratios	16
China AI TAM analysis: scenario-based framework	19
1. AI Coding: monetising developer productivity	19
2. Consumer AI revenues: penetration and ARPU expansion	20
3. Enterprise applications: anchoring AI to the digital economy	21
4. Multi-modality: tapping into China's content economy	22
Linking the China framework to SpaceX's global AI TAM.....	24
From better models to a larger intelligence market.....	24
Why we like Z.AI and MiniMax.....	26
Risks	26
Company Section	
Z.AI	28
Minimax Group.....	56

Premium vs. Efficiency: both co-exist and are relevant

Following rapid advances in model capability and cost efficiency, the global LLM industry is moving from model development towards broader application, agent and enterprise adoption. Currently, the key debate in the market is whether the premier frontier models will continue to outperform the open-source models. In this section, we share our view and key analysis.

Capability will continue to win the price premium

We do not see this as a winner-takes-all contest. Frontier capability and cost efficiency address different customer needs and support two coexisting monetisation routes: leading models can command premium pricing in high-value workloads, while efficient models can drive broader adoption and greater usage volume.


We therefore believe AI pricing should increasingly be assessed through two lenses: **price per token** and **price per unit of intelligence, or ultimately, cost per completed task**.

Historically, most industry discussions focused on token pricing. However, as models become more capable, token price alone becomes a less relevant measure of value. What matters increasingly is how effectively a model converts tokens into useful work.

Indeed, we see growing discussions arguing that **cost per completed task** is a more meaningful metric than cost per token, as it incorporates not only model intelligence but also inference efficiency, cache engineering, infrastructure utilisation, latency, and workflow execution capability.

On one hand, we believe that **frontier models are likely to preserve their pricing premium for the foreseeable future**. The reason is straightforward: the intelligence frontier continues to advance rapidly. Scaling laws, reinforcement learning, synthetic data generation, reasoning improvements, and agentic architecture etc., are collectively pushing model capability boundaries higher every quarter. As models become capable of handling increasingly complex coding, research, and enterprise workflows, total token consumption per task also rises. The resulting value creation enables frontier providers to charge premium prices in high-value scenarios such as software engineering, enterprise automation, specialised knowledge work and more in exploration.

Total parameters of the largest known frontier model (as of August 2026)



Source: Company, Financial Times, Epoch AI, Daiwa estimates


LLM market is big enough for multiple winners

A second reason the premium should endure is that **fewer companies can afford to compete for it**. Frontier training costs are rising exponentially, while platform players continue to push up the scale curve, with ByteDance reportedly training at up to 10tn parameters, Alibaba's Qwen 3.8 Max reaching 2.4tn parameters, and Z.AI's GLM-5.2 at 744bn versus GLM 5.5's reported 1tn. As a result, leadership may rotate quickly, but increasingly within a small group of players that can continuously fund larger models and compute infrastructure.

Model leadership window is likely short

However, **the premium window is likely short**. Track records suggest that frontier leadership shifts quickly and is increasingly difficult and expensive to achieve and maintain. A small number of long-term winners may ultimately emerge, but visibility on the timeline remains limited.

Frontier language model intelligence overtime



Source: Artificial Analysis (as of 31 July 2026)

Cost efficiency will be the key to adoption and scalability

At the same time, we expect **cost efficiency to become an increasingly important competitive factor, which benefits Chinese AI models with lower cost and competitive pricing**. As model performance becomes sufficient for a broader range of mainstream tasks, incremental intelligence gains become less valuable relative to lowering the cost and barriers to execution. Once mid-tier models can reliably complete common agentic workflows, enterprises should increasingly route workloads according to performance-to-price, making pricing competition difficult to avoid.

API-based usage charge and subscription based monthly charge are two mainstream charging models for LLMs today. Currently, Chinese leading models are mostly open source, easy to deploy, and offer a substantial API price advantage over leading US frontier models. Based on our blended price methodology, as of early August, **the median price of US models (USD6.6/m token) under global top 50 scored is c.8x higher than that of China models (USD0.8/m token)**. For flagship models, the best-scored Chinese model Kimi K3 (ranking No.5 worldwide) costs USD6.6/m tokens, which is around **70%** below Claude Fable 5 (ranking No.1), while their intelligence score difference is less than 4%. The discount is wider for other flagship models: GPT 5.4 Pro is charging **c.390 times** more than DeepSeek V4 flash 0731.

Global frontier LLM API pricing comparison (as of August 2026)

Rank*	Model	Creator	Overall Score	Input price [^]	Output price [^]	Blended price [^]	Coding	Agentic	Multimodality
50	DeepSeek V4 Flash	DeepSeek	60	0.1	0.3	0.2	42	43	-
19	Hy3	Tencent	68	0.1	0.5	0.3	63	61	-
25	GPT-5.6 Luna	OpenAI	67	0.2	1.2	0.5	72	58	67
28	GPT-5.4 nano	OpenAI	66	0.2	1.3	0.5	43	41	17
17	MiMo-V2.5-Pro	Xiaomi	69	0.4	0.9	0.6	55	42	-
18	MiniMax M3	MiniMax	69	0.3	1.2	0.6	48	40	48
40	MiniMax M2.7	MiniMax	63	0.3	1.2	0.6	46	35	-
49	DeepSeek V4 Pro	DeepSeek	60	0.4	0.9	0.6	42	43	-
27	Qwen3.7 Plus	Alibaba	66	0.3	1.3	0.6	58	43	73
46	GLM-4.7	Z.AI	60	0.4	1.8	0.8	51	51	-
34	Qwen3.6 Plus	Alibaba	64	0.3	2.0	0.8	53	37	68
32	Gemini 3.5 Flash-Lite	Google	65	0.3	2.5	1.0	52	40	64
36	Grok 4.3	xAI	64	1.3	2.5	1.6	38	6	65
31	GLM-5	Z.AI	65	1.0	3.2	1.7	58	54	-
29	GLM-5-Turbo	Z.AI	66	1.2	4.0	2.0	-	-	-
44	GLM-5V-Turbo	Z.AI	62	1.2	4.0	2.0	57	-	-
7	Muse Spark 1.1	Meta	76	1.3	4.3	2.2	64	68	80
16	Muse Spark	Meta	70	1.3	4.3	2.2	63	58	77
22	GLM-5.1	Z.AI	67	1.4	4.4	2.3	56	49	-
41	GLM-5.2	Z.AI	63	1.4	4.4	2.3	64	57	-
13	Qwen3.7 Max	Alibaba	72	1.5	4.4	2.4	54	40	-
24	InKling	Thinking Machines Lab	67	1.9	4.7	2.7	43	46	53
8	Grok 4.5	xAI	76	2.0	6.0	3.2	57	61	69
9	Gemini 3.6 Flash	Google	75	1.5	7.5	3.3	64	46	75
37	Gemini 3.5 Flash	Google	64	1.5	9.0	3.8	61	45	85
33	Claude Sonnet 5	Anthropic	65	2.0	10.0	4.4	68	59	80
12	GPT-5.6 Terra	OpenAI	72	2.0	12.0	5.0	65	73	76
21	Gemini 3 Pro	Google	67	2.0	12.0	5.0	61	58	74
30	GPT-5.3 Codex	OpenAI	66	1.8	14.0	5.4	63	61	-
10	GPT-5.4	OpenAI	73	2.5	15.0	6.3	58	56	69
5	Kimi K3	Moonshot AI	80	3.0	15.0	6.6	77	66	90
35	Claude Sonnet 4.6	Anthropic	64	3.0	15.0	6.6	58	51	58
2	Claude Opus 5	Anthropic	83	5.0	25.0	11.0	78	72	89
6	Claude Opus 4.8	Anthropic	77	5.0	25.0	11.0	70	63	90
14	Claude Opus 4.7 (Adaptive)	Anthropic	72	5.0	25.0	11.0	64	64	51
15	Claude Opus 4.7	Anthropic	71	5.0	25.0	11.0	68	55	-
20	Claude Opus 4.6	Anthropic	68	5.0	25.0	11.0	61	56	62
38	Claude Opus 4.6 (Adaptive)	Anthropic	63	5.0	25.0	11.0	60	56	-
39	Claude Opus 4.5	Anthropic	63	5.0	25.0	11.0	61	36	29
4	GPT-5.6 Sol	OpenAI	81	5.0	30.0	12.5	78	76	85
11	GPT-5.5	OpenAI	72	5.0	30.0	12.5	71	64	69
1	Claude Mythos 5	Anthropic	83	10.0	50.0	22.0	80	66	86
3	Claude Fable 5	Anthropic	83	10.0	50.0	22.0	80	65	64
26	GPT-5.2 Pro	OpenAI	66	25.0	150.0	62.5	-	-	-
42	GPT-5.5 Pro	OpenAI	63	30.0	180.0	75.0	-	60	-
48	GPT-5.4 Pro	OpenAI	60	30.0	180.0	75.0	-	57	89

Source: BenchLM, companies, Daiwa estimates

Note: Our Blended price = input × 0.7 + output × 0.3. The 70/30 ratio reflects a generalised production workload.

*Ranking of intelligence score per BenchLM, we selected Top 49, and added DeepSeek V4 Flash 0731 as No.50.

[^]USD/m token

This gives enterprises who want to embrace AI at economic cost a broad range of Chinese alternatives, from premium domestic models such as Kimi K3 to high-volume, cost-efficient offerings from DeepSeek, MiniMax, Alibaba, Tencent, and Xiaomi.

A tiered domestic market with a c.33 times price differential


API Pricing also differentiates Chinese providers from one another. We estimate that the **GLM 5.2 API pricing** is almost **one-third of Kimi K3** and approximately **3x that of MiniMax M3** as of early August. MiniMax, Qwen, Tencent, and DeepSeek compete more directly on value for money, with blended prices generally below USD1 per 1m tokens. This creates a tiered domestic market in which stronger capability can still earn a premium, while lower-cost models compete for token volume and broader adoption. The gap between the highest priced Kimi K3 and the lowest priced DeepSeek V4 Flash 0731 is **c.33 times**.

The same pattern is visible in subscriptions, although usage allowances are not directly comparable across providers. Entry-level Chinese plans range from around USD3 per month for Kimi Code and USD9 for Alibaba's Qoder-CN to USD17 for Z.AI's GLM Coding Plan Lite, compared with USD20 for the standard individual plans offered by Anthropic and OpenAI. At the higher end, Z.AI's USD79 Pro and USD159 Max plans remain below the

USD100 and USD200 high-usage tiers offered by both Anthropic and OpenAI. Lower subscription prices reduce adoption barriers for individual developers and smaller enterprises, while API pricing remains the more important determinant of economics at scale.

Recent US price reductions suggest that Chinese providers' cost advantage is beginning to influence global pricing behaviour. However, lower pricing alone does not guarantee commercial success – in our view, efficiency primarily determines user adoption willingness: the sustainable winners will be those that combine competitive intelligence with lower serving costs and better user experience, sufficient inference capacity and product-level distribution and maintenance, and especially those who can build up user switching cost.

Frontier language model by country



Source: Artificial Analysis (as of 31 July 2026)

In our view, long-term AI competition will be determined not only by who builds the smartest model, but also by who can deliver intelligence most efficiently at scale.

Global frontier LLM subscription/coding plan pricing comparison – full tier detail (USD per user per month, as of August 2026)

Provider	HQ	Coding product	Segment	Tier	Currency	List price before promo (USD/mo)	List price before promo (USD/mo)	Seat minimum	Coding agent included?	Published usage allowance	Notes
Z.AI (Z.ai)	China	GLM Coding Plan	Individual	Lite	CNY	118	17	1	Yes	c.80 prompts / 5h; c.400 / week; 100 MCP calls / mo	30% intro promo takes monthly to USD12.60; annual USD151.20
Z.AI (Z.ai)	China	GLM Coding Plan	Individual	Pro	CNY	538	79	1	Yes	c.400 prompts / 5h; c.2,000 / week; 1,000 MCP calls / mo	c.5x Lite allowance; annual USD604.80
Z.AI (Z.ai)	China	GLM Coding Plan	Individual	Max	CNY	1,078	159	1	Yes	c.1,600 prompts / 5h; c.8,000 / week; 4,000 MCP calls / mo	c.20x Lite allowance; annual USD1,344
Z.AI (Z.ai)	China	GLM Coding Team Plan	Team / Enterprise	Standard Seat	CNY	598	88	2	Yes	15,000 credits / 5h; 66,000 / week	n/a – seat price not published in docs; shown at Z.ai checkout only
Z.AI (Z.ai)	China	GLM Coding Team Plan	Team / Enterprise	Premium Seat	CNY	1,198	176	2	Yes	35,000 credits / 5h; 155,000 / week	n/a – price not published. Adds early model access + peak-hour priority
Anthropic	USA	Claude Code (in Claude plans)	Individual	Pro	USD	20	20	1	Yes	5-hour rolling session + weekly cap	Official page lists USD17/mo on annual (USD200 billed up front)
Anthropic	USA	Claude Code (in Claude plans)	Individual	Max 5x	USD	100	100	1	Yes	5x Pro usage per 5-hour session	Monthly billing only
Anthropic	USA	Claude Code (in Claude plans)	Individual	Max 20x	USD	200	200	1	Yes	20x Pro usage per 5-hour session	Monthly billing only
Anthropic	USA	Claude Code (in Claude plans)	Team / Enterprise	Team – Standard seat	USD	25	25	2	Yes	More usage than Pro	Team plan covers 2–150 people; USD20/seat on annual
Anthropic	USA	Claude Code (in Claude plans)	Team / Enterprise	Team – Premium seat	USD	125	125	2	Yes	5x Standard seat usage	USD100/seat on annual; seat types can be mixed
Anthropic	USA	Claude Code (in Claude plans)	Team / Enterprise	Enterprise	USD	20	20	n/a	Yes	500k context on default model; usage billed on top	USD20/seat/mo PLUS usage at API rates. Annual billing only
OpenAI	USA	Codex (in ChatGPT plans)	Individual	Go	USD	8	8	1	Limited	Lightweight coding tasks only	Cheapest paid entry; Codex access is capped
OpenAI	USA	Codex (in ChatGPT plans)	Individual	Plus	USD	20	20	1	Yes	15–80 GPT-5.5 local msgs / 5h; 10–60 cloud tasks	Practical individual entry point for Codex
OpenAI	USA	Codex (in ChatGPT plans)	Individual	Pro 5x	USD	100	100	1	Yes	80–400 GPT-5.5 local msgs / 5h	'From USD100/mo'; promo has run at 10x
OpenAI	USA	Codex (in ChatGPT plans)	Individual	Pro 20x	USD	200	200	1	Yes	300–1,600 GPT-5.5 local msgs / 5h	20x Plus limits
OpenAI	USA	Codex (in ChatGPT plans)	Team / Enterprise	Business seat	USD	25	25	2	Yes	Plus-level per-seat limits + purchasable credits	USD20/user/mo on annual. Replaced the old Team plan (Apr 2026)
Kimi (Moonshot AI)	China	Kimi Code (in Kimi membership)	Individual	Moderato	CNY	19	3	1	Yes	60 agent credits; entry-level Kimi Code credits	Annual billing c.20% off
Kimi (Moonshot AI)	China	Kimi Code (in Kimi membership)	Individual	Allegretto	CNY	39	6	1	Yes	Higher credit + quota allocation; 1M context	Annual billing c.20% off
Kimi (Moonshot AI)	China	Kimi Code (in Kimi membership)	Individual	Allegro	CNY	99	15	1	Yes	Agent Swarm; Kimi Code at 5x	Annual billing c.20% off
Kimi (Moonshot AI)	China	Kimi Code (in Kimi membership)	Individual	Vivace	CNY	199	29	1	Yes	720 agent credits; Kimi Code at 30x	USD159/mo confirmed on annual billing
Kimi (Moonshot AI)	China	Kimi Business	Team / Enterprise	Business seat	CNY	50	7	5	Yes	Allegretto-level benefits per seat	USD599 per seat per year, annual billing only, 5-seat minimum
Qoder-CN (Alibaba)	China	Native coding tool, Qwen model only	Individual	Pro	CNY	59	9	1	Yes	2000 credits/month	Only Qwen models, support VSCode/JetBrains/Qoder-CN
Qoder-CN (Alibaba)	China	Native coding tool, Qwen model only	Individual	Pro+	CNY	169	25	1	Yes	6000 credits/month	Only Qwen models, support VSCode/JetBrains/Qoder-CN
Qoder-CN (Alibaba)	China	Native coding tool, Qwen model only	Team / Enterprise	Teams	CNY	99	15	1	Yes	3000 credits/seat/month, can be shared among team	Support IDE, CLI with seat management and usage monitoring
Qoder-CN (Alibaba)	China	Native coding tool, Qwen model only	Team / Enterprise	Enterprise VPC	CNY	199	29	50	Yes	3000 credits/seat/month, can be shared among team	Support IDE, CLI with seat management and usage monitoring
Bailian Coding Plan (Alibaba)	China	Coding tool supporting various models	Individual	Lite	CNY	40	6	1	Yes	suspended since April	Comprise Qwen, GLM, Kimi, Minimax top models, limited supply availability
Bailian Coding Plan (Alibaba)	China	Coding tool supporting various models	Individual	Pro	CNY	200	29	1	Yes	6k request/5h; 45k requests/week; 90k requests/month	Comprise Qwen, GLM, Kimi, Minimax top models, limited supply availability

Source: companies, compiled by Daiwa

Note: all figures are publicly listed prices, gathered 4 August 2026.

ARR trends: the AI market continues to outgrow its own forecasts

Before sizing the market, it is worth looking at what AI companies are actually earning today. Annualised recurring revenue (ARR) is the metric the industry reports on, the metric private rounds are priced against, and widely regarded as the most current read on demand available to us. It is also the number that has most consistently embarrassed published market forecasts.

ARR keeps overtaking published market forecasts

Most published forecasts are built from currently visible applications, users and business models, which understates the opportunity whenever capability and commercialisation advance faster than expected. That has been the consistent pattern. In MiniMax's late-2025 prospectus, it estimated a **USD22bn global LLM market in 2025 and USD41.4bn in 2026**. By July 2026, OpenAI and Anthropic alone reported ARR of **USD41.3bn and USD74.1bn, respectively** — a combined run-rate already more than 2.5x the full-year global market forecast, even before counting Google, xAI, Chinese developers and other application providers.

As AI evolves from answering questions to executing workflows, its addressable market could also expand beyond software budgets into labour, operating and transaction pools. If AI eventually approaches AGI-level capability, conventional sizing based on existing jobs, wages or software spending may become increasingly incomplete, and its economic value may be better framed as a share of GDP.

US: still compounding, but the slope weakened after June


US frontier labs have generated the vast majority of the total ARR generated to date. OpenAI and Anthropic ARR combined rose from **USD30bn at end-2025 to USD115bn in July 2026**, about three times growth in seven months, and roughly 19 times the entire visible Chinese LLM model MaaS industry's run-rate.

US vs. China: visible ARR comparison

USD115bn	c.USD6bn	c.19x
OpenAI and Anthropic combined ARR Jul-26	Visible China LLM model ARR Jul-26	US visible China scale

Source: Daiwa

Open AI vs. Anthropic: monthly ARR and growth trend of the two US frontier labs (USDbn)



Source: The Information, Reuters, CNBC, TickerTrends, company disclosures; Daiwa estimates

Note: ARR is an annualised run-rate, not recognised revenue.

Anthropic reports MaaS-reseller revenue gross of partner commission; OpenAI reports net. On a net-comparable basis Anthropic's July figure is estimated nearer USD50-55bn.

Sequential growth accelerated before April 2026. Combined monthly ARR growth ran at 14-15% a month through 1Q26 and then stepped up to 31-33% a month between April and June, as coding and agentic workloads scaled. July broke the pattern. Combined growth

US visible Jul-26 ARR scale is 19x that of China's

slowed to 13.0% MoM, less than half the preceding three-month average. OpenAI added 3.2% vs. 21.2% in June, compared to Anthropic's 19.4% vs. 40.9%.

We do not overread a single month's softening momentum as an indicator of industry growth slowdown. ARR is a run-rate, not recognised revenue, and is reported on inconsistent bases across companies, with one print not making a trend. But **the direction is corroborated by pricing behaviour, as enterprises tighten AI budgets and route suitable workloads towards cheaper models**, Chinese providers are gaining usage through stronger performance-to-price. Chinese models accounted for c.46% of US enterprise token usage on OpenRouter by July 2026, while OpenAI's subsequent price cuts suggest rising competitive pressure. This does not imply weaker demand, as token volume still increased more than ten-fold during the first 30 weeks of 2026. Rather, revenue per token is moderating as lower prices unlock greater consumption. And on balance we think the near-vertical phase of US ARR growth has probably passed before the next high-value scenario is observed.


One caveat on comparability, which matters if the two are read side by side: according to *Forbes* and *Financial Times*, **Anthropic books cloud-reseller revenue gross of the commission retained by AWS and Google Cloud, while OpenAI books net of Microsoft's share**. According to *The Verge*, OpenAI argued in April 2026 that this overstated Anthropic's then-USD30bn run-rate by c.USD8bn. On a like-for-like net basis, we estimate that Anthropic's July figure is probably nearer USD50-55bn. The growth comparison is unaffected, but the level gap is narrower than the headline suggests.

China: faster growth from a smaller base

Steep curve of picking up Chinese LLM ARR growth in 2026

While US frontier labs continue to lead in absolute revenue scale, Chinese players are exhibiting similarly steep growth trajectories from a smaller base. **MiniMax achieved a 32% monthly ARR CAGR between December 2025 and May 2026, while Z.AI delivered a 41% monthly ARR CAGR between March and July 2026**. We see this as evidence that commercial AI adoption is expanding rapidly across the ecosystem, with long-term revenue leadership increasingly determined by a combination of model capability, pricing competitiveness, product execution, and distribution, rather than technology leadership alone.

Chinese frontier LLM providers: latest ARR (USD bn)



Source: Company disclosures, *The Information*, *TechCrunch*, *kr-asia*, *Digitimes*, *SCMP*; Daiwa estimates


Note: Platform-company figures cover the group large-model business and are not directly comparable with the independent labs' API-only ARR;

#DeepSeek does not disclose revenue and should be treated as indicative order of magnitude only;

*Daiwa estimates, given *STAR Market Daily* reported a USD300m ARR in June 2026 and 6 times daily usage growth post K3 release;

^per *Digitimes*, ByteDance's AI revenue hit USD4b run-rate on 31 July; we excluded an estimated ARR of USD1.7bn of Seendance 2.0 whose revenue is at CNY1bn/month here for finding LLM MaaS revenue.

Z.AI vs. MiniMax: historical ARR trends



Source: Company

Overall, **AI remains in a hyper-growth phase without clear consolidation.** US frontier labs lead in absolute revenue and premium workloads, while Chinese providers are gaining scale through cost efficiency, faster model iteration and broader distribution. We therefore expect multiple business models to coexist, with capability winning value and efficiency winning volume. This should allow emerging providers such as Z.AI and MiniMax to gain commercial traction alongside global leaders, while supporting TAM expansion beyond current market forecasts.

Volume trends: token volumes surge, but leadership remains fluid

Following the pricing comparison above, usage data provides a second lens through which to assess the role of model economics. According to OpenRouter data, aggregate token consumption increased by more than **10x** between the first week of 2026 and week 30 (ending 26 July).

In our view, this provides further evidence that 1) AI application TAM is expanding faster than previously expected, with wider usage, penetration and deployment of agentic AI; and 2) declining intelligence costs enables and attracts more usage, which forms a “technology-commercialisation flywheel”. This means that as models become increasingly capable, each dollar spent generates greater economic value, driving higher usage volumes.

Usage leadership changes within weeks


Usage leadership shifts as fast as model leadership

However, beneath the rapid usage expansion, competition among model providers remain intense, with durable leadership proving difficult to sustain. **Similar with model score leadership, token usage leadership is shifting even faster, meaningfully disrupted by price discounts when new models launch.** For example, HY3 Preview rose to the top of the usage rankings shortly after launch but retained the leading position for only 3 weeks before being displaced by DeepSeek V4 Flash, which itself remained No.1 for 7 weeks before being overtaken by newer entrants.

We believe this demonstrates that foundation models have become highly contestable, with **limited user lock-in** and workloads migrating rapidly toward models offering superior cost-performance economics.

The speed of workload migration also highlights low switching costs associated with foundation models. New model launches frequently trigger immediate shifts in token consumption as developers redirect workloads toward models offering better capability-cost trade-offs. However, these gains often prove temporary as newer competing models enter the market. As a result, developer behaviour in the LLM market appears highly pragmatic and increasingly driven by real-time performance and economics.


Weekly usage of models across OpenRouter



Source: OpenRouter, Daiwa

Note: OpenRouter is a third-party model access provider who discloses token usage on its server, often used as an indicator of global token usage trend and model vendor market share, though it accounts for about mid to low single-digit % of actual global token usage per our channel checks.

China model market share by volume on OpenRouter



Source: OpenRouter, compiled by Daiwa

Efficiency is the clearest common factor among volume leaders

This naturally raises the key question: what ultimately drives leadership? In our view, the answer may lie in **cost-performance ratio**.

First, Chinese models dominate visible usage rankings. In June 2026, seven of OpenRouter's 10 most-used models were developed by Chinese companies, collectively accounting for 84% of token consumption among the Top 10. The dataset is weighed towards individual developers, who are typically more price-sensitive than large enterprises, but it still indicates that commercial adoption is determined not only by absolute benchmark performance, where US frontier labs retain an advantage, but also by the ability to deliver sufficient intelligence at an attractive price.

Second, only a small number of providers maintain a top usage position across multiple model cycles, most notably MiniMax and Xiaomi. MiniMax models ranked among OpenRouter's top 10 for 21 of the first 30 weeks of 2026, while Xiaomi models appeared for 18 weeks. Both providers emphasise cost efficiency, with blended API prices equivalent to only around 3% of Claude Fable 5's price as of July 2026. Their persistence therefore reinforces our view that model economics is an important driver of sustained usage.

Overall, we think if measured in tokens processed, the winning formula is **"good enough plus competitive pricing"**: as AI adoption expands across a broader economic base, volume leadership should accrue to whoever delivers the most intelligence per dollar. The same period in which Chinese models took seven of the top ten slots on OpenRouter was also the period in which US frontier labs added more absolute ARR than the entire visible Chinese model industry has generated in total, because the highest-value workloads still "pay up" in terms of capability.

Ecosystem related user experience might form durable advantages in long-run

We therefore do not read this as one side displacing the other. Both capability and efficiency are durable positions, and they serve different customers, cost-sensitive developers, and SMEs on one side, enterprises running high-stakes, high-token-intensity workloads on the other. **What is specifically favourable for Chinese players is that operating under cost discipline from day one is a genuine structural advantage on the efficiency route** and, as Z.AI's pricing demonstrates, it does not preclude competing on the capability route as well.

Durable advantages may form in ecosystem development

At the same time, as AI adoption expands beyond experimentation into real-world applications and agentic workflows, **distribution, user engagement, and ecosystem development are becoming increasingly important sources of competitive advantage**. In this environment, sustainable leadership depends not only on technological innovation, but also on a company's ability to translate model capabilities into scaled adoption with profitability and leveraged commercial deployment to continuously strengthen its technology.


Companies with direct commercial touchpoints, including AI agents and AI-native applications, are better positioned to speed up the flywheel. **Real-world usage generates valuable feedback, behavioural data and revenue streams that accelerate model iteration and product optimisation**. Stronger model capabilities then drive broader adoption, deeper engagement and higher monetisation, reinforcing a virtuous cycle between technological progress and commercial scale. We believe companies that successfully establish this feedback loop will be better positioned to sustain technology leadership and commercialisation over the long term.

Pricing trends: frontier model premium vs. falling price of intelligence

LLM pricing is both increasing and decreasing

We believe the industry has exhibited two seemingly contradictory trends in prices since 2022. On one hand, leading frontier models have generally maintained strong pricing power, with pricing moving higher alongside capability improvements since 2024. While the sector experienced pricing volatility from 2Q26, leading vendors continue to command premium pricing when capability leadership is demonstrated. For example, Z.AI's GLM series has consistently raised API prices alongside model upgrades, with the blended token price of GLM-5.2 (launched in June 2026) being roughly thrice that of GLM-4.7 (launched in December 2025).

Token price index: average blended cost of SOTA model




Source: AI Token Price Index, Daiwa collects

Intelligence per unit is getting cheap

On the other hand, the price of intelligence has fallen dramatically. According to Artificial Analysis, the cost per intelligence point has declined sharply across capability tiers since 2022. The cheapest model, achieving an intelligent score above 40 became approximately **88% cheaper**, falling from roughly USD0.09 to USD0.01 per intelligence point from late 2025 to mid-2026 per LLM Stats.

Frontier language model price by intelligence index band over time



Source: Artificial Analysis (as of 31 July 2026).

The divergence highlights an important distinction: **token prices and intelligence prices are moving in opposite directions**. Premium frontier models remain elevated on a per-token basis, while the amount of intelligence delivered per dollar continues to improve rapidly.

Recent pricing actions from OpenAI and Anthropic suggest that industry price competition is focused on commoditised and mid-tier models rather than frontier intelligence. OpenAI cut GPT-5.6 entry-tier pricing by 80% in July 2026 while keeping flagship pricing unchanged, and Anthropic made Claude Sonnet 5's discounted USD2/10 pricing


permanent in August 2026, highlighting continued cost deflation in mainstream AI services but relative pricing resilience at the high end

In the first section, we analyse in detail the price tier landscape and price gap between different players and how the above trend reflects and impacts user behaviour and market development. Overall, our conclusion is that **the market is big enough for multiple winners**. Capability chasers will maintain their premium pricing for more reliable performance on high-value, complex tasks, while efficiency presenters may compete on performance for price ratio while enlarging the addressable user base and market value.

Deep dive on capability vs. efficiency: discussion on scaling impact and activation ratios

We would like to highlight the scaling law impact on the LLM landscape. It is obvious that **progressively fewer companies can afford to compete on the premium side**. Scaling still works for model training, and the industry is spending as though it does. The largest disclosed model has gone from 175bn parameters in 2020 to Kimi K3's 2.8tn in July 2026, and ByteDance is reported to be pre-training at up to 10tn. As a result, frontier training costs have compounded at c.5x a year since 2020. A leading run today costs a few hundred million dollars. On trend it passes a billion in 2027, not to mention the compute capacity to support its hosting and massive usage.

Cost of a leading training run



Source: Epoch AI, Daiwa collates and estimates:

This arithmetic does the sorting. Frontier capability is no longer won by a research idea that a well-funded startup can execute; it is won by a balance sheet that can absorb a 9-figure training run and then do it again 18 months (or even shorter) later. Every doubling removes participants. **We expect the premium tier to narrow to a handful of names with sustainable cash flow**, which is precisely why we think premium might still exist to support the sustainability.

This is also why we do not read a short premium window as a weak one. Leadership changes hands quickly and will continue to do so. However, it now rotates inside a shrinking club, and the club's pricing power comes from the cost of entry rather than from holding rank in any given quarter. The question for the next five years is not who leads today but who can survive in 3, 5 or 10 years.

By this measure, platforms are becoming increasingly aggressive on the scale race. ByteDance is reported to be training at up to 10tn parameters against GLM-5.2 at 744bn, and Alibaba's recent Qwen 3.8 Max at 2.4tn ranks No.2 among released open-source models, larger than GLM 5.5's reported 1tn to be released in August 2026. The independents do not need to win the scale race. However, they do need a capability

Cost of training and hosting is building up moat and pricing power for the premium

position that does not depend on winning it, which is why we value them on workload leadership rather than on rank.

Why lower activation ratios do not guarantee cheaper tokens

Another popular explanation for China's low AI pricing is that Chinese models are more sparse. We believe this view is incomplete. While **activation ratios** affect inference efficiency, they **do not directly determine token pricing and should always be interpreted alongside the absolute number of activated parameters**, which is a much better proxy for serving cost.

What is the ratio? Every flagship model is now a mixture of experts (MoE): it holds a very large number of weights but fires only a small slice of them on any given token. The activation ratio is this slice: activated parameters over total parameters. It matters because it breaks the link between how big a model is and what it costs to run.

A 2.4tn-parameter model that activates 4% of its weights roughly does the arithmetic of a 95bn-parameter dense one. **That is the whole reason frontier models kept growing while inference prices kept falling.** However, the landscape has converged: **every flagship worth naming now activates 3-5% of its weights, against roughly 10% two years ago.** Moreover, the gap between the sparsest and the least sparse is smaller than the gap between two versions of the same model. Sparsity is the entry ticket to serving a trillion-parameter model for all. It is no longer an edge over anyone who also holds one.

A lower ratio does not mean a cheaper token. This is where the popular explanation for cheap Chinese tokens goes wrong. If sparsity sets price, the sparsest models would be the cheapest; however, they are the most expensive. Kimi K3 fires 3.7% of its weights and is the most expensive model in our set, whereas Tencent's Hunyuan Hy3 is the least sparse flagship at 7.1% and the second cheapest to buy. Across the six priced flagships, the ranking is close to inverted. This is because the ratio is mostly a story about its own denominator: a low activation ratio is what a lab buys to afford a 2.4tn model, not a discount it hands to the customer. What it charges afterwards is a margin decision.

Activation ratio

Absolute activated parameters reflect serving cost more than activation ratio


Model	Provider	Total params (bn)	Activated params (bn)	Activation ratio (%)	Blended price (USD/m)	Rank: ratio (sparsest =1)	Rank: activated (lowest =1)	Rank: price (cheapest =1)
Kimi K3	Moonshot	2,800	104	3.7%	6	1	6	6
Qwen3.8-Max	Alibaba	2,400	95	4.0%	3	2	5	5
DeepSeek-V4-Flash-0731	DeepSeek	284	13	4.6%	0.18	3	1	1
MiniMax-M3	MiniMax	428	22.5	5.3%	1.05	4	3	3
GLM-5.2	Z.ai / Z.AI	744	40	5.4%	2.15	5	4	4
Hunyuan Hy3	Tencent	295	21	7.1%	0.26	6	2	2

Source: vendor model cards and pricing pages, Daiwa collects and estimates

Read it with the absolute number. Serving cost per token scales with activated parameters, not with the ratio. The total count is a fixed memory cost, while the activated slice is the arithmetic actually paid for on every token generated. Ranked this way, the picture is orderly rather than perverse: all six priced flagships take the exact same order on activated parameters as on price, from DeepSeek-V4-Flash at 13bn to Kimi K3 at 104bn. Thus, the ratio indicates how ambitious a lab was allowed to be with its total parameter count, while the absolute activated number reveals what the token costs to produce. Price follows the second, then moves by however much margin or subsidy the lab chooses to add. This residual choice is still wide, spanning more than four times per billion of activated parameter across the same six models. Architecture is common property; the pricing is not.

Activation ratio vs. absolute activated parameters

Sparsity is not what makes a token cheap – the amount of compute you fire is



Priced flagship models only. Blended price = 0.75 × input + 0.25 × output. GLM-5.2 (Z.ai) in red.

Source: vendor model cards and pricing pages; aggregator pricing for GLM-5.2, MiniMax M3 and Hunyuan Hy3; Daiwa collates. Data as of 9 August 2026.

Source: companies, Daiwa collects and estimates

We will discuss how big the market could be in the following section.

China AI TAM analysis: scenario-based framework

We estimate China's AI application market size through two complementary approaches: **bottom-up TAM on coding and consumer applications**, and **top-down addressable value analysis on enterprise and multi-modality applications**.

The methodology partly follows the economic-value approach used in SpaceX's IPO materials, given that xAI and Chinese foundation-model providers share broadly similar product architectures across consumer applications, model APIs and enterprise services. Despite varied focus markets, the underlying addressable-market logic is comparable: AI monetisation can be framed against the consumer, labour and enterprise economic pools that models may serve.

We estimate a CNY14tn theoretically addressable market value for China AI model vendors

We use **2025 market revenue as the starting point**, **2028 as an interim adoption scenario**, and **2030 as the long-term TAM scenario**. Because the two approaches partially overlap, the results should not be added mechanically.

Enterprise AI applications addressable economic-value analysis

Scenarios	2025	2028E	2030E
AI Coding (TAM, CNYbn)	0.4	132	264
Consumer AI applications (TAM, CNYbn)	1.1	58	116
Enterprise applications (Addressable economic pool, CNY tn)	10		
Multi-modality applications (Addressable economic pool, CNY tn)	4		

Source: Daiwa estimates

Our conclusion suggests a **CNY264bn** China AI coding TAM and **CNY116bn** consumer AI TAM by 2030E as the base case visible market and a **broader economic value (EV)** that AI may ultimately reshape at **CNY10tn** for enterprise applications and **CNY4tn** for multimodality application. A more detailed analysis follows.

1. AI Coding: monetising developer productivity

AI coding evolves from code generation into a broader software-production platform, covering requirement analysis, development, testing, debugging, deployment and maintenance. Coding agents can reduce the amount of human effort required for software development, while creating incremental value by improving developer efficiency, shortening delivery cycles, raising code quality and enabling broader software output.

Accordingly, we size the AI coding TAM using a combined value capturing framework that captures both:

- the addressable personnel-cost pool represented by software and IT services employees in developer-related roles; and
- the incremental value created through efficiency and productivity enhancement.

Under this approach, we estimate the addressable pool as:

Software & IT services sector employment × annual personnel cost × addressable developer-role ratio
+ Software & IT services sector employment × annual personnel cost × value added via efficiency enhancement

This implies:

$$11\text{m} \times \text{CNY}100,000 \times 50\% + 11\text{m} \times \text{CNY}100,000 \times 30\% = \text{CNY}880\text{bn}$$

Key assumptions applied in China coding TAM analysis

Assumption	Estimate	Basis
Software & IT services sector employment (mn)	11	Industry-employment proxy, MIIT, 2025 number
Addressable developer roles	50%	Daiwa assumption
Value added via efficiency enhancement	30%	Industry survey suggests 30-50% efficiency enhancements among developers
Addressable employees	5.5	11m*50%
Annual personnel cost (CNY)	100,000	Daiwa blended assumption on NBS data
Total addressable pool	880	Addressable personnel-cost pool plus value-added market size
AI value capture in 2028	10% / 15% / 20%	Daiwa scenarios
AI value capture in 2030	20% / 30% / 40%	Daiwa scenarios

Source: Ministry of Industry and Information Technology (MIIT), National Bureau of Statistics, Daiwa collects and estimates

The CNY100,000 cost assumption is deliberately below official industry wage data. China's National Bureau of Statistics reported 2025 average annual wages of CNY248,752 for non-private employees and CNY128,166 for private employees in information transmission, software and IT services. We use a lower blended figure to account for role mix, geographic differences, junior employees and the inclusion of lower-cost software-service personnel.

China coding TAM scenarios

CNYbn	2025 market	2028E	2025-28E CAGR	2030E	2025-30E CAGR
Bearish	0.4	88	504%	176	238%
Base	0.4	132	591%	264	266%
Bullish	0.4	176	661%	352	288%

Source: Ministry of Industry and Information Technology (MIIT), National Bureau of Statistics, Daiwa collects and estimates

IDC estimates China's AI coding market at **CNY399m in 2025**, rising to **CNY1.17bn in 2026**. We use CNY399m as the 2025 observable revenue base.

Our 2030 AI coding TAM range increases to **CNY176-352bn**, based on AI capturing **20-40%** of the CNY880bn addressable value pool. The implied CAGR remains exceptionally high because the 2025 market is still at an early commercialisation stage.

A 20-40% value-capture assumption does not imply equivalent workforce replacement. It represents the share of productivity value potentially monetised through subscriptions, API usage, agent execution and enterprise software. Developers and enterprises are likely to retain part of the benefit, while AI providers capture the remainder through pricing and increased consumption.

2. Consumer AI revenues: penetration and ARPU expansion

For consumer AI, we use a population-based approach:

$$\text{Addressable population} \times \text{paid penetration} \times \text{monthly ARPU} \times 12$$

The National Bureau of Statistics reported **968.5m people aged 15-64 in 2025**, which we use as a conservative proxy for the addressable population for consumer AI services. We believe this cohort represents the primary user base for intensive AI-native applications, including search, content creation, education, entertainment, shopping assistance and personal productivity tools.

Consumer AI can be monetized through various sources in China

Subscriptions are currently the most visible monetisation model. That said, we expect consumer AI companies to monetise users through a combination of channels, including advertising, transaction commissions, agent-service fees, revenue sharing, premium features and other emerging business models. Therefore, we frame the opportunity using **ARPU per user across monetisation sources**, rather than subscription fees alone.

We believe this approach is particularly relevant in China. Chinese consumers have historically demonstrated lower willingness to pay for standalone digital subscriptions than some overseas markets. However, Chinese internet platforms have also proven highly effective at monetising large user bases through localised business models, including

advertising, commerce integration, transaction-based services and ecosystem monetisation. As AI becomes increasingly embedded within consumer workflows, we believe user lifetime value (LTV) can be unlocked through a broad range of monetisation mechanisms beyond direct subscription payments.

Key assumptions applied in consumer AI TAM analysis

Scenario	Addressable population (m)	2028E paid penetration	2030E paid penetration	Monthly ARPU (CNY)
Bearish	968.5	5%	10%	20
Base	968.5	10%	20%	50
Bullish	968.5	15%	30%	100

Source: National Bureau of Statistics, Daiwa collects and estimates

The ARPU assumptions represent the average monthly revenue generated per monetised user across revenue streams rather than a specific subscription price. The penetration assumptions reflect the proportion of users that ultimately generate meaningful economic value through AI services such as direct payment, advertising-supported engagement, transaction activity or other monetisation mechanisms.

Consumer AI TAM scenario analysis

CNY bn	2025 market size	2028E	2025-28E CAGR	2030E	2025-30E CAGR
Bearish	1.1	12	119%	23	84%
Base	1.1	58	275%	116	154%
Bullish	1.1	174	441%	349	216%

Source: National Bureau of Statistics, Daiwa collects and estimates

We use the **CNY1.1bn 2025 consumer LLM market forecast** as the starting point for our analysis. Under our framework, China's consumer AI market could reach **CNY23-349bn by 2030**, depending on user adoption, engagement intensity and monetisation effectiveness.

Our base-case scenario assumes **20% monetised-user penetration** and **CNY50 monthly ARPU**, implying a **CNY116bn** market opportunity by 2030. The bullish scenario assumes AI evolves into a mainstream consumer platform integrated across productivity, content, search, commerce and personal-agent use cases, resulting in **30% monetised-user penetration** and **CNY100 monthly ARPU**, corresponding to a **CNY349bn** TAM.

Importantly, we view these estimates as a measure of the long-term consumer value pool rather than a forecast of subscription spending alone. As AI capabilities continue to expand, new monetisation models are likely to emerge. Consequently, future consumer AI revenues may be increasingly driven by a combination of subscriptions, advertising, commerce enablement, transaction commissions and AI-agent services, rather than any single revenue model.

3. Enterprise applications: anchoring AI to the digital economy

For enterprise applications, we anchor AI TAM to the broader digital economy rather than the current size of the AI non-hardware market. Our view is that the long-term opportunity for foundation-model providers, agent platforms and AI-native applications should be measured by the economic activities that can be augmented, automated or executed by AI. Under this framework:

Enterprise AI application is visibly the key battlefield, in our view

Enterprise AI TAM = Digital-economy value × AI application value-capture ratio

According to the **National Development and Reform Commission (NDRC)**, China's digital economy was **CNY49tn in 2025**, representing **c.35% of GDP**. We use this as our starting point for assessing the enterprise AI opportunity.

A key debate in sizing AI TAM is determining what proportion of economic activity may ultimately become AI-addressable. While some market participants, such as SpaceX, assume AI could eventually capture c.100% of the digital economy, we believe a significant portion of digital spending will continue to reside in hardware, infrastructure and other

segments outside the addressable scope of foundation model companies such as Z.AI and MiniMax. Accordingly, we first exclude these non-addressable components and estimate that roughly **c.50% of the digital economy is potentially AI application addressable**.

Within this addressable economic pool, we estimate that AI applications can progressively automate, enhance or optimise a growing share of enterprise activities, including software development, customer service, content generation, knowledge management, workflow automation, data analysis and business-process execution. Notably, this framework incorporates the AI coding opportunity discussed earlier, as software development represents one component of the broader enterprise AI value pool rather than a standalone market. We therefore apply **20%/40%/60% penetration assumptions to this AI-addressable portion**, representing conservative, base-case and bullish scenarios for long-term AI adoption and monetisation.

Scenario analysis: enterprise AI applications addressable economic-value

CNY trn	2025 digital economy value	AI application value capture	AI application addressable value
Bearish	49	20%	5
Base	49	40%	10
Bullish	49	60%	15

Source: National Development and Reform Commission, Daiwa collects and estimates

AI applications can progressively automate, enhance or optimise a growing share of enterprise activities

Our **base-case scenario assumes 40% AI value capture**, implying an addressable economic opportunity of approximately **CNY10tn**. This assumption reflects our view that AI applications are likely to become deeply embedded across knowledge work, software development, business operations and decision-making processes over time.

Importantly, these figures represent the broader economic value that AI applications may help create, optimise or mediate, rather than the revenue opportunity available to any individual AI company. Similar to how internet platforms ultimately monetised a meaningful portion of the economic activity they facilitated, AI vendors are likely to capture a fraction of this addressable value.

We believe this framework reflects the monetisation opportunity for AI better than conventional market-size estimates that focus solely on current spending on AI software. As AI evolves from a productivity tool into an execution layer for enterprise workflows, the relevant addressable market may progressively expand beyond software budgets towards the broader digital economy itself.

4. Multi-modality: tapping into China's content economy

Multi-modality is increasingly augmenting content-centric creative and production workflows, including content generation, ad creatives, marketing materials, and creative software workflows (eg, graphic design, image/video editing). In essence, multi-modal AI serves as an enabling technology that expands the range of tasks AI applications and agents can perform, accelerating AI adoption rather than representing a standalone application category. As such, we believe the opportunity is best assessed through the lens of the broader content economy, where AI-generated content (AIGC) can create meaningful economic value.

We size the multi-modal AI opportunity based on China's cultural new economy revenue pool, which encompasses digital content creation, online media, internet entertainment, digital publishing, advertising and other content-centric activities. We believe these segments represent the most immediate commercialisation opportunities for AIGC technologies, given the central role of content in their value creation and monetisation models. Under this framework, we estimate multi-modality AI's addressable value pool as:

Multi-modal AI TAM = China cultural new economy revenue × AI value-capture rate

China multi-modality AI addressable economic-value analysis

CNY trn	2025 China cultural new economy value	Multi-modality AI value capture	Multi-modality AI addressable value
Bearish	7.7	30%	2.3
Base	7.7	50%	3.9
Bullish	7.7	70%	5.4

Source: National Bureau of Statistics, Daiwa collects and estimates

We use China's 2025 cultural new economy revenue pool of CNY7.7tn, reported by the NBS, as the starting point of our analysis. As model capabilities evolve, we expect AI adoption to expand across the content value chain, supporting richer, more realistic and immersive content experiences while enhancing production efficiency and scalability. We apply long-term AI penetration rates of 30%/50%/70% under our bear, base and bull case scenarios to assess the potential impact of multi-modal AI.

Our base-case scenario assumes 50% AI value capture, representing an addressable economic value pool of **CNY3.9tn**. While the pace of value realisation depends on adoption rates, use-case breadth and commercialisation progress, we believe multi-modal AI has substantial long-term commercial potential. In our view, the inherently cross-modal nature of human communication underpins our belief that multi-modal capabilities will become a foundational layer of AI applications embedded across most content-related workflows.

Summary on China AI application TAM

In summary, our AI application framework assesses the **visible monetisable market (TAM)** and the **broader economic value (EV)** that AI may ultimately reshape.

Enterprise AI applications addressable economic-value analysis

Scenarios	2025	2028E	2030E
AI Coding (TAM, CNYbn)	0.4	132	264
Consumer AI applications (TAM, CNYbn)	1.1	58	116
Enterprise applications (Addressable economic pool, CNY tn)	10		
Multi-modality applications (Addressable economic pool, CNY tn)	4		

Source: Daiwa estimates

Our AI application framework highlights a substantial gap between today's market size and the long-term opportunity.

Under our base case, we estimate China's **AI coding TAM** could expand from **CNY0.4bn in 2025 to CNY264bn by 2030**, implying a **266% CAGR**, as AI increasingly enhances software-development productivity and captures value from labour efficiency and incremental output.

Meanwhile, **consumer AI TAM** could grow from **CNY1.1bn to CNY116bn by 2030**, representing a **154% CAGR**, driven by rising user penetration and improving monetisation through subscriptions, advertising, commissions, agent services and other emerging business models.

We have made a cross-check: the visible China ARR of USD6bn as of July 2026 is c.11% of our estimated coding + consumer AI TAM in 2030 and 0.4% of our estimated enterprise application addressable market. While this 10% may appear elevated, we believe it primarily reflects the conservative nature of our TAM framework, which assumes AI captures only 30% of the underlying addressable spend by 2030. In addition, ARR is a forward-looking run-rate metric rather than recognised revenue, meaning it can scale ahead of realised monetisation. As a result, the comparison likely overstates market penetration and should not be interpreted as AI monetisation having already reached a mature stage.

Beyond these visible revenue opportunities, we estimate AI could ultimately influence approximately **CNY10tn of enterprise digital-economic activity** and **CNY4tn of the cultural new economy value creation** under our base-case assumptions. While only a fraction of this economic value is likely to be directly monetised by AI vendors, we believe

We see a substantial gap between today's market size and the long-term opportunity

the addressable economic pool provides a better representation of AI's long-term impact than conventional software-market sizing alone.

Note that the estimates cannot be added directly. Coding activity is part of enterprise applications, while some consumer subscriptions may also be purchased for professional use. The analysis is intended to show the potential order of magnitude under different monetisation frameworks, rather than produce a single point estimate.


Considering this, our analysis suggests that current AI revenues remain at a very early stage relative to the scale of economic activity that AI applications may eventually augment, automate and reshape.

Our estimated China AI TAM is c.9% of SpaceX's global AI TAM estimates

Linking the China framework to SpaceX's global AI TAM

SpaceX's IPO presentation estimates a **USD26.5tn global AI TAM**, comprising USD22.7tn of enterprise applications, USD2.4tn of AI infra, USD760bn of consumer subscriptions and USD600bn of digital advertising. Excluding AI infra, our estimated China TAM of CNY14tn is **c.9%** of its estimated global TAM.

Space X's estimated TAM by segment



Source: Space X

The presentation effectively treats AI as an intelligence layer addressing existing economic pools, rather than defining TAM as the current revenue of model companies. Enterprise applications account for most of the estimate, reflecting the possibility that AI participates in a broad range of knowledge work, decisions and business processes.

We apply a similar logic to China because xAI and Chinese LLM providers commercialise comparable categories of products, and both business models ultimately seek to monetise intelligence across consumer and enterprise activities.

The comparison also reinforces why our China estimates distinguish between **current market revenue** and the **broadier addressable economic pool**. A top-down EV can illustrate the long-term opportunity, but it should not be treated as revenue achievable by a single provider or within a fixed period and a bottom-up TAM reveals short-term growth momentum.


From better models to a larger intelligence market

Since ChatGPT's emergence in 2022, frontier models have advanced rapidly across reasoning, coding, multimodal understanding and agentic capabilities. As model intelligence continues to improve while inference costs decline, AI is increasingly moving beyond chatbot interactions towards the execution of economically valuable work.

More high-value scenarios for AI to unfold

We believe the next phase of AI adoption will be driven by the transition from answering questions to completing tasks. In particular, advances in reasoning, memory, tool use and workflow orchestration are enabling agents capable of performing increasingly complex and multi-step workloads. This expands AI's addressable market from information retrieval and content generation into higher-value domains from coding to office productivity, healthcare, scientific research and workflow automation in high-end manufacturing and industrials, where the economic value per task is significantly higher.

Intelligence index moving up fast since 2023



Source: Artificial Analysis (as of 31 July 2026)

Meanwhile, the emergence of agents could materially increase AI consumption intensity. Unlike traditional chatbot interactions that generate a single response, agentic workflows may involve multiple rounds of reasoning, planning, tool calls and execution. As a result, a single user request can consume substantially more tokens and computing resources than conventional AI usage.

We therefore expect foundation-model TAM expansion to be driven by three reinforcing trends:

- **Broader adoption:** increasing penetration among consumers, developers and enterprises
- **Higher-value workloads:** expansion from content generation into coding, professional work, healthcare and scientific discovery
- **Greater consumption intensity:** agentic workflows require repeated reasoning, tool calling and task execution, resulting in significantly higher model usage per user


As AI systems evolve from information tools into digital workers, the opportunity is no longer limited to the software market alone but increasingly linked to the much larger market for human intelligence and knowledge work. This creates a favourable backdrop for foundation-model developers with leading capabilities, scalable infrastructure and strong commercialisation execution.

Why we like Z.AI and MiniMax

Two leading listed China AI labs to watch

Everything above points to a market that does not need a single winner. Frontier capability and low cost are two products sold to two different buyers, and the pool that both address is large enough for multiple winners such that the two routes compound alongside each other for years rather than quarters. Z.AI and MiniMax each own one of them outright: Z.AI is the premium case, having raised price through a full cycle and kept volume; MiniMax is the efficiency case, competitive intelligence at a fraction of the price.

Leading models by AI lab: the highest Intelligence Index achieved by each AI lab (global top-15 models, as of 10 August 2026)



Source: Artificial Analysis, Daiwa compiles

Note: companies we pick is highlighted in colours while China models are in red (Dated 10 Aug 2026).

We still rank Alibaba above both, because the binding constraint is moving. For the past three years, it was model capability, and Alibaba lagged behind. Now it is gradually becoming the ability to fund compute through a cycle and to put models in front of users.

As Alibaba is catching up quickly on its capability gap, we do not believe the labs will close the funding and distribution gap on the same timescale. Both must raise before they earn, and that risk weighs more heavily on MiniMax, whose price-led strategy caps gross margin and converts the same funding need into more dilution.

Hence, **our pecking order is Alibaba, Z.AI and MiniMax.** We still like both labs. For Z.AI, the market has confused losing a benchmark with losing a franchise; for MiniMax, the market has priced its cost leadership at nothing.

Risks

Slower AI adoption and monetisation is the primary risk. If usage, paid penetration or enterprise adoption lag, revenue growth and ARR conversion could miss our forecasts.

Intensified competition. Rapid model iteration and low switching costs could shorten leadership cycles, weaken user retention and limit pricing power.

Pricing and margin pressure. Open-source models and subsidised offerings could drive down API prices, delaying margin improvement despite rising token volumes.

Compute constraints and cost inflation. Restricted chip access, export controls and higher infrastructure costs could slow model development and raise capital intensity.

Execution risk. Weak R&D allocation or failure to convert model capabilities into scalable products could lead to market-share losses of China AI players.

Financing and cash-burn risk. High R&D and compute spending may require further funding if monetisation progresses more slowly than expected.

Regulatory and geopolitical risks. Tighter AI, data and technology restrictions could raise compliance costs and constrain domestic or overseas deployment.

Copyright and IP risks. Litigation or tighter licensing requirements could increase costs and restrict model training or product features.

Z.AI (2513 HK)

 Target price: **HKD1,500.00**

 Share price (11 Aug): **HKD1,192.00** | Up/downside: **+25.8%**

5 4 3 2 1


Buy
 (initiation)

Initiation: buy the cadence, not the crown

- Repeatable model iteration sustains premium intelligence and pricing
- Capacity increase to unlock unmet demand and monetisation potential
- Initiating with a Buy (1) rating and 12M TP of HKD1,500

John Choi

(852) 2773 8730

john.choi@hk.daiwacm.com

Lily Liu

(852) 2848 4435

lily.liu@hk.daiwacm.com



Investment case: We initiate on Z.AI, also known as Zhipu, an independent Chinese foundation-model developer providing proprietary GLM models & products, with a Buy (1) rating. July's sell-off exposed the central risk in owning a frontier-model developer; investors were reminded by Kimi K3 that model leadership can change within weeks. Yet we believe the investment case does not depend on Z.AI holding every ranking. **The durable advantage is its ability to remain near the frontier** through a Tsinghua-rooted research base, frequent model iterations, coding and agent feedback from real workloads, full-stack inference engineering and domestic chip compatibility. Z.AI's repeated price increases provide stronger commercial evidence than any single benchmark, while demand continues to exceed available inference capacity.


The longer-term case is the conversion of this technical cadence into scalable revenue and eventual profitability. Annual cloud recurring revenue (ARR) reached USD1bn in July 2026, five months ahead of the company's target, while we forecast group revenue to rise from CNY724m in 2025 to CNY47.5bn in 2029E, implying a 185% CAGR. We believe cloud-based model-as-a-service (MaaS) will become the dominant revenue contributor, with group GPM recovering from a forecast trough of 31% in 2026E to 40% in 2029E as pricing, inference efficiency and utilisation improve, with adjusted profitability to be met in 2029E. **We therefore view the recent volatility as an entry point into a multi-year scaling story.** That said, we believe 1H26 results will provide the first meaningful test on ARR conversion, compute expansion and cloud-margin management, and therefore present a clearer basis for rerating.

Catalysts: 1) Launch of next flagship model with notable breakthroughs; 2) additional compute capacity and/or optimisation of task economics; 3) sustained growth in paid token and effective revenue conversion; 4) positive traction for proprietary AI products; 5) re-acceleration of on-premise with solid margins; and 6) evidence of scalable overseas business.

Valuation: We initiate coverage of Z.AI with a Buy (1) rating and a 12M TP of HKD1,500, based on a discounted EV/ARR methodology, applying a 15x multiple to our 2030E average ARR forecast of CNY82.5bn, and discounting it back by 3.5 years at a 14% cost of equity. Our multiple is below its peers, which currently trade at c.65x EV/ARR on street estimates, reflecting a discount given lack of high-quality disclosure, and conservatism on rapidly changing sector metrics and limited financial visibility.

Risks: (1) Failure to sustain frontier capability; (2) share-supply pressure in Jan 2027; (3) slower ARR growth/conversion; and (4) pricing pressure.

Share price performance



12-month range	131.50-2,410.00
Market cap (USDbn)	70.74
3m avg daily turnover (USDm)	776.73
Shares outstanding (m)	466
Major shareholder	Lhasa Junqi (11.2%)

Financial summary (CNY)

Year to 31 Dec	26E	27A	28E
Revenue (m)	5,026	13,971	27,002
Operating profit (m)	(7,682)	(12,951)	(11,121)
Net profit (m)	(6,320)	(11,006)	(8,248)
Core EPS (fully-diluted)	(13.572)	(23.173)	(17.026)
EPS change (%)	n.a.	n.a.	n.a.
Daiwa vs Cons. EPS (%)	n.m.	n.m.	n.m.
PER (x)	n.a.	n.a.	n.a.
Dividend yield (%)	0.0	0.0	0.0
DPS	0.000	0.000	0.000
PBR (x)	n.a.	n.a.	n.a.
EV/EBITDA (x)	n.a.	n.a.	n.a.
ROE (%)	n.a.	n.a.	n.a.

Source: FactSet, Daiwa forecasts

5 4 3 2 1

How do we justify our view?

Buy
(initiation)

Growth outlook

✓ ✓ ✓ ✓ ✓

Valuation

✓ ✓ ✓ ✓ ✓

Earnings revisions

✓ ✓ ✓ ✓ ✓

Growth outlook

✓ ✓ ✓ ✓ ✓

We forecast Z.AI's revenue to increase to CNY47.5bn in 2029E from CNY724m in 2025, implying a 185% CAGR, driven mainly by the conversion of rapid MaaS ARR growth into Open Platform and API revenue. Under our industry TAM base case analysis, China's AI coding TAM is expected to expand at a 266% CAGR in 2025-30E to CNY264bn, while consumer AI TAM rises at a 154% CAGR to CNY116bn, with further addressable value across enterprise and multimodal applications. These estimates remain scenario-based, but the repeated upward revisions imply faster-than-expected AI usage and monetisation.

Although the fast-moving industry landscape limits visibility, we see significant growth upside for Z.AI.


Valuation

✓ ✓ ✓ ✓ ✓

We value Z.AI on a discounted EV/ARR. As ARR is the better base while token volume and pricing are both resetting quarterly, and it is how Z.AI's mostly private peers are marked. We apply 15x to our 2030E average ARR of CNY82.5bn, add net cash, and discount 3.5 years at a 14% cost of equity to reach our target price of HKD1,500.

Our target multiple is at a discount to its peers, which currently trade at c.65x EV/ARR on street estimates, and is close to Anthropic's 16x, which we regard as the fairer comparable. Implied multiples compress as ARR scales, so marks struck at a few hundred million of ARR are not a fair anchor for a business past USD1bn.

Z.AI: revenue and revenue growth (CNYm)



Source: Company, Daiwa forecasts

Z.AI: valuation table

	2030E
Sales (CNYm)	78,061
Average ARR (CNYm)	82,504
Target EV/ARR multiple	15x
Total Enterprise Value (CNYm)	989,306
Net cash (CNYm)	7,427
Total Equity Value (CNYm)	996,733
Shares outstanding (m)	494
Fair value in 2029E (CNY)	2,017
Cost of Equity	14%
-Equity market premium (high growth stock)	7.0%
-Risk free rate (US 10-year bond yield)	4.7%
-Beta (Daiwa assumption)	1.400
Years to discount	3.5
12M Target Price (HKD)	1,500

Source: Daiwa estimates and forecasts


Earnings revisions

✓ ✓ ✓ ✓ ✓

Bloomberg consensus has revised Z.AI's 2026E revenue growth sharply higher since April, reflecting faster ARR expansion, strong paid-token demand and expectations for additional inference capacity. In contrast, GPM and adjusted EPS growth forecasts have remained subdued as higher compute and R&D requirements offset stronger monetisation.

Our 2026E revenue remains 10% above consensus, and we see further upside from capacity additions and new flagship model release, although pricing pressure and thin cloud margins mean revenue upgrades may not translate directly into earnings upgrades. The 1H26 results will be the first key test of ARR conversion and cloud economics.

Z.AI: consensus revision history on YoY earnings forecasts



Source: Bloomberg

Financial summary

Key assumptions

Year to 31 Dec	2021	2022	2023	2024	2025	2026E	2027E	2028E
Average ARR YoY % (USD m)	n.a.	n.a.	n.a.	n.a.	292.7	2,783.3	150.0	100.0
Effective blended token price YoY % (USD/m token)	n.a.	n.a.	n.a.	n.a.	(87.0)	216.7	20.0	10.0
Token usage YoY % (bn token/day)	n.a.	n.a.	n.a.	n.a.	2,915.0	810.5	108.3	81.8

Profit and loss (CNYm)

Year to 31 Dec	2021	2022	2023	2024	2025	2026E	2027E	2028E
Open platform and API	n.a.	3	12	48	190	3,992	12,351	24,702
Enterprise-level agents	n.a.	0	0	47	166	430	774	1,200
Other Revenue	n.a.	55	113	216	368	604	846	1,100
Total Revenue	n.a.	57	125	312	724	5,026	13,971	27,002
Other income	n.a.	2	10	19	15	101	279	540
COGS	n.a.	(26)	(44)	(137)	(428)	(3,486)	(9,446)	(17,391)
SG&A	n.a.	(47)	(168)	(521)	(896)	(2,513)	(4,331)	(6,210)
Other op.expenses	n.a.	(104)	(618)	(2,492)	(3,486)	(6,809)	(13,425)	(15,061)
Operating profit	n.a.	(119)	(695)	(2,818)	(4,071)	(7,682)	(12,951)	(11,121)
Net-interest inc./(exp.)	n.a.	(6)	(26)	(38)	(74)	(60)	(60)	(60)
Assoc/forex/extraord./others	n.a.	(19)	(67)	(102)	(573)	416	799	1,485
Pre-tax profit	n.a.	(144)	(788)	(2,958)	(4,718)	(7,325)	(12,212)	(9,696)
Tax	n.a.	0	0	0	0	0	0	0
Min. int./pref. div./others	n.a.	(0)	0	(2)	(20)	0	0	0
Net profit (reported)	n.a.	(144)	(788)	(2,960)	(4,738)	(7,325)	(12,212)	(9,696)
Net profit (adjusted)	n.a.	(97)	(621)	(2,467)	(3,202)	(6,320)	(11,006)	(8,248)
EPS (reported)(CNY)	n.a.	(1.081)	(5.218)	(18.511)	(29.702)	(15.732)	(25.714)	(20.015)
EPS (adjusted)(CNY)	n.a.	(0.731)	(4.110)	(15.430)	(20.070)	(13.572)	(23.173)	(17.026)
EPS (adjusted fully-diluted)(CNY)	n.a.	(0.731)	(4.110)	(15.430)	(20.070)	(13.572)	(23.173)	(17.026)
DPS (CNY)	n.a.	0.000	0.000	0.000	0.000	0.000	0.000	0.000
EBIT	n.a.	(119)	(695)	(2,818)	(4,071)	(7,682)	(12,951)	(11,121)
EBITDA	n.a.	(99)	(626)	(2,538)	(3,787)	(7,325)	(12,212)	(9,696)

Cash flow (CNYm)

Year to 31 Dec	2021	2022	2023	2024	2025	2026E	2027E	2028E
Profit before tax	n.a.	(144)	(788)	(2,958)	(4,718)	(7,325)	(12,212)	(9,696)
Depreciation and amortisation	n.a.	20	69	280	281	356	739	1,425
Tax paid	n.a.	0	0	0	0	0	0	0
Change in working capital	n.a.	10	(96)	(0)	704	2,274	938	(8)
Other operational CF items	n.a.	45	167	433	1,487	1,065	1,266	1,508
Cash flow from operations	n.a.	(68)	(648)	(2,245)	(2,246)	(3,629)	(9,269)	(6,771)
Capex	n.a.	(32)	(507)	(126)	(11)	(1,005)	(2,794)	(4,050)
Net (acquisitions)/disposals	n.a.	(11)	(202)	(23)	(52)	99	(10)	(10)
Other investing CF items	n.a.	75	(76)	101	(312)	0	0	0
Cash flow from investing	n.a.	33	(785)	(49)	(376)	(907)	(2,804)	(4,060)
Change in debt	n.a.	0	0	137	552	400	300	150
Net share issues/(repurchases)	n.a.	203	2,406	3,011	1,579	36,300	0	0
Dividends paid	n.a.	0	0	0	0	0	0	0
Other financing CF items	n.a.	(12)	57	164	483	(60)	1,940	1,940
Cash flow from financing	n.a.	191	2,463	3,312	2,614	36,640	2,240	2,090
Forex effect/others	n.a.	0	0	0	(5)	0	0	0
Change in cash	n.a.	156	1,030	1,019	(13)	32,104	(9,833)	(8,741)
Free cash flow	n.a.	(100)	(1,155)	(2,371)	(2,258)	(4,635)	(12,063)	(10,821)

Source: FactSet, Daiwa forecasts

Financial summary continued ...

Balance sheet (CNYm)

Year to 31 Dec	2021	2022	2023	2024	2025	2026E	2027E	2028E
Cash & short-term investment	n.a.	251	1,408	2,312	2,636	34,740	24,907	16,166
Inventory	n.a.	10	29	32	126	341	788	1,440
Accounts receivable	n.a.	28	416	667	699	4,131	10,335	17,755
Other current assets	n.a.	10	10	5	110	50	140	270
Total current assets	n.a.	300	1,863	3,016	3,571	39,263	36,170	35,630
Fixed assets	n.a.	41	788	866	656	1,318	3,389	6,031
Goodwill & intangibles	n.a.	21	94	90	90	87	82	74
Other non-current assets	n.a.	0	115	404	536	536	536	536
Total assets	n.a.	362	2,860	4,376	4,854	41,204	40,176	42,272
Short-term debt	n.a.	0	0	137	605	725	815	860
Accounts payable	n.a.	26	288	603	1,329	6,686	12,940	19,059
Other current liabilities	n.a.	506	3,321	7,097	10,473	10,516	12,605	14,736
Total current liabilities	n.a.	532	3,609	7,838	12,406	17,927	26,360	34,655
Long-term debt	n.a.	0	0	0	85	365	575	680
Other non-current liabilities	n.a.	10	234	493	474	1,044	2,379	4,323
Total liabilities	n.a.	542	3,843	8,331	12,965	19,335	29,313	39,657
Share capital	n.a.	15	28	36	40	36,340	36,340	36,340
Reserves/R.E./others	n.a.	(195)	(1,011)	(3,991)	(8,151)	(14,471)	(25,477)	(33,726)
Shareholders' equity	n.a.	(180)	(983)	(3,955)	(8,111)	21,869	10,863	2,614
Minority interests	n.a.	0	0	0	0	0	0	0
Total equity & liabilities	n.a.	362	2,860	4,376	4,854	41,204	40,176	42,272
EV	n.a.	476,912	475,742	474,787	474,878	443,174	453,308	462,199
Net debt/(cash)	n.a.	(251)	(1,408)	(2,175)	(1,947)	(33,651)	(23,518)	(14,626)
BVPS (CNY)	n.a.	n.a.	n.a.	n.a.	n.a.	n.a.	n.a.	n.a.

Key ratios (%)

Year to 31 Dec	2021	2022	2023	2024	2025	2026E	2027E	2028E
Sales (YoY)	n.a.	n.a.	116.9	150.9	131.9	593.9	178.0	93.3
EBITDA (YoY)	n.a.	n.a.	n.a.	n.a.	n.a.	n.a.	n.a.	n.a.
Operating profit (YoY)	n.a.	n.a.	n.a.	n.a.	n.a.	n.a.	n.a.	n.a.
Net profit (YoY)	n.a.	n.a.	n.a.	n.a.	n.a.	n.a.	n.a.	n.a.
Core EPS (fully-diluted) (YoY)	n.a.	n.a.	n.a.	n.a.	n.a.	n.a.	n.a.	n.a.
Gross-profit margin	n.a.	54.6	64.6	56.3	41.0	30.6	32.4	35.6
EBITDA margin	n.a.	n.a.	n.a.	n.a.	n.a.	n.a.	n.a.	n.a.
Operating-profit margin	n.a.	n.a.	n.a.	n.a.	n.a.	n.a.	n.a.	n.a.
Net profit margin	n.a.	n.a.	n.a.	n.a.	n.a.	n.a.	n.a.	n.a.
ROAE	n.a.	n.a.	n.a.	n.a.	n.a.	n.a.	n.a.	n.a.
ROAA	n.a.	n.a.	n.a.	n.a.	n.a.	n.a.	n.a.	n.a.
ROCE	n.a.	n.a.	n.a.	n.a.	n.a.	n.a.	n.a.	n.a.
ROIC	n.a.	27.5	49.2	66.1	50.3	70.3	n.a.	90.2
Net debt to equity	n.a.	139.2	143.3	55.0	24.0	n.a.	n.a.	n.a.
Effective tax rate	n.a.	0.0	0.0	0.0	0.0	0.0	0.0	0.0
Accounts receivable (days)	n.a.	n.a.	651.1	632.8	344.1	175.4	189.0	189.8
Current ratio (x)	n.a.	0.6	0.5	0.4	0.3	2.2	1.4	1.0
Net interest cover (x)	n.a.	n.a.	n.a.	n.a.	n.a.	n.a.	n.a.	n.a.
Net dividend payout	n.a.	n.a.	n.a.	n.a.	n.a.	n.a.	n.a.	n.a.
Free cash flow yield	n.a.	n.a.	n.a.	n.a.	n.a.	n.a.	n.a.	n.a.

Source: FactSet, Daiwa forecasts

Company profile

Z.AI, commonly known as Zhipu AI, is an independent Chinese foundation-model developer providing proprietary GLM models & products, mainly cloud-based MaaS services and on-premise enterprise AI solutions.

From a Tsinghua research lab to leading frontier foundation AI model company

Overview: enterprise AI, coding, and agents form the core commercial focus

Z.AI: China's listed claim on frontier intelligence

Z.AI Co., Ltd ("Z.AI") was founded in 2019 and originated from Tsinghua University's Knowledge Engineering Group (KEG). The company focuses on the R&D and commercialisation of foundation models, with its primary growth driver transitioning from project-based deployments towards usage-driven, cloud-based AI consumption. Over time, the company has developed model capabilities spanning language, coding and multimodal and agentic AI applications.

Compared with several domestic peers that place greater emphasis on consumer-facing applications, Z.AI's commercialisation efforts have historically been concentrated in enterprise, government and developer-oriented markets. Coding, enterprise workflows and agent-related complex and high-value use cases are frequently highlighted by management as key areas of focus.

Since listing in Hong Kong, the company has become one of the few publicly traded companies offering direct exposure to China's domestic foundation model sector.

History: Business Evolution

The company's development can broadly be divided into three stages.

1) Academic incubation stage (2019-22)

During its initial years, the company's primary focus was model research, open-source ecosystem development and technology validation. Open-source model releases helped establish an early developer community and increased market awareness of the company's model capabilities. Commercialisation efforts remained relatively limited during this period.

2) Enterprise deployment stage (2023- late-2025)

As domestic demand for localised AI capabilities increased, the company expanded its enterprise deployment business, particularly among financial institutions, state-owned enterprises and public-sector clients. Revenue growth during this period was driven primarily by project-based and on-premise deployments.

3) Cloud commercialisation stage (late-2025-present)


Beginning in late-2025, the company increasingly shifted resources toward cloud-based MaaS services. While on-premise deployment remains an important source of customer acquisition and recurring enterprise relationships, cloud-based revenue from API usage leveraging its strong coding model and agent ability will likely be the main revenue drivers in 2026 and onwards.

In our view, this transition is important because it changes the company's revenue model from largely project-based income toward a more scalable, usage-based framework.

Key developments shaped eye-catching stock performance since listing

We believe the following key issues shaped Z.AI's eye-catching stock performance since it listed on 8 January.

Z.AI: YTD stock performance catalysts



Source: FactSet, Daiwa

Note: green/red/blue represents upside/downside/neutral catalysts

Share price hit HKD2,980 in June 2026, up 24.6x from the IPO price

Model and industry upgrades drove the continuous rerating

12 February 2026 | GLM-5 launch: Ranked No.4 globally and No.1 by open-source on coding benchmarks, on par with Claude Opus 4.5. Stock +29% close.

7 April 2026 | Sector catalyst: Anthropic hits USD30bn ARR, surpassing OpenAI and claimed to turn profit in 2Q26. Rerated global LLM monetisation upside, supporting Z.AI's coding-led valuation premium.

8 April 2026 | GLM-5.1 launch: First open-source model with 8hr autonomous task capability, outperforming Claude Opus 4.6 on SWE-bench Pro. Stock +14% close.

17 June 2026 | GLM-5.2 launch: 1M-token context window, No. 2 globally and No.1 commercially accessible on Code Arena, benchmarked vs Claude Opus 4.8. Stock peaked at HKD2,980 on Jun 22, market cap >HKD1.3tn.

Late-June 2026 | Global AI beta derating on concern over AI capex sustainability and pre-emptive profit-taking ahead of unlock of first-batch locked-up shares.

9 July 2026 | HKD31.4bn placement: 19.78mn new H shares at HKD1,588 (13% discount) to fund R&D and compute; strong institutional oversubscription. Stock +13.4% on 8 July, +11.3% on 9 July.

16 July 2026 | Moonshot launches Kimi K3, overtaking GLM-5.2 as Code Arena #1. Z.AI -28.5% on July 17, -19.6% on 20 July on leadership risk repricing.

21 Jul 2026 | 1GW domestic compute: announced 1GW level all-domestic-chip data center + XCore Sigma acquisition, reinforcing supply chain self-reliance. Shares rebounded +36.9% on day.

Well-balanced shareholding structure and scientist-led corporate governance

Shareholder structure: unusually balanced and well-sponsored

In our view, Z.AI's shareholding structure is deliberately constructed, benefiting from both fragmentation and composition. Post-placement, no external shareholder exceeds 5.2% of enlarged capital, and the top ten non-insider holders account for under 30%, limiting single-sponsor-exit risk.


Shareholder composition matters. **Notably, government and university-backed capital comprise c.23.5% of shareholders**, spanning national vehicles (eg, the NSSF Zhongguancun Innovation Fund) and municipal sponsors across Haidian, Daxing, Tianjin Wuqing, Hangzhou, Pudong, Zhuhai and Chengdu. We view this cohort as **more duration-tolerant than conventional financial sponsors**, given mandates weighted toward industrial policy rather than fund-life IRR, which in our view should lower the odds of coordinated selling into the January 2027 unlock.

A further c.14.5% is owned by internet and technology capital, including Meituan, Ant, Tencent, Alibaba, Shunwei, Lenovo, and Saudi Aramco's Prosperity7. This cohort supplies industry input that a purely financial register cannot. Against CNY1.6bn of compute procurement in 2024 vs. CNY312m of revenue, having major domestic cloud operators on the register is a supply-side asset as much as a validation signal, as nine of China's ten largest internet companies already use GLM models as of March 2026.

Underlying this is the Tsinghua lineage: the Huakong complex at c.5.2%, plus Li Juanzi and Tang Jie as faculty co-founders. We regard this as the most durable structural advantage, **underwriting recruitment and research throughput from the KEG lab** that peers must otherwise replicate through compensation alone.

Employee ownership is also deep. The Huihui and Zhideng platforms hold a combined c.14.5% of enlarged capital for 451 employees and advisers, against a headcount of c.1,100. **Close to half the workforce carries direct economic exposure, a stronger retention tool** than a standard option pool in a market bidding aggressively for frontier talent.

Estimated shareholding structure by investor type



Source: Company, Bloomberg, Daiwa estimates

Regarding **voting rights**, control is contractual, not structural, which distinguishes Z.AI from the dual-class template most Chinese internet issuers adopt in Hong Kong. Although the company is listed under Chapter 18C, which permits weighted voting rights, it did not adopt them; there is no "-W" suffix and each share carries one vote. Control instead rests

Shareholder composition is fragmented and well-balanced

on an acting-in-concert agreement between Beijing Lianpai, two employee platforms, and five individuals, collectively controlling 33.03% pre-IPO, 30.22% at listing and, by our estimates, c.28.9% currently.

Research direction and commercial execution are led in parallel

Scientist-led research combined with commercial execution team

Z.AI's governance structure combines scientist-led research with an executive team responsible for commercial execution. Founder and Chief Scientist Dr. Tang Jie oversees technology strategy, frontier research, core-model development, AI safety and the R&D organization. He is also a Tsinghua professor and Director of its Fundamental Model Research Center, having published over 300 academic papers.

Chairman Dr. Liu Debing is responsible for strategic planning and overall management, while CEO Dr. Zhang Peng oversees business development, R&D and daily operations. Both have extensive research or working experience at Tsinghua. Therefore, there is no strict division between scientists and professional managers. Rather, Dr. Tang sets the longer-term technology direction, while the executive team connects research, product development and commercialisation.

Z.AI: core management team

Executive	Position	Relevant background and responsibilities
Dr. Tang Jie	Founder and Chief Scientist	Tsinghua professor; leads technology strategy, frontier research and core-model innovation
Dr. Liu Debing	Co-founder and Chairman	Former Tsinghua senior engineer; leads strategy and overall management
Dr. Zhang Peng	Co-founder and CEO	Worked at Tsinghua from 2005 to 2020; oversees business development, R&D and operations
Wang Shaolan	Co-founder and Deputy General Manager	Oversees business development, financing and operations

Source: Company, Daiwa

Tsinghua-linked talent supports a high iteration cadence

Z.AI's Tsinghua origins provide an important talent pipeline. The company had 657 R&D employees at end-2024, representing 74% of total staff. At end-2025, total permanent headcount reached 1,094, with an R&D team of c.700 per our estimate, supplemented by c.**600 interns**, many of whom are working on a full-time basis and include postgraduate and doctoral-level researchers.

This talent base supports work across model architecture, coding, reinforcement learning, synthetic data and inference engineering. We notice that Z.AI has historically released minor model upgrades every 1-2 months and major iterations every 7-8 months. We view this repeatable iteration capability as more important than leadership in any individual benchmark.

Z.AI: product launch log, 2021 to 2026 YTD

Date	Type	Launch	What it is
2021	Foundation models	GLM-10B	First-generation bilingual pre-trained model on the in-house GLM architecture
2021	Products & monetisation	MaaS platform	Model-as-a-service API platform; the original commercial channel
Aug-22	Foundation models	GLM-130B	130bn-parameter bilingual base model; established the open-weight posture
Sep-22	Coding	CodeGeeX	Multilingual code-generation model and IDE plug-in; earliest vertical product
May-23	Multimodal generation	VisualGLM	First multimodal release; open image-text model
Aug-23	Products & monetisation	ChatGLM (Z.AI Qingyan) app	Consumer assistant app; first direct-to-user surface
Jan-24	Foundation models	GLM-4	First flagship generation shipped with a full API and tooling stack
Aug-24	Multimodal generation	GLM-4V	Vision-language variant of the GLM-4 flagship
Oct-24	Agents	AutoGLM	Phone / GUI agent executing multi-step tasks directly on the user's device
Oct-24	Multimodal generation	GLM-4-Voice	End-to-end speech-in / speech-out model
Oct-24	Products & monetisation	AutoGLM app	Consumer-facing agent app built on the AutoGLM stack
Apr-25	Agents	GLM-Z1-Rumination	Deep-reasoning model aimed at research-style, long-deliberation tasks
Jul-25	Foundation models	GLM-4.5	Flagship refresh explicitly tuned for tool-use and agentic workloads
2H25	Products & monetisation	GLM Coding Plan	Coding subscription; priced at a deep discount to Western equivalents
2025	Agents	AutoGLM 2.0	Second-generation agent stack with broader app and browser coverage
2025	Multimodal generation	CogVideoX-3	Third-generation text and image-to-video model
2025	Multimodal generation	GLM-Realtime	Low-latency real-time voice and video interaction model
Dec-25	Foundation models	GLM-4.7	Incremental flagship upgrade; also served as the coding tier
Dec-25	Multimodal generation	GLM-TTS	Dedicated speech-synthesis model
Feb-26	Agents	GLM-5-Turbo	Speed and cost-optimised GLM-5 tier for high-volume agent loops
Feb-26	Foundation models	GLM-5	New flagship generation; ~744bn-parameter MoE, roughly 2x the GLM-4.5 scale
Feb-26	Multimodal generation	GLM-Image	Image-generation model, completing the modality set
1Q26	Coding	GLM-5V-Turbo	Speed-optimised vision-capable tier serving coding and UI-aware agent work
Mar-26	Products & monetisation	AutoClaw	One-click local agent client for macOS/Windows, shipping 50+ pre-built skills
Apr-26	Foundation models	GLM-5.1	Long-horizon iteration; open-weighted about two weeks after subscription release
Apr-26	Products & monetisation	Claw Plan	Subscription tier attached to the AutoClaw agent stack
May-26	Products & monetisation	GLM-5.1-HighSpeed	Latency-optimised variant of GLM-5.1 for throughput-sensitive deployments
May-26	Products & monetisation	AutoClaw self-evolution	Self-improving / skill-learning upgrade to the AutoClaw agent
Jun-26	Foundation models	GLM-5.2	Current flagship; 1m-token context, open weights under a permissive licence
Jun-26	Products & monetisation	ZCode	Agentic development environment positioned as the native harness for GLM-5.2

Source: Company, compiled by Daiwa

R&D expenses were CNY3.18bn in 2025, driven by higher staff costs and computing-service expenditure. This was **more than four times the reported revenue**, highlighting that Z.AI remains a research-intensive frontier lab rather than a mature business. Equity incentives help retain core researchers, although competition for leading AI talent remains a material risk.

Independent positioning supports enterprise adoption

High model and product iteration cadence

Unlike models developed within major internet ecosystems, Z.AI operates as an independent provider without a captive or overlapping traffic platform. This allows the company to work across and serve internet platforms, large enterprises, device manufacturers and domestic chip suppliers with limited conflicts of interest or customer overlap. Adoption among large domestic customers is not only broad but unusually fast: ByteDance, Alibaba and Tencent all officially integrated GLM-5 **within 24 hours of release**, and **nine of China's ten** largest internet companies have embedded GLM in daily operations as of March 2026.

This positioning is also relevant for enterprise and public-sector customers seeking data control, local deployment and infrastructure compatibility. However, Z.AI lacks the captive cloud resources and distribution available to internet-platform-backed peers. Its independence will only become a durable advantage if the company can convert model capability into deeper workflow integration, proprietary usage data and higher customer switching costs.

Beyond China, Z.AI's coding, agent and enterprise model services now reach **218 countries and territories**. Distribution outside China has broadened well beyond Z.AI's own endpoint. GLM models are served on Google Vertex AI, AWS Bedrock, Fireworks AI and Cerebras, as well as through aggregators including OpenRouter and Vercel, where GLM is among the top-ranked paid models. GLM has also become the default model on coding platforms such as Windsurf and OpenCode. We regard third-party distribution as strategically double-edged: it accelerates adoption and brand formation, but usage routed through other hosts does not necessarily generate revenue for Z.AI.

Business structure, product portfolio and growth strategy

Cloud-based MaaS(API) is becoming the primary growth engine

By business segments, Z.AI generates revenue from three principal business lines: Open Platform and API services, enterprise-level agents, and enterprise foundation-model deployments. In 2025, enterprise foundation models remained the largest contributor at CNY365.7m, or 50.4% of revenue, followed by enterprise agents at CNY165.7m, or 22.9%, and Open Platform and API at CNY190.4m, or 26.3%. However, Open Platform and API grew by 293% YoY, materially faster than the 71% YoY growth of enterprise foundation-model deployments.

By market, as of 1H25, Z.AI generated **90% of revenue from China and <10% from overseas markets**. Given model competence relative to cost, benchmarked to its overseas peers, we see great potential from overseas markets. However, due to Z.AI's open-source strategy and service capacity, meaningful revenue contribution will take more patience but will be a key catalyst if it happens.

Z.AI: 2025 financials and key business segments

Business	Revenue model	Revenue (CNYm)	% of revenue	GPM	Strategic role
Open Platform and API	Token consumption and subscriptions	203	26%	19%	Primary scalable growth engine
Enterprise-level agents	Customized workflow and agent solutions	90	23%	52%	Deeper integration and higher switching costs
Enterprise foundation models	Mainly on-premise project deployment	150	50%	47%	Established enterprise customer base
Technical services and others	Consulting and technical support	1	0%	92%	Small and non-core

Source: Company, Daiwa

Surging cloud-based MaaS is reshaping business structure

We believe this mix is changing rapidly with fast expansion of API ARR in 1H26 following high-ranking model intelligence scoring of its GLM series from February. By revenue mode, we estimate API services contributed more than **90% of cloud ARR** as of July 2026, with GLM Coding Plan subscriptions **under 10%**. That is close to the structure of the business Z.AI is trying to become: Anthropic derives **c.80% of its revenue from API calls**. Subscription products are better understood as developer acquisition and brand formation than as a revenue line, though their reach is real, with the Claw Plan reaching 100,000 subscribers within two days of its March 2026 launch and 400,000 within 20 days. And LatePost just reported on 10 Aug, that Z.AI's developer tool "ZCode", benchmarking Code X of Claude, has accumulated over 1 million users within one month of release.


Open Platform and API: scalable, but dependent on model leadership, computing capacity and token efficiency

The Open Platform converts Z.AI's foundation-model capabilities into standardized API and subscription based MaaS services. Usage-based revenue is recognised as customers consume computing resources, while subscription revenue is recognized over the contract period. Compared with project-based deployments, we believe **this model can support more customers without requiring a similar increase in delivery personnel or project time**, although incremental usage also creates corresponding inference costs.

Cloud ARR increased from c.USD4m in early 2025 to USD250m by March 2026, and to USD1bn in July 2026, hitting its year-end ARR target five months in advance; token consumption rose **c.15x over six months**. In our view, the speed of this increase confirms strong demand, but ARR should not be treated as equivalent to recognized revenue


because it annualises the prevailing revenue run rate rather than representing completed historical sales.

Z.AI ARR trend and forecasts



Source: Company, Daiwa estimates

Blended price of Z.AI LLM API



Source: Company, Daiwa estimates

Note: Our blended price = input × 0.7 + output × 0.3. The 70/30 ratio reflects a generalised production workload.

MaaS price increases are being absorbed, which is the clearest evidence of pricing power in the sector. Z.AI raised pricing across its GLM Coding Plan and API services during 2026, with YTD API price hike of **over two-fold**.

Z.AI: model API pricing log

Within 6 months, Z.AI's frontier model API pricing doubled, while coding pricing increased c.6-fold

Model	launch time	input price (per 1m tokens)	output price (per 1m tokens)	blended price (USD/per 1m tokens)	% change vs. last ver.
GLM-4*	Jan-24	CNY100	CNY100	14.29	
GLM-4-Air*	5-Jun-24	CNY1.0	CNY1.0	0.14	-99%
GLM-4-AirX*	5-Jun-24	CNY 10	CNY 10	1.43	900%
GLM-4-Plus*	29-Aug-24	CNY 50	CNY 50	7.14	400%
GLM-4.5	28-Jul-25	CNY0.8	CNY2.0	0.17	-98%
GLM-4.5-Air	28-Jul-25	USD0.2	USD1.1	0.47	184%
GLM-4.6	30-Sep-25	USD0.6	USD2.2	1.08	130%
GLM-4.7	22-Dec-25	USD0.6	USD2.2	1.08	0%
GLM-5	11-Feb-26	USD1.0	USD3.2	1.66	54%
GLM-5 Turbo	16-Mar-26	USD1.2	USD4.0	2.04	23%
GLM-5.1	8-Apr-26	USD1.4	USD4.4	2.30	13%
GLM-5.2	16-Jun-26	USD1.4	USD4.4	2.30	0%

Source: Company, Daiwa collects

Note: * for these versions a single price based on total processed tokens, rather than separate input and output rates were applied

We estimate that the **GLM 5.2 API pricing** is now almost **one-third of Kimi K3** and **c.3x that of MiniMax M3** as of early August. The company's ability to sustain user growth following price increases indicates that customers are willing to pay for stronger coding and agent performance, although this pricing power remains conditional on Z.AI maintaining relative model quality.

Global frontier LLM API pricing comparison (as of Aug 2026)

Rank*	Model	Creator	Overall Score	Input price^	Output price^	Blended price^	Coding	Agentic	Multimodality
50	DeepSeek V4 Flash	DeepSeek	60	0.1	0.3	0.2	42	43	-
19	Hy3	Tencent	68	0.1	0.5	0.3	63	61	-
25	GPT-5.6 Luna	OpenAI	67	0.2	1.2	0.5	72	58	67
28	GPT-5.4 nano	OpenAI	66	0.2	1.3	0.5	43	41	17
17	MiMo-V2.5-Pro	Xiaomi	69	0.4	0.9	0.6	55	42	-
18	MiniMax M3	MiniMax	69	0.3	1.2	0.6	48	40	48
40	MiniMax M2.7	MiniMax	63	0.3	1.2	0.6	46	35	-
49	DeepSeek V4 Pro	DeepSeek	60	0.4	0.9	0.6	42	43	-
27	Qwen3.7 Plus	Alibaba	66	0.3	1.3	0.6	58	43	73
46	GLM-4.7	Z.AI	60	0.4	1.8	0.8	51	51	-
34	Qwen3.6 Plus	Alibaba	64	0.3	2.0	0.8	53	37	68
32	Gemini 3.5 Flash-Lite	Google	65	0.3	2.5	1.0	52	40	64
36	Grok 4.3	xAI	64	1.3	2.5	1.6	38	6	65
31	GLM-5	Z.AI	65	1.0	3.2	1.7	58	54	-
29	GLM-5-Turbo	Z.AI	66	1.2	4.0	2.0	-	-	-
44	GLM-5V-Turbo	Z.AI	62	1.2	4.0	2.0	57	-	-
7	Muse Spark 1.1	Meta	76	1.3	4.3	2.2	64	68	80
16	Muse Spark	Meta	70	1.3	4.3	2.2	63	58	77
22	GLM-5.1	Z.AI	67	1.4	4.4	2.3	56	49	-
41	GLM-5.2	Z.AI	63	1.4	4.4	2.3	64	57	-
13	Qwen3.7 Max	Alibaba	72	1.5	4.4	2.4	54	40	-
24	InKling	Thinking Machines Lab	67	1.9	4.7	2.7	43	46	53
8	Grok 4.5	xAI	76	2.0	6.0	3.2	57	61	69
9	Gemini 3.6 Flash	Google	75	1.5	7.5	3.3	64	46	75
37	Gemini 3.5 Flash	Google	64	1.5	9.0	3.8	61	45	85
33	Claude Sonnet 5	Anthropic	65	2.0	10.0	4.4	68	59	80
12	GPT-5.6 Terra	OpenAI	72	2.0	12.0	5.0	65	73	76
21	Gemini 3 Pro	Google	67	2.0	12.0	5.0	61	58	74
30	GPT-5.3 Codex	OpenAI	66	1.8	14.0	5.4	63	61	-
10	GPT-5.4	OpenAI	73	2.5	15.0	6.3	58	56	69
5	Kimi K3	Moonshot AI	80	3.0	15.0	6.6	77	66	90
35	Claude Sonnet 4.6	Anthropic	64	3.0	15.0	6.6	58	51	58
2	Claude Opus 5	Anthropic	83	5.0	25.0	11.0	78	72	89
6	Claude Opus 4.8	Anthropic	77	5.0	25.0	11.0	70	63	90
14	Claude Opus 4.7 (Adaptive)	Anthropic	72	5.0	25.0	11.0	64	64	51
15	Claude Opus 4.7	Anthropic	71	5.0	25.0	11.0	68	55	-
20	Claude Opus 4.6	Anthropic	68	5.0	25.0	11.0	61	56	62
38	Claude Opus 4.6 (Adaptive)	Anthropic	63	5.0	25.0	11.0	60	56	-
39	Claude Opus 4.5	Anthropic	63	5.0	25.0	11.0	61	36	29
4	GPT-5.6 Sol	OpenAI	81	5.0	30.0	12.5	78	76	85
11	GPT-5.5	OpenAI	72	5.0	30.0	12.5	71	64	69
1	Claude Mythos 5	Anthropic	83	10.0	50.0	22.0	80	66	86
3	Claude Fable 5	Anthropic	83	10.0	50.0	22.0	80	65	64
26	GPT-5.2 Pro	OpenAI	66	25.0	150.0	62.5	-	-	-
42	GPT-5.5 Pro	OpenAI	63	30.0	180.0	75.0	-	60	-
48	GPT-5.4 Pro	OpenAI	60	30.0	180.0	75.0	-	57	89

Source: BenchLM, companies, Daiwa collect and estimates

Note: Our blended price = input × 0.7 + output × 0.3. The 70/30 ratio reflects a generalised production workload.

*Ranking of intelligence score per BenchLM, we selected Top 49 and added DeepSeek V4 Flash 0731 as No.50.

^USD/m token

On the GLM Coding Plan side, it raised prices by 30% in February 2026 and simultaneously withdrew first-purchase discounts, then meaningfully lifted prices in April, and in July again, then simultaneously removed purchase limits; its July Lite GLM Coding Plan price (CNY118/month) is **c.6x** higher than Feb (CNY20/month). That said, paid developers still surpassed 242,000. We think this is more important than any benchmark position – a price increase that holds is the only unambiguous test of whether a capability premium is real.

Global frontier LLM subscription/coding plan pricing comparison – full tier Detail (USD per user per month, as of Aug 2026)

Provider	HQ	Coding product	Segment	Tier	Currency	List price before promo (USD/mo)	List price before promo (USD/mo)	Seat minimum	Coding agent included?	Published usage allowance	Notes
Z.AI (Z.ai)	China	GLM Coding Plan	Individual	Lite	CNY	118	17	1	Yes	c.80 prompts / 5h; c.400 / week; 100 MCP calls / mo	30% intro promo takes monthly to USD12.60; annual USD151.20
Z.AI (Z.ai)	China	GLM Coding Plan	Individual	Pro	CNY	538	79	1	Yes	c.400 prompts / 5h; c.2,000 / week; 1,000 MCP calls / mo	c.5x Lite allowance; annual USD604.80
Z.AI (Z.ai)	China	GLM Coding Plan	Individual	Max	CNY	1,078	159	1	Yes	c.1,600 prompts / 5h; c.8,000 / week; 4,000 MCP calls / mo	c.20x Lite allowance; annual USD1,344
Z.AI (Z.ai)	China	GLM Coding Team Plan	Team / Enterprise	Standard Seat	CNY	598	88	2	Yes	15,000 credits / 5h; 66,000 / week	n/a – seat price not published in docs; shown at Z.ai checkout only
Z.AI (Z.ai)	China	GLM Coding Team Plan	Team / Enterprise	Premium Seat	CNY	1,198	176	2	Yes	35,000 credits / 5h; 155,000 / week	n/a – price not published. Adds early model access + peak-hour priority
Anthropic	USA	Claude Code (in Claude plans)	Individual	Pro	USD	20	20	1	Yes	5-hour rolling session + weekly cap	Official page lists USD17/mo on annual (USD200 billed up front)
Anthropic	USA	Claude Code (in Claude plans)	Individual	Max 5x	USD	100	100	1	Yes	5x Pro usage per 5-hour session	Monthly billing only
Anthropic	USA	Claude Code (in Claude plans)	Individual	Max 20x	USD	200	200	1	Yes	20x Pro usage per 5-hour session	Monthly billing only
Anthropic	USA	Claude Code (in Claude plans)	Team / Enterprise	Team – Standard seat	USD	25	25	2	Yes	More usage than Pro	Team plan covers 2–150 people; USD20/seat on annual
Anthropic	USA	Claude Code (in Claude plans)	Team / Enterprise	Team – Premium seat	USD	125	125	2	Yes	5x Standard seat usage	USD100/seat on annual; seat types can be mixed
Anthropic	USA	Claude Code (in Claude plans)	Team / Enterprise	Enterprise	USD	20	20	n/a	Yes	500k context on default model; usage billed on top	USD20/seat/mo PLUS usage at API rates. Annual billing only
OpenAI	USA	Codex (in ChatGPT plans)	Individual	Go	USD	8	8	1	Limited	Lightweight coding tasks only	Cheapest paid entry; Codex access is capped
OpenAI	USA	Codex (in ChatGPT plans)	Individual	Plus	USD	20	20	1	Yes	15–80 GPT-5.5 local msgs / 5h; 10–60 cloud tasks	Practical individual entry point for Codex
OpenAI	USA	Codex (in ChatGPT plans)	Individual	Pro 5x	USD	100	100	1	Yes	80–400 GPT-5.5 local msgs / 5h	'From USD100/mo'; promo has run at 10x
OpenAI	USA	Codex (in ChatGPT plans)	Individual	Pro 20x	USD	200	200	1	Yes	300–1,600 GPT-5.5 local msgs / 5h	20x Plus limits
OpenAI	USA	Codex (in ChatGPT plans)	Team / Enterprise	Business seat	USD	25	25	2	Yes	Plus-level per-seat limits + purchasable credits	USD20/user/mo on annual. Replaced the old Team plan (April 2026)
Kimi (Moonshot AI)	China	Kimi Code (in Kimi membership)	Individual	Moderato	CNY	19	3	1	Yes	60 agent credits; entry-level Kimi Code credits	Annual billing c.20% off
Kimi (Moonshot AI)	China	Kimi Code (in Kimi membership)	Individual	Allegretto	CNY	39	6	1	Yes	Higher credit + quota allocation; 1M context	Annual billing c.20% off
Kimi (Moonshot AI)	China	Kimi Code (in Kimi membership)	Individual	Allegro	CNY	99	15	1	Yes	Agent Swarm; Kimi Code at 5x	Annual billing c.20% off
Kimi (Moonshot AI)	China	Kimi Code (in Kimi membership)	Individual	Vivace	CNY	199	29	1	Yes	720 agent credits; Kimi Code at 30x	USD159/mo confirmed on annual billing
Kimi (Moonshot AI)	China	Kimi Business	Team / Enterprise	Business seat	CNY	50	7	5	Yes	Allegretto-level benefits per seat	USD599 per seat per year, annual billing only, 5-seat minimum
Qoder-CN (Alibaba)	China	Native coding tool, Qwen model only	Individual	Pro	CNY	59	9	1	Yes	2000 credits/month	Only Qwen models, support VSCode/JetBrains/Qoder-CN
Qoder-CN (Alibaba)	China	Native coding tool, Qwen model only	Individual	Pro+	CNY	169	25	1	Yes	6000 credits/month	Only Qwen models, support VSCode/JetBrains/Qoder-CN
Qoder-CN (Alibaba)	China	Native coding tool, Qwen model only	Team / Enterprise	Teams	CNY	99	15	1	Yes	3000 credits/seat/month, can be shared among team	Support IDE, CLI with seat management and usage monitoring
Qoder-CN (Alibaba)	China	Native coding tool, Qwen model only	Team / Enterprise	Enterprise VPC	CNY	199	29	50	Yes	3000 credits/seat/month, can be shared among team	Support IDE, CLI with seat management and usage monitoring
Bailian Coding Plan (Alibaba)	China	Coding tool supporting various models	Individual	Lite	CNY	40	6	1	Yes	suspended since April 2026	Comprise Qwen, GLM, Kimi, Minimax top models, limited supply availability
Bailian Coding Plan (Alibaba)	China	Coding tool supporting various models	Individual	Pro	CNY	200	29	1	Yes	6k request/5h; 45k requests/week; 90k requests/month	Comprise Qwen, GLM, Kimi, Minimax top models, limited supply availability

Source: companies, Daiwa

Note: all figures are publicly listed prices, gathered 4 August 2026.

Enterprise agents: a route towards deeper workflow integration

Enterprise agents combine foundation models with customer knowledge bases, task planning, memory and external tools. In 2025, Revenue from this business increased by 249% YoY to CNY165.7m, while gross margin rose to 52.3% from 49.3%. The relatively stable margin profile suggests that enterprise agents can provide better economics than the current cloud API business, in our view, although the higher degree of customisation may limit scalability.

On-premise deployment contributed 73% of the company's revenue in 2025

We view agents as strategically important because they can move Z.AI beyond interchangeable model calls towards fuller workflow integration. **Coding and enterprise agents can generate feedback from real tasks**, including whether code executes successfully or whether a multi-step process is completed. This provides more structured reinforcement signals than general conversation data and may support subsequent model improvement. We believe this coding-data feedback loop will enhance Z.AI's competitive positioning.

Agents may also increase customer switching costs over time. An enterprise can change API providers relatively easily if models are accessed through standardised interfaces. Switching becomes more complex when the model is integrated with proprietary data, tools, memory and workflow controls. However, Z.AI's agent business remains at an early stage, and the company has not disclosed sufficient retention or customer-cohort data to quantify the resulting improvement in revenue durability.

Enterprise foundation models deployment: a stable base rather than the main growth driver

Enterprise foundation models revenue increased by 70% YoY to CNY366m in 2025 and represented 50% of total revenue. The business principally serves customers requiring models and related services to be deployed within their own infrastructure, including institutions with data-security, localisation or customisation requirements. Revenue is generally recognised when the model and related services are delivered and accepted by the customer.

While the segment provides established enterprise relationships, its economics weakened in 2025. Segment gross margin declined from 69.5% in 2024 to 46.9% in 2025 as Z.AI committed more delivery resources to increasingly diverse customer requirements. We also believe that rising customisation costs, higher hardware costs within fixed customer budgets and strategic pricing for selected projects are reasons behind margin pressure; we believe this pressure will persist in 2026.

We therefore expect on-premise deployment to remain an important entry point for financial institutions, SOEs and public-sector customers, but not the primary source of long-term operating leverage. The main value of the business is likely to lie in customer access, local deployment experience and potential migration towards agents or cloud consumption, rather than project revenue growth alone.

What are Z.AI's competitive advantages?

No.1: coding and agent capabilities anchor the product strategy


Z.AI is considered to be the coding champion of China LLM labs

The GLM family covers language, coding, multimodal and agent applications, but **coding has become Z.AI's clearest area of differentiation**. Management has positioned coding as an entry point into broader agentic workloads because software execution provides relatively objective feedback and allows models to interact with external tools and digital environments.

Z.AI's coding advantage is supported by post-training in over 10,000 live software-engineering environments, using its proprietary importance-sampling algorithm and Slime asynchronous reinforcement-learning framework. Unlike conversational data, code provides objective and scalable reward signals because outputs can be verified through compilation and testing. We therefore view coding as a research and data strategy rather than only a product category.


In addition, management frames Z.AI's strategy as **AGI commercial value = intelligence ceiling × token consumption scale**. Capability supports pricing power, while inference efficiency broadens affordable usage. In our view, their roadmap suggests progression from coding assistants to agentic engineering and ultimately autonomous digital engineers capable of completing long-horizon tasks. Concepts such as LLM-OS and Token Architecture Capability indicate that Z.AI intends to position GLM as an intelligence infrastructure layer rather than only an application provider.

Z.AI's GLM-5.2 vs. global frontier models on coding and agentic benchmarks



Source: Company, respective benchmark entities, compiled by Daiwa

Frontend code arena scores, July 2026



Source: Arena.ai, Daiwa

Sustained coding & agent capability leadership

In late-2025, GLM-4.7 marked Z.AI's first meaningful coding capability breakthrough among global developers, positioning the company as a credible alternative to Anthropic's Claude series in coding and agent workflows rather than merely a domestic challenger to GPT-class models. Then the cadence iteration of the GLM 5 series proved that the stand-out of GLM 4.7 was not a one-off performance.

Launched in Feb 2026, GLM-5 ranked No.4 globally on coding benchmarks, on par with Claude Opus 4.5. Launched in April 2026, GLM-5.1 improved performance in coding and long-horizon tasks and is ranked No.4 again on Artificial Analysis as of May 2026. GLM-5.2 subsequently introduced a 1m-token context window and was ranked No. 2 globally on Code Arena shortly after release in June 2026.

These rankings supported Z.AI's position as a leading Chinese coding-model provider before Kimi K3 moved ahead on the frontend coding benchmark in July.

We do not treat benchmark leadership as a permanent moat. Rankings can change rapidly following competitor releases and may not fully capture reliability, latency or cost in commercial workloads. More relevant indicators, in our view, are whether Z.AI can repeatedly produce competitive models, maintain paid usage and convert developer adoption into enterprise workflows.

No.2: inference engineering and domestic chip compatibility address the compute constraint

Z.AI's technical investment extends beyond model training into inference efficiency, with disclosed optimisation across KV-cache management, sparse attention, prefill-decode separation, distributed communication, speculative decoding, GPU utilisation and domestic-chip adaptation. For example, dynamic sparse attention reportedly halves deployment costs, while GLM-5.2's IndexShare reduces per-token FLOPs by 2.9x at a 1mn-token context. Together, these optimisations aim to increase throughput and reduce the unit cost of serving long-context coding and agent workloads.

Disclosed engineering optimising techniques

Technique	What it does	Disclosed effect
Muon Split + MLA-256	Splits the attention computation path; multi-level attention with 256-dimension feature optimisation	Materially lower KV-cache footprint at performance parity with GQA-8; more stable training
Dynamic sparse attention	Dynamically selects the critical attention heads and positions for long sequences	Deployment cost cut to 50% of the prior level, stated as lossless
Prefill-Decode (PD) separation	Separates the prefill and decode stages of inference	Higher GPU utilisation
Slime framework	Asynchronous RL; fully decouples generation from training	Removes GPU idling on long-horizon agent tasks
Heartbeat fault tolerance	Periodic signals detect and recover training failures	Training robustness at scale
Lightning Indexer	Custom fused kernel optimising memory access	Higher inference efficiency specifically on domestic chips
FlashComm	Communication optimisation across distributed nodes	Lower training and inference communication latency
IndexShare (GLM-5.2)	Long-context efficiency	2.9x lower per-token FLOPs at 1mn context
MTP + KVShare speculative decoding	Faster token generation	c.20% higher acceptance length

Source: Company, GLM-5.2 model documentation, Daiwa

This is commercially important because inference, rather than training, was the main compute bottleneck as of July 2026 for the business to scale. Most incremental compute was being allocated to inference, and the company relied primarily on rented NVIDIA resources while increasing its use of domestic chips. Per our channel checks, Z.AI does not currently purchase GPUs directly, limiting capex but leaving the business exposed to computing service availability and pricing.


Domestic-chip adaptation converts supply risk into a supply option. Z.AI has expanded from basic model porting towards software-hardware co-design, including customised kernels and communication paths for domestic accelerators.

In July 2026, management claimed its planned training and inference capacity for the year had reached the **1GW level**, supported by a diversified accelerator pool and the acquisition of XCore Sigma (Zhongke Jiahe) to strengthen middleware and compiler capabilities.

The company also has the option to deploy a dedicated cluster of c.50,000 domestic AI accelerators for inference and large-scale optimisation. This could broaden supply and provide greater ARR headroom in 2H26, particularly while access to advanced overseas chips remains constrained.

That said, management did not confirm that a completed, wholly domestic 1GW-level data centre was operational. The deployment timing, installed capacity, utilisation and economics remain undisclosed.

US hyperscalers: data centre capacity availability (GW)



Source: Dolphin Research, Daiwa

However, we consider the 1GW-level compute capacity to be a milestone capacity expansion, as even US peers with easy access to premium chips like SpaceX only secured **0.7GW** available compute power in 2025.

However, in our view, the economic benefit remains dependent on chip availability, utilisation and actual inference efficiency, and management's claimed performance parity with overseas hardware has not been independently verified.

No.3: compact architecture strengthens serving economics, though task-level costs remain the test

MoE sparsity was used as an indicator of model economics, but the activation ratio should be read together with absolute activated parameters, in our view, which suggests that, among Chinese flagship LLMs, GLM achieves frontier capability with a compact active compute.

Activated vs. total parameters (MoE sparsity), by provider

Provider	Model	Released	Total params (B)	Activated params (B)	Activation ratio
DeepSeek	DeepSeek-V4-Pro	Apr 2026	1,600.0	49.0	3.1%
DeepSeek	DeepSeek-V4-Flash	Apr 2026	284.0	13.0	4.6%
Kimi (Moonshot)	Kimi K3	Jul 2026	2,800.0	104.0	3.7%
Kimi (Moonshot)	Kimi K2.6	Apr 2026	1,000.0	32.0	3.2%
GLM (Z.ai / Z.AI)	GLM-5.2	Jun 2026	744.0	40.0	5.4%
GLM (Z.ai / Z.AI)	GLM-5.1	Apr 2026	744.0	40.0	5.4%
GLM (Z.ai / Z.AI)	GLM-5	Feb 2026	744.0	40.0	5.4%
MiniMax	MiniMax-M3	Jun 2026	428.0	22.5	5.3%
MiMo (Xiaomi)	MiMo-V2.5-Pro	Apr 2026	1,020.0	42.0	4.1%
MiMo (Xiaomi)	MiMo-V2-Pro	Mar 2026	1,000.0	42.0	4.2%
MiMo (Xiaomi)	MiMo-V2-Flash	Dec 2025	309.0	15.0	4.9%
Qwen (Alibaba)	Qwen3.8-Max-Preview	Aug 2026	2,400.0	95.0	4.0%
Qwen (Alibaba)	Qwen3.5-397B-A17B	Feb 2026	397.0	17.0	4.3%
Qwen (Alibaba)	Qwen3.6-35B-A3B	Apr 2026	35.0	3.0	8.6%
Hunyuan (Tencent)	Hunyuan Hy3 (Hunyuan 3.0)	Jul 2026	295.0	21.0	7.1%
Hunyuan (Tencent)	Hunyuan Hy3-preview	Apr 2026	295.0	21.0	7.1%

Source: Companies, Daiwa

Considering the activation ratio: GLM is the lean model, not the heavy one. GLM-5.2 activates 5.4% of its parameters per token, a high ratio among Chinese flagships, but the ratio flatters the peers rather than penalising Z.AI.

Z.AI has more room on price than MoE sparsity suggests

Serving cost tracks absolute activated parameters, and on that measure **GLM-5.2 is the cheapest of the three large flagships to run**: 40bn active, against 49bn for DeepSeek-V4-Pro and 104bn for Kimi K3. **GLM's ratio is high only because its denominator is small**. In our view, the correct reading is the opposite of the headline: Z.AI buys its capability with a compact model worked hard, not with more compute per token, and it **therefore has more room on price than the sparsity table implies, not less**. The architecture does not cap intelligence, and the small model size is a feature.

Z.AI has strong domestic chip adaptability

Sparsity sets a ceiling, and post-training determines how much of it is realized. Z.AI's coding and agent scores come from reinforcement learning across more than 10,000 live software-engineering environments, where the reward signal is self-verifying because code either compiles or does not. A 744bn-parameter model also fits in far less accelerator memory than a 1.6–2.8tn one, so a serving replica spans fewer chips and less inter-chip traffic, which matters most precisely where interconnect, not raw FLOPs, is the binding constraint, as on domestic silicon. **This also echoes why Z.AI has stronger domestic chip adaptability.**

The hedge is on the output side: reasoning models are billed per token generated, so if GLM sustains its performance by thinking longer using the **CoT technique**, cost per completed task could rise even as cost per token falls.

Other competitive advantages and current limitations

Key strengths

- **Research talent and iteration capability:** Z.AI's Tsinghua roots support access to foundation-model researchers, while more than 70% of employees work in R&D. The company has historically maintained frequent minor and major model updates.
- **Coding-data feedback loop:** Usage in real coding and enterprise tasks generates clearer success and failure signals, supporting model post-training and further product improvement.
- **Enterprise customer foundation:** Z.AI has accumulated experience in localized deployment, data-security requirements and customized enterprise workflows. Nine of China's ten largest internet companies had integrated GLM as of March 2026.
- **Stronger funding capacity:** The July placement provides additional resources for compute, model development and commercialisation, reducing the immediate financing constraint in a capital-intensive industry.
- **Independent positioning:** Z.AI can work across internet platforms, enterprises and infrastructure providers without being tied to a single downstream ecosystem.

Current limitations

- **Low API switching costs:** Customers using standardized APIs can test and reroute workloads across different models, making demand sensitive to changes in relative performance and pricing. We believe product integration, proprietary workflows and data lock-in as important conditions for more durable API revenue.
- **Cloud margins remain immature:** Open Platform and API GPM was 18.9% in 2025, below the margins reported by established overseas model providers in the broker references.
- **Compute availability constrains monetisation:** Inference capacity is the main operational bottleneck, implying that reported demand cannot be fully converted into revenue without additional resources.
- **Benchmark leadership can be temporary:** Fast iteration of model leaderships demonstrates that a competitor can rapidly overtake Z.AI on a prominent coding benchmark.
- **The open-source model supports adoption but limits direct monetisation:** Third-party clouds can deploy GLM without necessarily routing usage through Z.AI's own API, meaning model adoption may exceed reported commercial revenue.

Financial outlook

Revenue by business: API segment expansion drives scale

Z.AI is entering a financial transition as its growth engine shifts from project-based enterprise deployment towards usage-driven MaaS services. We expect Open Platform and API revenue to drive a sharp acceleration in group revenue from 2026E, supported by rising enterprise token consumption and the progressive conversion of cloud ARR into recognized revenue. However, in our view, rapid API growth will not translate immediately into profitability, as the company must continue investing heavily in inference capacity, model training and R&D. The key financial question is therefore whether cloud scale, pricing and improving inference efficiency can ultimately offset the substantial compute and research costs required to maintain model competitiveness.

API expansion to drive up revenue scale over our forecast horizon

Open Platform and API segment to drive a jump in revenue scale


Z.AI's revenue increased by 132% YoY to CNY724m in 2025, extending the strong growth recorded since 2022, although the historical revenue base continued to reflect the company's project-led enterprise deployment model. We forecast revenue to rise to CNY5.0bn in 2026E and CNY14.0bn in 2027E, primarily driven by Open Platform and API services as rapid MaaS ARR growth progressively converts into reported revenue.

Group revenue is forecast to reach CNY47.5bn by 2029E, representing a **185% CAGR over 2025-29E**. YoY growth is expected to peak at 594% in 2026E before moderating in following years with limited visibility.

The forecast group-revenue inflection is driven almost entirely by Open Platform and API segment revenue, which increased by 293% YoY to CNY190mn in 2025 and is forecast to reach CNY4.0bn in 2026E. Accordingly, the segment's contribution to total revenue is expected to rise sharply **from 26% in 2025 to 79% in 2026E**, marking Z.AI's transition from project-based enterprise deployment towards a predominantly usage-driven MaaS model, thanks to booming token usage and model price in 2026E with GLM 5 & GLM 5.2 standout in global AI LLM score war.


We forecast Open Platform and API revenue to reach CNY44.5bn by 2029E, representing a **291% CAGR over 2025-29E** and 94% of group revenue. Annual growth rate is peaked at 1,997% in 2026E with low base, with the forecast primarily dependent on continued enterprise token consumption volume along with price hike. Going forward, strong volume growth is expected to continue but with speed moderating due to scaling effect. The growth rate will also be restricted by the company's compute capacity and the ability to maintain model leadership. While price trend is hard to forecast under the current situation, given fast evolving technologies, products and competitive landscape of AI LLM sector.

Z.AI: revenue and YoY growth rate



Source: Company, Daiwa estimates and forecasts

Z.AI: Open Platform and API revenue and YoY growth rate



Source: Company, Daiwa estimates and forecasts


Enterprise agents to lead growth among the remaining segments

We forecast Enterprise-level agent revenue to increase from CNY166m in 2025 to CNY1.7bn in 2029, representing a **78% CAGR**. We expect growth to moderate from 160% in 2026E to 40% in 2029E as the segment scales. Agents will become the fastest growing non-API business, in our view, supported by enterprise demand for model capabilities integrated with proprietary data, tools and workflows, although customised implementation constrains scalability relative to standardised APIs.

We forecast enterprise-level foundation-model revenue to rise from CNY366m in 2025 to CNY1.3bn in 2029, implying a **38% CAGR**. We expect growth to slow YoY from 50% in 2026E as Z.AI prioritises scalable MaaS growth over labour-intensive private deployments, especially enterprise foundation model, under which Z.AI has limited bargaining power and differentiation. The fast iteration of open-source domestic LLM is also making the model less stable.


We forecast technical services and other revenue to increase from CNY3m in 2025 to CNY11m in 2029E, representing a **44% CAGR**, but expect it to remain negligible relative to group revenue throughout the forecast period.

Z.AI: revenue composition by business



Source: Company, Daiwa forecasts

Non-cloud-based segments: revenue and YoY growth rate



Source: Company, Daiwa forecasts

Revenue and profitability by deployment method


Open Platform and API is equivalent to cloud-based deployment

Revenue can be decomposed by deployment model, with Open Platform and API representing cloud-based deployment and the remaining businesses classified as on-premise deployment, as explained under the business section. We forecast cloud-based revenue to increase from CNY190m in 2025 to CNY44.5bn in 2029E, representing a **291% CAGR**, compared with a **54% CAGR** for on-premise revenue from CNY534m to CNY3.0bn.

Open Platform and API segment = cloud-based deployment


The widening growth gap reflects Z.AI's strategic preference for cloud delivery. In our view, on-premise projects require customer-specific customisation, implementation and respective delivery personnel, which limits scaling, while the payable customer base, especially in China, is concentrated among large domestic enterprises and public-sector clients with stronger bargaining power. Cloud-based coding and agent workloads also provide more direct usage feedback from high-value tasks, supporting model iteration.

Z.AI: revenue composition by deployment



Source: Company, Daiwa forecasts

Z.AI: company and segment GPM



Source: Company, Daiwa forecasts

Cloud mix initially dilutes group GPM before scale supports recovery


The company's GPM declined from 65% in 2023 to 41% in 2025 and we expect this to bottom at 31% in 2026E, mainly because lower-margin cloud-based deployment has risen from 26% to 82% of revenue over that period. We forecast Cloud GPM to improve from 19% in 2025 to 39% in 2029E, supported by model pricing, higher inference utilisation and lower unit serving costs as API volume scales. We expect on-premise GPM to remain broadly stable at 48-51%, reflecting its higher delivery and customisation requirements. As cloud unit economics improve, we forecast group GPM to expand to 40% by 2029E.

We forecast R&D expense ratio to reach 35% in 2029E vs. 439% in 2025

R&D-led opex ratios to normalise with scaling despite continued investment


R&D remained Z.AI's largest operating expense at 439% of revenue in 2025, reflecting substantial model-training, compute, and technical-talent investment relative to its small revenue base. We forecast the R&D ratio to decline to 127% in 2026E and 35% in 2029E as revenue scales, although absolute R&D spending will likely continue to increase. We also forecast G&A and S&M ratios to fall from 70% and 54% in 2025 to 4% and 5% in 2029E, respectively, as corporate infrastructure, financing admin costs and customer-acquisition costs are expected to grow much slower than revenue. In our view, the projected reduction in operating expenses is relatively aggressive and requires tight cost control alongside rapid revenue delivery.

Z.AI: operating expense ratio trends



Source: Company, Daiwa forecasts

Z.AI: adjusted net (loss)/profit and adjusted NPM trends



Source: Company, Daiwa forecasts

Adjusted net losses set to break even in 2029E

We forecast adjusted net loss to widen from CNY3.2bn in 2025 to CNY6.3bn in 2026E and peak at CNY11.0bn in 2027E, **as continued R&D and computing investment proportional to much faster revenue growth outweighs the initial increase in gross profit**. Nevertheless, the gap will likely gradually narrow. We forecast Z.AI's adjusted NPM to improve from negative 439% in 2025 to negative 79% in 2027E and negative 31% in 2028E, reflecting increasing operating leverage with scaling. We forecast adjusted net

profit of CNY552m and a 1% adjusted NPM in 2029E, supported by improving cloud GPM and operating leverage. The timing of break-even remains sensitive to API pricing, inference efficiency and the level of investment required to maintain model competitiveness.


Leasing-led compute expansion keeps capex relatively contained

We forecast compute cost to account for 46% of revenue in 2029E vs. 321% in 2025

We expect Z.AI to rely mainly on leased computing capacity, which reduces upfront capital requirements, provides faster access to scarce chips and transfers part of the infrastructure-management and obsolescence risk to service providers. We forecast capex to rise from CNY11m in 2025 to CNY7.1bn in 2029E, but to remain at 15% of revenue.

Computing service fees, including training, inference, and hosting costs recorded under COGS and R&D, to increase from CNY2.3bn in 2025 to CNY21.7bn in 2029E, on our forecast. We expect their share of revenue to decline from 321% to 46%, respectively, as API revenue scales and compute utilisation improves. Therefore, lower computing-service cost intensity, rather than limited capex alone, is the key assumption supporting Z.AI's margin recovery.

Z.AI: capex and computing service fee and their respective % of revenue



Source: Company, Daiwa forecasts

Note: *Computing service fee refers to the total model training, inference, fine-tuning and hosting costs under COGS and R&D

US CSPs and LLM companies: Capex (USD bn) and capex as % of revenue

Company	Fiscal basis	Metric	2023	2024	2025	2026 YTD	2026 capex guidance
Microsoft	FY ends 30 June	Capex	28	45	65	116	175
		Capex % of revenue	13%	18%	23%	35%	
Alphabet	Calendar year	Capex	32	53	91	81	195-205
		Capex % of revenue	11%	15%	23%	35%	
Amazon	Calendar year	Capex	53	83	132	98	220
		Capex % of revenue	9%	13%	18%	26%	
Meta	Calendar year	Capex	28	39	72	51	130-145
		Capex % of revenue	21%	24%	36%	43%	

Source: Companies, Daiwa

We note from US CSPs and LLM peers that their capex ratio hit the c.35-40% level YTD, therefore we consider our assumptions of significant computing resource cost to be fair.

Ever stronger financing access mitigates industry-wide cash burn concerns

Financing remains an important consideration given the industry's high cash intensity and Z.AI's expected investment cycle. However, we believe the company's funding position is materially stronger than that of most independent foundation-model peers.

First, Z.AI has demonstrated exceptional fundraising capability. The company raised more than CNY35bn in 2026 alone, with the latest strategic financing significantly exceeding the size of its IPO proceeds and completed at a substantially higher valuation. We view this as

evidence of strong investor confidence in the company's technology positioning, commercialisation prospects, and strategic relevance.

**Financing risk mitigated
after the valuation
re-rating after following
its listing**

Second, we believe Z.AI is among the best-positioned beneficiaries of policy and ecosystem support. AI has become a national strategic priority, while Z.AI remains one of China's leading domestic foundation-model developers. With approximately 24% ownership by government- and university-backed investors, together with its stated commitment to domestic-chip adoption, the company appears well aligned with the broader China policy objectives and therefore is well placed to receive continued support from both financing institutions and strategic & financial investors.

Third, while our forecasts imply continued cash burn, the scale of historical operating cash outflows and capex before 2025 has remained substantially below recent fundraising levels since listing. In our view, the capital raised in 2026 should provide meaningful runway for continued model development, infrastructure investment, and commercialisation efforts. As a result, we believe the key question is not access to capital, but rather the pace at which Z.AI converts its technology leadership into sustainable monetisation.

Overall, we acknowledge financing risk as an inherent feature of the foundation-model industry. However, given Z.AI's technology leadership, strategic importance, shareholder base, and demonstrated fundraising capability, we believe it is better positioned than most independent LLM peers to secure the capital required to sustain long-term development.

Valuation: discounted EV/ARR could better reflect Z.AI's rapidly evolving business

We value Z.AI using a discounted EV/ARR methodology. ARR provides a more timely measure of commercial momentum than reported revenue, as the company's MaaS competence, token consumption, and pricing are changing rapidly. It also offers better comparability with Z.AI's closest foundation-model peers, most of which remain privately held and are commonly benchmarked against their latest disclosed or estimated ARR.

Frontier AI companies: LLM EV/ARR valuation comparison table

Company	Valuation date	Valuation (USD bn)	Nearest ARR date	ARR (USD bn)	adj ARR (USD bn)	Implied EV/ARR (x)
OpenAI	Oct-2024	157	FY2024	3.7	n.a.	42.4x
OpenAI	Mar-2025	300	Dec-2024	6.0	n.a.	50.0x
OpenAI	Mar-2026	852	Feb-2026	25.0	n.a.	34.1x
OpenAI	Jun-2026	1,000	Jun-2026	40.0	n.a.	25.0x
Anthropic	Mar-2025	62	Dec-2024	1.0	1	73.8x
Anthropic	Sep-2025	183	Aug-2025	5.0	4	43.9x
Anthropic	Feb-2026	380	Feb-2026	14.0	12	32.6x
Anthropic	May-2026	965	May-2026	47.0	39	24.6x
Anthropic	Jun-2026	965	Jul-2026	74.0	62	15.6x
Moonshot	May-2026	20	Mar-2026	0.1	n.a.	200.0x
Moonshot	Jul-2026	35	Jun-2026	0.3	n.a.	116.7x
Moonshot*	Aug-2026	50	Jun-2026	0.3	n.a.	166.7x
Moonshot	Aug-2026	50	Jul-2026	0.8	n.a.	62.5x
DeepSeek	Jun-2026	50	mid-2025	0.2	n.a.	227.3x
DeepSeek	Jul-2026	71	Jul-2026	0.5	n.a.	157.8x
Average of the latest						65.2x

Source: Reuters, CNBC, Sacra, Bybit, The information, TickerTrends, Trace Cohen, Daiwa estimates.

Note: We apply a 20% discount to Anthropic's ARR to enhance comparability with OpenAI, as Anthropic reports ARR gross of partner commissions, whereas OpenAI recognises revenue net of Microsoft channel commissions. * = Daiwa estimates.

We apply a **15x EV/ARR multiple to our 2030E average ARR** forecast of CNY82.5bn, using average rather than year-end ARR as it better reflects the revenue base throughout the year and reduces volatility. Our target multiple represents a substantial discount to the **65x average EV/ARR** which its peers currently trade at on street estimates, due to our following views: 1) disclosure quality could significantly impact valuation fairness, so we should treat comparable valuations as a reference, not a starting point. For example, estimates on DeepSeek's ARR could be more distorted than that on Anthropic or OpenAI, given differences in scale, funding structure, and philosophy of operation, while the latest Moonshot observations could be distorted by a relatively small ARR base facing significant appreciation for the new model capability; 2) the valuation reflects differences in model positioning and growth stage and visibility, and we believe Z.AI's positioning in China is better benchmarked to Anthropic in the US, which is only valued at **16x** for the latest implied EV/ARR; 3) we also noticed a trend of declining EV/ARR multiple on same single LLM company as their valuation and ARR getting higher, and we believe Z.AI is one of the leaders in China LLM monetisation scale; and 4) we are being deliberately conservative in factoring in fast changing sector and company metrics.

Z.AI: valuation

	2030E
Sales (CNYm)	78,061
Average ARR (CNYm)	82,504
Target EV/ARR Multiple	15x
Total Enterprise Value (CNYm)	989,306
Net cash (CNYm)	7,427
Total Equity Value (CNYm)	996,733
Shares outstanding (m)	494
Fair value in 2029E (CNY)	2,017
Cost of Equity	14%
-Equity market premium (high growth stock)	7.0%
-Risk free rate (US 10-year bond yield)	4.7%
-Beta (Daiwa assumption)	1.400
Years to discount	3.5
12M Target Price (HKD)	1,500

Source: Daiwa forecasts

We apply a 15x EV/ARR multiple to our 2030E average ARR forecast

The methodology derives a 2030E enterprise value of CNY989bn. After adding forecast net cash of CNY7.4bn, we estimate an equity value of CNY996bn, equivalent to CNY2,017 per share. Discounting this value by 3.5 years at a 14% cost of equity results, we derive our **12-month target price of HKD1,500**. We initiate coverage with a Buy (1) rating.

Upside catalysts and key risks

74% of Z.AI's H-shares are set to be unlocked in January 2027

Upside catalysts

- **Sustained frontier model leadership:** Further releases that remain competitive in coding and long-horizon agent tasks could support usage and pricing.
- **Faster inference-capacity expansion:** Additional compute could allow Z.AI to monetise demand currently constrained by supply.
- **Improving API GPM:** Higher prices and lower unit inference costs would strengthen the relationship between ARR growth and gross profit.
- **Deeper enterprise-agent adoption:** Integration with customer data, tools and workflows could increase retention and reduce API substitution.
- **International or sovereign-AI contracts:** Overseas adoption could broaden the addressable market, although disclosed commercial evidence remains limited.
- **Monetisation beyond Z.AI-hosted inference:** Commercial licensing or revenue-sharing arrangements for third-party GLM deployments could create less compute-intensive revenue, but this remains a potential rather than an established model.

Key downside risks

- **Rapid model competition:** Leadership can narrow quickly following launches from Moonshot, MiniMax, DeepSeek, and major internet platforms.
- **Pricing pressure:** API customers can switch providers, and renewed price competition would delay margin improvement.
- **Compute constraints and cost inflation:** Limited access or higher rental prices could restrict revenue and compress margins.
- **ARR deceleration or weak revenue conversion:** Slower usage growth would challenge the assumptions embedded in the company's valuation.
- **Sustained cash burn:** Z.AI requires continued R&D and compute investment before achieving profitability.
- **Talent retention:** The loss of senior researchers could affect model architecture, release cadence and engineering execution.
- **Geopolitical and regulatory restrictions:** Z.AI's inclusion on the US BIS Entity List, alongside potential future export-control measures affecting advanced chips, cloud services or overseas operations, could increase costs, constrain model development and limit international expansion.
- **Share-supply pressure:** The January 2027 expiry of lock-up restrictions on founders, employee platforms, and selected pre-IPO investors of c.74% of its H-shares, although expiry does not necessarily imply disposal.

ESG analysis

ESG opportunities

Opportunities	Management	Analyst comments
S AI-enabled productivity	2	Zhipu's AI solutions help institutional and individual users improve productivity, optimise resource allocation and reduce operating costs. Its programming, agent and enterprise-model services covered 218 countries and regions and more than 4mn SMEs and developers in 2025.
S Computing efficiency	2	Zhipu is improving model-training and inference efficiency through algorithm and engineering optimisation and is exploring green-computing approaches. However, the company does not disclose quantified energy-efficiency or emissions-reduction targets.

ESG risks

Risks	Management	Analyst comments
G Executive/board quality	2	The nine-member Board comprises three executive, three non-executive and three independent non-executive directors, with the chairman and CEO roles separated. The Board includes two female directors and has experience spanning AI research, management, investment and accounting.
G Capital management	3	Zhipu remained loss-making and had no distributable reserves at end-2025; no dividend was declared or paid. IPO proceeds are primarily allocated to foundation-model R&D and MaaS infrastructure, reflecting the company's ongoing investment stage.
G Related party & transaction	2	No 2025 connected transaction required disclosure under Chapter 14A, and no disclosed related-party transaction required shareholder approval, annual review or connected-transaction disclosure.
S Data security	2	Zhipu applies vendor and client security reviews, minimum-access controls, audit trails, data desensitisation and cybersecurity incident-response drills. The company reported no material data-security non-compliance during the track-record period.
G Systemic risk management	2	The Board oversees risk policy, senior management implements controls, and dedicated risk-management and internal-audit departments coordinate risk identification, assessment and response. The framework covers market, credit, operational, compliance and strategic risks, although it can provide reasonable rather than absolute assurance.
S Employee engagement, diversity & inclusion	2	Zhipu has equal-employment and non-discrimination policies and uses employee ownership platforms to align core technical and management personnel. Female employees represented approximately 32% of the workforce at end-2025.
S Customer privacy	2	Zhipu requires customer data to have legitimate sources and appropriate consent, restricts processing to authorised purposes, conducts privacy reviews before product launches and material updates, and obtains regular external legal reviews of privacy practices.
Company specific	Management	Analyst comments
G Responsible AI and model safety	2	Foundation-model reliability and responsible deployment are material company-specific risks because model failures, hallucinations or prolonged service disruption could affect customers and regulatory compliance. Zhipu is developing hierarchical memory and self-evaluation mechanisms intended to reduce hallucination and improve logical verification, though company provides limited quantified disclosure on model-safety testing, harmful-output rates or external assurance.

Note: Management score represents a company's ability to manage/benefit from certain ESG topics. The scores range from 1 to 3, with 1 being the strongest.

Update Date: 11 Aug 2026

Source: Daiwa, Company

Minimax Group (100 HK)

 Target price: **HKD530.00**

 Share price (11 Aug): **HKD328.00** | Up/downside: **+61.6%**

5 4 3 2 1


Buy
 (initiation)

Initiation: fuelling China's AI commercialisation boom

- Dual strengths in capability and efficiency drive monetisation potential
- Multi-modality strategy broadens commercial applicability
- Initiating with a Buy (1) rating and 12-month TP of HKD530

John Choi

(852) 2773 8730

john.choi@hk.daiwacm.com



Investment case: We initiate coverage of MiniMax, a leading foundation model developer in China, with a Buy (1) rating. Through its dual focus on competitive model capabilities and cost efficiency, MiniMax is one of the few listed pure-play beneficiaries of China's AI commercialisation trend.

Competitive capability, fractional cost. MiniMax has steadily narrowed the capability gap with leading Western model developers while maintaining a strong cost-efficiency advantage. Despite significantly lower R&D spend and fewer resources, its flagship M3 model currently ranks competitively on global benchmarks, underscoring the company's strength in model development and execution. We believe MiniMax's differentiated performance-to-cost proposition positions it well to capture growing AI demand, particularly among cost-conscious enterprises and developers.


Multi-modality driving the B2B-B2C flywheel. We believe MiniMax's full-modality strategy broadens its applicability across AI agents and next-gen AI applications. Its integrated B2C and B2B ecosystem creates a commercialisation flywheel, where consumer products drive model iteration and monetisation, while APIs/MaaS provide a scalable pathway for enterprise deployment and recurring revenue growth. We forecast total revenue to rise at a 180% CAGR over 2025-29E, with revenue from Open Platform and AI-based enterprise services (232% 2025-29E CAGR) outpacing that of AI-native products (138% 2025-29E CAGR). We also expect continuous margin improvement over 2026-28E, supported by operating leverage as commercialisation pace improves.

Catalysts: 1) Positive usage traction and ranking performance of video generation model H3; 2) successful launch of large-scale text model M3 Pro with a narrowed capability gap with frontier players; 3) new AI-native products gaining user and monetisation traction; 4) a step-up in annual recurring revenue (ARR); and 5) easing chip supply constraints in China.

Valuation: We initiate coverage of MiniMax with a Buy (1) rating and 12M TP of HKD530, derived by applying a 2029E EV/ARR of 5x, and discounting back 2.5 years at a discount rate of 15%. While its share price has fallen over 70% from its peak of HKD1,238 in March amid intensifying competition, we believe the sell-off has been excessive relative to the company's fundamentals and growth prospects. In our view, the recent weakness presents an attractive entry point, with upside catalysts that should support a valuation rerating in the next 12 months.

Risks: Intensifying competition among AI model developers and slower-than-expected execution of R&D investments.

Share price performance



12-month range	193.10-1,238.00
Market cap (USDbn)	13.11
3m avg daily turnover (USDm)	242.02
Shares outstanding (m)	314
Major shareholder	Alisoft China Holding (14.8%)

Financial summary (USD)

Year to 31 Dec	26E	27E	28E
Revenue (m)	377	1,115	2,472
Operating profit (m)	(499)	(598)	(410)
Net profit (m)	(335)	(300)	(27)
Core EPS (fully-diluted)	(1.067)	(0.955)	(0.087)
EPS change (%)	n.a.	n.a.	n.a.
Daiwa vs Cons. EPS (%)	n.m.	n.m.	n.m.
PER (x)	n.a.	n.a.	n.a.
Dividend yield (%)	0.0	0.0	0.0
DPS	0.000	0.000	0.000
PBR (x)	n.a.	n.a.	n.a.
EV/EBITDA (x)	n.a.	n.a.	n.a.
ROE (%)	n.a.	n.a.	n.a.

Source: FactSet, Daiwa forecasts

5 4 3 2 1

How do we justify our view?

Buy
(initiation)

Growth outlook

✓ ✓ ✓ ✓ ✓

Valuation

✓ ✓ ✓ ✓ ✓

Earnings revisions


✓ ✓ ✓ ✓ ✓

Growth outlook

✓ ✓ ✓ ✓ ✓

We believe MiniMax is poised for strong growth over 2026-29E as global AI adoption scales. We forecast MiniMax's total revenue to grow at a 180% CAGR during 2025-29E, with revenue from Open Platform and AI-related services (232% 2025-29E CAGR) outpacing that of AI-native products (138% 2025-29E CAGR). In our view, B2B revenue is likely to see exponential growth as companies race to integrate AI into everyday workflows in a competitive efficiency race, supporting a scalable recurring revenue stream. Meanwhile, we expect consumer-facing products to maintain robust growth on rising user adoption. We also expect margin improvement driven by operating leverage as commercialisation pace improves.

MiniMax: revenue and revenue growth (USDm)



Source: Company, Daiwa forecasts

Valuation

✓ ✓ ✓ ✓ ✓

Our 12M TP of HKD530 is based on a 2029E EV/ARR multiple of 5x, and discounting back 2.5 years at a discount rate of 15%. Our 5x target multiple represents a substantial discount to its peers, including 25x based on last round EV/ARR on street estimates for OpenAI, 16x for Anthropic and 63x for MoonShot, as implied by reported transactions. As compared with leading players that focus on state-of-the-art intelligence, we believe MiniMax's strategic focus on commercial applicability warrants a relatively conservative valuation. While Minimax's stock price has fallen over 70% from its peak amid rising competition, we believe the sell-off has been excessive. We see the recent weakness creating an attractive entry point ahead of multiple upside catalysts that could support a rerating.

MiniMax: valuation summary

	2029E
Sales (USDm)	4,861
Average ARR (USDm)	5,715
Target EV/ARR multiple (x)	5x
Enterprise value (USDm)	28,575
Net debt (USDm)	-1,795
Equity value (USDm)	30,370
Shares outstanding (m)	314
Discount rate (%)	15%
Years to discount	2.5
Target price in 12M (HKD)	530


Source: Daiwa estimates

Earnings revisions

✓ ✓ ✓ ✓ ✓

The Bloomberg consensus has steadily raised MiniMax's 2026E revenue forecasts since its listing, driven by rapid AI market growth, rising global AI adoption, and the company's compelling performance-to-cost proposition, which positions it to benefit from growing enterprise AI spending. Meanwhile, 2026E EPS forecasts have been trending down amid concerns over rising competition and ongoing investment requirements. We see upside to current 2026E revenue and EPS expectations, as the benefits of the recent H3 launch and upcoming M3 Pro release are, in our view, not yet fully reflected in market forecasts. We believe both models have the potential to drive a meaningful step-up in ARR through stronger user adoption, engagement, and monetisation.

MiniMax: 2026 revenue and EPS forecast revisions



Source: Bloomberg

Financial summary

Key assumptions

Year to 31 Dec	2021	2022	2023	2024	2025	2026E	2027E	2028E
Total ARR (YoY %)	-	n.a.	n.a.	n.a.	n.a.	10.0	1.0	0.9
B2B ARR (YoY %)	-	n.a.	n.a.	n.a.	n.a.	21.0	1.2	1.1
B2C ARR (YoY %)	-	n.a.	n.a.	n.a.	n.a.	5.3	0.6	0.5

Profit and loss (USDm)

Year to 31 Dec	2021	2022	2023	2024	2025	2026E	2027E	2028E
AI-native products	-	0	1	22	53	181	428	891
Open Platform and other AI-based enterprise services	-	0	3	9	26	196	687	1,580
Other Revenue	-	0	0	0	0	0	0	0
Total Revenue	-	0	3	31	79	377	1,115	2,472
Other income	-	1	10	38	43	77	114	146
COGS	-	0	(4)	(27)	(59)	(251)	(676)	(1,399)
SG&A	-	(4)	(30)	(101)	(89)	(223)	(312)	(371)
Other op. expenses	-	(12)	(80)	(227)	(296)	(479)	(839)	(1,259)
Operating profit	-	(14)	(101)	(287)	(321)	(499)	(598)	(410)
Net-interest inc./(exp.)	-	0	8	20	9	33	44	31
Assoc/forex/extraord./others	-	(59)	(176)	(199)	(1,559)	31	31	31
Pre-tax profit	-	(74)	(269)	(465)	(1,872)	(435)	(523)	(349)
Tax	-	0	0	0	0	0	0	0
Min. int./pref. div./others	-	0	0	0	0	0	0	0
Net profit (reported)	-	(74)	(269)	(465)	(1,872)	(435)	(523)	(349)
Net profit (adjusted)	-	(12)	(89)	(244)	(251)	(335)	(300)	(27)
EPS (reported)(USD)	-	(0.737)	(2.556)	(4.282)	(17.226)	(1.388)	(1.666)	(1.112)
EPS (adjusted)(USD)	-	(0.121)	(0.846)	(2.248)	(2.309)	(1.067)	(0.955)	(0.087)
EPS (adjusted fully-diluted)(USD)	-	(0.121)	(0.846)	(2.248)	(2.309)	(1.067)	(0.955)	(0.087)
DPS (USD)	-	0.000	0.000	0.000	0.000	0.000	0.000	0.000
EBIT	-	(13)	(98)	(280)	(297)	(398)	(375)	(89)
EBITDA	-	(13)	(97)	(278)	(295)	(385)	(337)	(5)

Cash flow (USDm)

Year to 31 Dec	2021	2022	2023	2024	2025	2026E	2027E	2028E
Profit before tax	-	(74)	(269)	(465)	(1,872)	(435)	(523)	(349)
Depreciation and amortisation	-	0	1	2	3	13	38	84
Tax paid	-	0	0	0	0	0	0	0
Change in working capital	-	2	26	14	6	22	(28)	(101)
Other operational CF items	-	61	178	191	1,583	102	224	322
Cash flow from operations	-	(11)	(64)	(258)	(280)	(299)	(288)	(44)
Capex	-	(0)	(1)	(1)	(1)	(9)	(28)	(62)
Net (acquisitions)/disposals	-	(35)	(40)	(702)	47	0	0	0
Other investing CF items	-	0	0	271	26	(9)	(27)	(60)
Cash flow from investing	-	(35)	(40)	(431)	72	(19)	(55)	(122)
Change in debt	-	0	0	33	1	830	(830)	0
Net share issues/(repurchases)	-	50	307	740	426	1,839	0	0
Dividends paid	-	0	0	0	0	0	0	0
Other financing CF items	-	(0)	(1)	(2)	(4)	(1)	(1)	(1)
Cash flow from financing	-	50	306	771	423	2,668	(831)	(1)
Forex effect/others	-	0	0	1	3	0	0	0
Change in cash	-	4	202	83	219	2,351	(1,174)	(166)
Free cash flow	-	(11)	(65)	(259)	(281)	(308)	(316)	(105)

Source: FactSet, Daiwa forecasts

Financial summary continued ...

Balance sheet (USDm)

Year to 31 Dec	2021	2022	2023	2024	2025	2026E	2027E	2028E
Cash & short-term investment	-	5	206	289	508	2,858	1,684	1,518
Inventory	-	0	0	0	0	0	0	0
Accounts receivable	-	0	1	7	11	68	201	445
Other current assets	-	69	112	510	489	551	703	983
Total current assets	-	73	320	806	1,007	3,477	2,588	2,946
Fixed assets	-	0	1	1	2	7	24	62
Goodwill & intangibles	-	0	3	3	2	2	2	2
Other non-current assets	-	0	0	101	77	80	87	101
Total assets	-	74	324	911	1,088	3,567	2,702	3,111
Short-term debt	-	0	0	19	35	35	35	35
Accounts payable	-	2	17	51	58	120	256	459
Other current liabilities	-	148	646	1,637	3,640	3,722	3,851	4,084
Total current liabilities	-	150	663	1,708	3,734	3,878	4,143	4,579
Long-term debt	-	0	0	0	0	0	0	0
Other non-current liabilities	-	0	3	2	3	833	3	3
Total liabilities	-	150	666	1,710	3,737	4,710	4,145	4,582
Share capital	-	0	0	0	0	1,839	1,839	1,839
Reserves/R.E./others	-	(76)	(342)	(799)	(2,648)	(2,983)	(3,283)	(3,310)
Shareholders' equity	-	(76)	(342)	(799)	(2,648)	(1,144)	(1,444)	(1,471)
Minority interests	-	0	0	0	0	0	0	0
Total equity & liabilities	-	74	324	911	1,088	3,567	2,702	3,111
EV	-	13,107	12,906	12,843	12,640	10,289	11,463	11,630
Net debt/(cash)	-	(5)	(206)	(269)	(472)	(2,823)	(1,649)	(1,482)
BVPS (USD)	-	n.a.	n.a.	n.a.	n.a.	n.a.	n.a.	n.a.

Key ratios (%)

Year to 31 Dec	2021	2022	2023	2024	2025	2026E	2027E	2028E
Sales (YoY)	-	n.a.	n.a.	782.2	158.9	377.0	195.8	121.6
EBITDA (YoY)	-	n.a.	n.a.	n.a.	n.a.	n.a.	n.a.	n.a.
Operating profit (YoY)	-	n.a.	n.a.	n.a.	n.a.	n.a.	n.a.	n.a.
Net profit (YoY)	-	n.a.	n.a.	n.a.	n.a.	n.a.	n.a.	n.a.
Core EPS (fully-diluted) (YoY)	-	n.a.	n.a.	n.a.	n.a.	n.a.	n.a.	n.a.
Gross-profit margin	-	n.a.	n.a.	12.2	25.4	33.4	39.4	43.4
EBITDA margin	-	n.a.	n.a.	n.a.	n.a.	n.a.	n.a.	n.a.
Operating-profit margin	-	n.a.	n.a.	n.a.	n.a.	n.a.	n.a.	n.a.
Net profit margin	-	n.a.	n.a.	n.a.	n.a.	n.a.	n.a.	n.a.
ROAE	-	n.a.	n.a.	n.a.	n.a.	n.a.	n.a.	n.a.
ROAA	-	n.a.	n.a.	n.a.	n.a.	n.a.	n.a.	n.a.
ROCE	-	n.a.	n.a.	n.a.	n.a.	n.a.	n.a.	n.a.
ROIC	-	17.7	32.2	35.5	15.3	14.1	16.9	13.6
Net debt to equity	-	6.1	60.3	33.7	17.8	246.8	114.2	100.8
Effective tax rate	-	0.0	0.0	0.0	0.0	0.0	0.0	0.0
Accounts receivable (days)	-	n.a.	70.6	49.7	40.9	38.0	44.0	47.7
Current ratio (x)	-	0.5	0.5	0.5	0.3	0.9	0.6	0.6
Net interest cover (x)	-	n.a.	n.a.	n.a.	n.a.	n.a.	n.a.	n.a.
Net dividend payout	-	n.a.	n.a.	n.a.	n.a.	n.a.	n.a.	n.a.
Free cash flow yield	-	n.a.	n.a.	n.a.	n.a.	n.a.	n.a.	n.a.

Source: FactSet, Daiwa forecasts

Company profile

Founded in 2021, MiniMax is a leading Chinese foundation model developer focused on full-modality AI and global operations. The company monetises its capabilities through AI-native products, its MaaS platform (Open Platform) and AI-related enterprise solutions.

A pure-play on China's AI commercialisation

MiniMax is a leading Chinese foundation model developer focused on full-modality AI and global operation

Founded in 2021, MiniMax is a leading Chinese foundation model developer focused on full-modality AI and global operation. Reflecting its view that human interaction is inherently multi-modal, the company has built capabilities across text, video and audio from the outset. Leveraging its foundation models, MiniMax has developed a portfolio of AI-native products that serves consumers, developers and enterprises. This has helped it create a flywheel between its model capability and commercialisation, where user adoption generates data and feedback that support continuous model iteration and product optimisation.

Competitive capability, fractional cost


MiniMax is closing the capability gap with global developers

Model capability is the core asset of foundation model developers, underpinning user growth, engagement, pricing power and monetisation potential. By emphasising infrastructure efficiency, MiniMax has established a robust track record of developing competitive foundation models while maintaining strong scalability, inference economics and capital efficiency, creating a solid foundation for long-term commercialisation.

The company's latest offerings have steadily narrowed the performance gap with leading Western players across global benchmarks. According to Artificial Analysis, MiniMax's latest flagship model M3, launched in June 2026, scored 44 out of 100 on the Intelligence Index, compared with 55 for OpenAI, 54 for Anthropic and 50 for Google at the time. This marked a meaningful improvement from a year earlier, where M1 scored 18 vs. 33 for Open AI, 31 for Anthropic and 26 for Google. Notably, MiniMax achieved these capability gains with a relatively lean model architecture, underscoring the company's efficiency in model design and training. M3 contains 428bn parameters, substantially below the trillion-parameter scale of leading frontier models such as GPT (OpenAI), Claude (Anthropic) and Gemini (Google).

In our view, MiniMax's ability to narrow the capability gap despite GPU supply constraints in China and lower financial resources than global peers underscore the strength of its underlying technology and R&D execution. As a result, the company has emerged as one of China's leading model developers while becoming increasingly competitive globally.

Artificial Analysis: frontier-language model intelligence, December 2022- August 2026




Source: Artificial Analysis (as of 5 August 2026)


Cost efficiency as a key competitive advantage

While MiniMax's model performance is closing the gap with that of Western peers, what stands out most to us is the company's ability to achieve this while maintaining strong cost efficiency. MiniMax had only incurred R&D spend of USD522m during FY22-25 (USD253m in FY25), compared with an estimated USD19bn spent by OpenAI in FY25 alone. In addition, MiniMax operates with a lean team of 415 employees as of FY25, significantly

below OpenAI's c.4,500 employees and Anthropic's c.2,500. We view this as evidence of the company's disciplined AI development approach and efficient resource allocation.

MiniMax: R&D expense comparison, 2025 (USDm)


Source: Company data, Daiwa estimates


MiniMax: number of employee comparison, 2025


Source: Company data, Daiwa estimates

Through its dual strengths in model capability and R&D efficiency, MiniMax offers one of the most compelling performance-to-cost propositions in the AI industry. According to Artificial Analysis, M3 currently ranks No. 15 globally on the Intelligence Index yet ranks No.6 in terms of the lowest cost on a cost-per-task basis compared with peers. In addition, M3 is one of the few highly ranked models that supports native multi-modal inputs (ie, text, images, and videos), whereas most competing models are primarily optimised for text, enabling a wider range of use cases through a single model architecture.


Many AI workloads require reliable, scalable performance rather than frontier intelligence

We believe MiniMax's strong performance-to-cost profile is becoming increasingly important as AI adoption scales, with customers prioritising ROI and affordability over marginal capability improvements. Many workloads, including general chats, customer service, and content generation, do not require frontier-level intelligence but rather reliable and scalable performance at attractive economics. By delivering competitive performance at lower inference costs while supporting a broader range of applications, MiniMax improves customer ROI and facilitates deployment, driving growth in user adoption, token consumption, and revenue.

Artificial Analysis's Intelligence Index – MiniMax's latest flagship model M3 ranked 15th


Source: Artificial Analysis (as of 5 August 2026)

Artificial Analysis Cost-per-Intelligence Index



Source: Artificial Analysis (as of 5 August 2026)

Note: Weighted average cost (USD) per Artificial Analysis Intelligence Index task.

This dynamic is particularly relevant for SMEs and developers, where budget constraints often outweigh the need for state-of-the-art capabilities. We thus believe MiniMax's unique "good-enough performance + attractive economics" proposition positions it well to capture cost-conscious demand, with pricing advantages supporting customer acquisition and usage growth.


Global reach unlocking scale and efficiency

In 9M25, MiniMax generated 73% of total revenue from markets outside Mainland China

MiniMax has a meaningful global footprint built through early overseas expansion. In 9M25, the company derived 73% of total revenue from markets outside Mainland China, with Singapore, the US, and other overseas markets contributing 24%, 20%, and 28% of total revenue, respectively.

We believe MiniMax's established international presence expands its TAM, presenting significant monetisation potential supported by higher customer purchasing power and AI pricing that remain materially higher than in China. In addition to capturing emerging market opportunities from a larger, fast-growing TAM, exposure to higher-value overseas customers also provide opportunities in lifting ARPU to enhance overall monetisation efficiency. Together, these factors strengthen MiniMax's ability to scale profitably as global AI adoption expands.

MiniMax: revenue breakdown by geographic market, 9M25



Source: Company

Multi-modality driving the B2B-B2C flywheel

Since inception, MiniMax has pursued a full-modality strategy, developing foundation models across text, video, speech, and music. This has created a unified foundation for both consumer- and enterprise-facing AI applications, strongly positioning MiniMax in capturing emerging market opportunities from AI applications and agentic workflows, as users increasingly demand richer and more versatile AI experiences.

MiniMax's full-modality strategy provides strong leverage to the emerging agent economy

We believe MiniMax's multi-modal capabilities are a key differentiator, as AI agents evolve from answering prompts to executing real-world tasks. Unlike single-modality models that can only address partial workflows, agentic applications increasingly require the ability to understand and generate multiple content formats within a single interaction. For example, a travel-planning agent may need to process a user's voice request (speech), search for flights and hotels (text), compare destination images (images) and reviews (text), and generate an itinerary (text). By natively supporting these capabilities, MiniMax can address broader use cases through a single platform. Combined with its industry-leading cost-performance proposition, we believe MiniMax lowers barriers to AI agent development and deployment, supporting wider adoption of the agent economy.

MiniMax: model portfolio

Modality	Latest model	Launch date	Develop timespan from previous version	Model description	Ranking
Text	M3	June 2026	3 months (from M2.7)	Reasoning model 428bn parameters (vs. 230bn for M2.7)	No. 15 on the Artificial Analysis Intelligence Index
Video	H3	July 2026	8 months (from Hailuo 2.3)	Multi-modal video generation model Open-sourced	No. 2 on the Artificial Analysis text-to-video leaderboard No. 3 on the Artificial Analysis image-to-video leaderboard No. 1 on the Artificial Analysis video editing leaderboard
Speech	Speech 2.8	February 2026	4 months (from Speech 2.6)	Multilingual voice generation	No. 10 on the Artificial Analysis text-to-speech leaderboard
Music	Music 3.0	July 2026	3 months (from Music 2.6)	Music generation	No. 7 on the Artificial Analysis Instrumental leaderboard (Music 2.6) No. 5 on the Artificial Analysis Vocals leaderboard (Music 2.5+)
Image	Image-01	February 2025	n/a	Image generation	No. 102 on the Artificial Analysis text-to-image leaderboard

Source: Company, Artificial Analysis (as of 5 August 2026)

The B2B and B2C dual business model

Leveraging its multi-modal model suite, MiniMax has built 2B (Open Platform and other AI-based enterprise services) and 2C (AI-native products) business models. On the consumer-facing end, MiniMax's product portfolio includes Talkie/Xingye, Hailuo AI, MiniMax and MiniMax Audio, spanning entertainment, content creation and productivity use cases. As of FY25, MiniMax cumulatively served over 236m users across over 200 countries and regions, and AI-native products accounted for 67% of total revenue in FY25.

On the 2B side, MiniMax mainly delivers services through its Open Platform, a MaaS offering that enables enterprises and developers to integrate MiniMax's capabilities into their own products. Monetisation is primarily usage-based, with customers charged based on token, video and audio consumption, alongside subscription plans for recurring access. This creates a scalable revenue model directly linked to customer activity while expanding MiniMax's ecosystem reach across applications and industry verticals. The company had cumulatively served over 214,000 enterprise customers and developers as of FY25, with Open platform and other AI-based enterprise services contributing 33% of revenue.

Consumer ecosystem powering the technology-as-product strategy

We view MiniMax's consumer ecosystem as a key pillar of its technology-as-product strategy, creating a flywheel between innovation and commercialisation. By integrating model advancements directly into consumer-facing applications, MiniMax can rapidly monetise new capabilities without developing business models. These applications also serve as testing grounds for new features, generating user feedback and data that support product optimisation and accelerate model iteration. More importantly, they help MiniMax

build product know-how needed to make AI accessible and intuitive for end users. This creates a self-reinforcing cycle where technological advances improve user experience, while real-world usage enhances products and model performance.

MiniMax: AI-native product portfolio

Product	Launch date	Product description	Monthly subscription price	Underlying models	Monetisation models	MAU (m)	Paying users (m)	Paying ratio (%)	Revenue (USDm)
Talkie/Xingye	Sept 2023	Multi-modal entertainment platform	Talkie: USD1.49-29.99 Xingye: CNY11-35	Text, video, speech, and music	- Freemium - Subscriptions - In-app purchases - Online marketing services	20.05	1.39	6.9%	18.8
HailuoAI	Aug 2024	Visual generation platform	USD9.99-199.99	Video	- Freemium - Subscriptions - In-app purchases	5.65	0.31	5.5%	17.5
MiniMax Audio	May 2025	Audio generation tool	USD5-99	Speech and music	- Freemium - Subscriptions - In-app purchases	0.49	0.06	12.1%	1.1
MiniMax	Mar 2025	Intelligent agent application	USD19-69	Text, video, speech, and music	- Freemium - Subscriptions - In-app purchases	1.43	0.01	0.7%	0.8
MiniMax Code	Jun 2026	Coding agent	USD20-120	Text, video, speech, and music	- Freemium - Subscriptions - In-app purchases	n/a	n/a	n/a	n/a
MiniMax Hub	Jun 2026	Desktop AI workstation	USD10.5-215	Text, video, speech, and music	- Freemium - Subscriptions - In-app purchases	n/a	n/a	n/a	n/a

Source: Company

Note: MAU, paying user, paying ratio and revenue are based on 9M25 data.

The proliferation of agentic workflows is driving API/MaaS demand as enterprises deploy AI into business processes

B2B business scaling the commercialisation flywheel

On the B2B end, we expect the proliferation of agentic workflows to drive demand for API/MaaS as enterprises increasingly embed AI into business processes. In our view, MiniMax's full-modality model portfolio and improving capabilities should enable the company to capture a growing share of enterprise AI spending. We expect B2B revenue growth to outpace B2C over time, as enterprise customers typically consume AI at a larger scale and on a recurring basis. Rising token consumption from AI applications and agents should translate into sustained API/MaaS demand, supporting a more scalable revenue model and stronger long-term growth visibility.


Financial analysis

We expect B2B revenue growth to outpace 2C as enterprises scale AI deployment

We forecast a 180% revenue CAGR over 2025-29E


We forecast MiniMax's total revenue to grow at a 180% CAGR during 2025-29E, with revenue from Open Platform and other AI-based enterprise services (232% 2025-29E CAGR) outpacing that of AI-native products (138% 2025-29E CAGR). In our view, B2B revenue will likely see significant growth as companies race to integrate AI into everyday workflows in a competitive efficiency race, supporting a scalable recurring revenue stream. Meanwhile, consumer products should also maintain robust growth on rising user adoption.

MiniMax: revenue for 2023-29E (USDm)



Source: Company data, Daiwa forecasts

MiniMax: revenue split for 2023-29E (%)




Source: Company data, Daiwa forecasts

AI-native products: 138% revenue CAGR over 2025-29E

We expect AI-native product revenue to deliver a 2025-29E CAGR of 138%, driven by rising user adoption and monetisation. We see upside to our forecasts as MiniMax continues to advance capabilities and expand product portfolio. For instance, the upcoming release of M3 Pro (3tn parameter, target for Sep-Oct) could drive a step-up in ARR as technological breakthroughs further stimulate market demand. Beyond model upgrades, we do not expect the company's product offerings or business models to remain static given the rapid pace of industry evolution. Instead, we anticipate MiniMax will continue introducing new products, such as MiniMax Code (coding agent) and MiniMax Hub (desktop AI workstation), while exploring additional monetisation channels including advertising and transaction-based commissions alongside its existing subscription and virtual-item businesses.

MiniMax: AI-native product revenue for 2023-29E (USDm)




Source: Company data, Daiwa forecasts

Open Platform: 232% revenue CAGR over 2025-29E

We view MiniMax's enterprise-facing business as its primary growth driver over the coming years, supported by rising enterprise AI spending, improving model capabilities and global expansion. Fast adoption of consumer AI applications has validated real-world use cases while accelerating product refinement and reducing deployment barriers. As a result, enterprises are becoming increasingly willing to embed AI into business workflows, backed by clearer ROI and a growing pool of proven applications, driving higher API/MaaS demand and token consumption. We see particularly strong opportunities in overseas markets, where higher AI spending and greater willingness to pay create a favourable operating environment. We expect the 2B segment to outpace 2C growth over time, as enterprise revenue is characterized by large-scale, recurring, and strong stickiness.

MiniMax: open platform and other AI-based enterprise service revenue for 2023-29E (USDm)




Source: Company data, Daiwa forecasts

GPM to trend-up on inference efficiency gains and scale

We forecast gross margin expansion over our forecast horizon driven by technology-led efficiency gains and economies of scale

MiniMax's COGS mainly comprises cloud services costs related to inference activities, accounting for 93% of total COGS in 9M25. The company has achieved meaningful efficiency gains over time, with cloud service costs as a % of revenue declining from 118% in 2023 to 85% in 2024 and further to 71% in 9M25, driven by ongoing optimisation of model architecture and infrastructure, higher compute utilisation efficiency, and scale benefits. These improvements have been the primary driver of MiniMax's GPM expansion, with GPM improving from -25% in FY23 to 25% in FY25. Looking ahead, we expect continued investment in R&D to drive further advancements in model and inference efficiency, while increasing scale should provide additional cost leverage. As a result, we forecast GPM to expand to 46% by 2029E.

MiniMax: gross margin, 2023-29E (%)




Source: Company data, Daiwa forecasts

Adjusted net losses to narrow on operating leverage

We expect adjusted net margin to improve on operating leverage across R&D and selling and distribution expenses


We expect MiniMax to narrow losses during 2026-28E and turn profitable in 2029E as operating leverage improves, particularly across R&D and selling and distribution expenses. On R&D, while training costs (79% of total R&D expense in 9M25) are likely to remain elevated amid intense competition in model development, accelerating commercialisation should drive R&D as % of revenue lower, in our view, from 320% in 2025 to 32% in 2029E. MiniMax has a dedicated in-house AI infrastructure team and efficiency-focused training framework, which we believe will further enhance training efficiency over time. On the selling and distribution expense front, we expect increasing revenue scale to support operating leverage, although fluctuations may occur alongside new product launches that require additional marketing investment. We forecast adjusted NPM to improve from -317% in 2025 to +10% by 2029E.

MiniMax: opex to revenue ratio over 2023-29E (%)



Source: Company data, Daiwa forecasts

MiniMax: Adjusted net margin over 2023-29E (%)



Source: Company data, Daiwa forecasts

Valuation methodology and risks

Initiating coverage with a Buy (1) rating

We initiate coverage of MiniMax with a Buy (1) rating and 12-month TP of HKD530, as we see the company as a direct beneficiary of the AI commercialisation boom in China. We believe MiniMax is well suited to the next phase of AI adoption, where commercial viability and scalable deployment are key enterprise requirements.

We assign 5x EV/ARR to MiniMax's 2029E average ARR

We derive our 12-month TP by applying 2029E EV/ARR multiple of 5x and discounting it back 2.5 years at a discount rate of 15%. Our 2029E EV/ARR of 5x represents a substantial discount to peer valuations, including OpenAI (25x based on the last round EV/ARR on street estimates), Anthropic (16x), and MoonShot (63x), as implied by reported transactions. We believe frontier model developers focused on advancing state of the art intelligence command a valuation premium due to their strategic scarcity and technological leadership. In comparison, MiniMax's strategy is centred on commercial applicability, efficient deployment and monetisation, which we believe warrants a relatively conservative valuation framework.

While Minimax's stock price has declined more than 70% from its peak of HKD1,238 in March amid intensifying competition among foundation model developers and uncertainty surrounding monetisation timelines, we believe the sell-off has been excessive relative to the company's fundamentals and growth prospects. In our view, the recent weakness has created an attractive entry point ahead of multiple upside catalysts that support a valuation rerating in the next 12 months.

MiniMax: valuation summary

	2029E
Sales (USDm)	4,861
Average ARR (USDm)	5,715
Target EV/ARR multiple (x)	5x
Enterprise value (USDm)	28,575
Net debt (USDm)	-1,795
Equity value (USDm)	31,094
Shares outstanding (m)	314
Discount rate (%)	15%
Equity market premium	7.0%
Risk free rate	4.7%
Beta	1.5
Years to discount	2.5
Target price in 12M (HKD)	530

Source: Daiwa estimates

Key catalysts to watch

Following the release of M3 in June 2026, MiniMax continues to demonstrate strong competitiveness in both model performance and cost-efficiency, reinforcing our long-term investment thesis. In addition, its latest **video generation model H3 launch** in late-July has gained robust traction, with top rankings on Artificial Analysis (No.2 in text-to-video, No.3 in image-to-video, and No.1 in video editing). By open-sourcing H3, MiniMax can accelerate adoption and ecosystem development, helping lower customer acquisition costs while supporting revenue growth over time. Model line-up wise, management has teased a **3tn parameter model M3 Pro set to launch in Sep-Oct**, which could deliver meaningful capability gains and represent another technological milestone. We expect these model launches and continued traction to drive a meaningful step-up in ARR and boost investor sentiment, potentially leading to a rerating from current valuation levels. Notably, Z.AI announced in July that it had already achieved its annual ARR target of USD1bn, and we believe MiniMax could be on a similar trajectory.

Risks to our call

Intensified competition among AI model developers is the primary risk to our call

Intensified competition among AI model developers – primary risk. The foundation model industry remains highly competitive, with rapid technological advancements, continuous product innovation and evolving customer demands. Failure to sustain rapid model iteration and competitive model performance could weaken MiniMax's market position, leading to market share erosion, weaker customer adoption and declining pricing power.

Execution risk in R&D investments. MiniMax's key competitive advantage is built on its ability to continuously advance model capabilities while maintaining strong economic efficiency. Sustaining this advantage requires disciplined execution across compute utilization, model development, and R&D resource allocation. Any inefficiencies in these areas could lead to cost escalation, reducing profitability and weakening cash flow.

Compute supply and cost risks. Computing capacity is a critical resource for AI model developers, underpinning both model training and inference deployment. Constraints in GPU supply, including those arising from export controls and trade restrictions, as well as rising computing infrastructure costs, could slow MiniMax's model iteration pace, increase operating expenses, and constrain the scalability of its products and services.

Financing risk. Given MiniMax's history of operating losses and negative operating cash flow since inception, the company may require additional funding if monetisation and revenue scaling are slower than expected. That said, we believe near-term financing risk is limited following the company's strong capital raising activities, including c.USD600m from its IPO and a subsequent c.USD2bn equity and convertible bond financing in July, which significantly strengthened its liquidity position. We believe the current cash buffer provides substantial funding runway to support continued R&D investment and commercialisation efforts.

Copyright and licensing risks. MiniMax faces copyright infringement lawsuit in the US related to Hailuo AI, with major Hollywood studios, including Disney, Universal, and Warner Bros. Discovery, alleging unauthorised use and generation of copyrighted characters and content. An adverse ruling could result in financial penalties and/or operational restrictions on products or features. More broadly, as generative AI continues to evolve, intellectual property rights remain a key legal uncertainty, with the potential for additional claims from content owners and rising compliance and licensing costs across the industry.

ESG analysis

ESG opportunities

Opportunities	Management	Analyst comments
G Strong ESG commitment	2	MiniMax integrates ESG values into its corporate strategy and operations. The company has identified 24 key ESG issues, such as environmental compliance and employee welfare, and regularly assess their impact long-term development. MiniMax also emphasizes suppliers' ESG performance, aiming to build sustainable partnerships and promote ESG standards across its ecosystem.

ESG risks

Risks	Management	Analyst comments
G Executive/board quality	2	33% of the board comprises independent directors, while directors bring extensive experience across leading internet, technology and investment firms.
G Capital management	2	Given its loss-making status, the company has yet to initiate dividend payout or share buybacks.
Related party & transaction	1	There is no evidence of any suspicious related-party transactions.
S Data security	2	MiniMax maintains a multi-layered data security framework covering governance, technology and operations, complies with cybersecurity and data protection regulations, limits data collection to necessary purposes, and safeguards data throughout its lifecycle via strict protection controls.
S Customer privacy	2	MiniMax embeds privacy-by-design, maintains dedicated data protection teams, implements data breach response mechanisms, and safeguards user privacy through compliance measures.
Company specific	Management	Analyst comments
G Intellectual property protection	2	MiniMax faces copyright infringement litigation related to the alleged unauthorized use of copyrighted content. However, the company has implemented comprehensive intellectual property management systems and dedicated legal oversight to mitigate IP risks throughout the product development lifecycle.

Note: Management score represents a company's ability to manage/benefit from certain ESG topics. The scores range from 1 to 3, with 1 being the strongest.

Update Date: 11 Aug 2026

Source: Daiwa, Company

Daiwa's Asia Pacific Research Directory

HONG KONG		
Hiroyuki GOTO	(852) 2773 8870	hiroyuki.goto@hk.daiwacm.com <i>Regional Head of Equity, Asia & Oceania; Regional Head of Asia Pacific Research</i>
John HETHERINGTON	(852) 2773 8787	john.hetherington@hk.daiwacm.com <i>Co-Head of Asia Pacific Research and Head of Research Publications</i>
Li XIONG	(852) 2773 8878	li.xiong@hk.daiwacm.com <i>Strategy (Hong Kong/China)</i>
Yue TAN	(852) 2848 4947	yue.tan@hk.daiwacm.com <i>Strategy (Greater China) and Macro Economics (China)</i>
Kelvin LAU	(852) 2848 4467	kelvin.lau@hk.daiwacm.com <i>Regional Head of Automobiles, Industrials and Transportation</i>
Evelyn ZHANG	(852) 2848 4970	evelyn.zhang@hk.daiwacm.com <i>Automobiles and Components (Hong Kong/China)</i>
Frank YIP	(852) 2773 8842	frank.yip@hk.daiwacm.com <i>Industrials and Transportation (Hong Kong/China)</i>
Carol XIA	(852) 2532 4349	carol.xia@hk.daiwacm.com <i>Head of China Consumer; Consumer (Hong Kong/China)</i>
Steven NIE	(852) 2848 4464	steven.nie@hk.daiwacm.com <i>Consumer (Hong Kong/China); Global Cryptocurrency</i>
Siman CHEN	(852) 2532 4350	siman.chen@hk.daiwacm.com <i>Consumer (Hong Kong/China)</i>
Jing YANG	(852) 2532 4308	jing.yang@hk.daiwacm.com <i>Head of China Healthcare</i>
John CHOI	(852) 2773 8730	john.choi@hk.daiwacm.com <i>Head of China Internet; Regional Head of Emerging Opportunities</i>
Dennis IP	(852) 2848 4068	dennis.ip@hk.daiwacm.com <i>Head of Asia ex-Japan Power, Utilities, Renewables & ESS (PURE) Research; US & CN Power Equipment (Gas, Nuclear, Solar, Grid, Battery & Rare Earth Materials); Tech Hardware – AIDC Power, PCB, Optical Fiber (China)</i>
Derek ZHANG	(852) 2532 4341	derek.zhang@hk.daiwacm.com <i>Power, Utilities, Renewables and ESS (PURE) – Gas & Gas Equipment, Rare Earth Materials (China), Utilities (Hong Kong), Power Equipment (US)</i>
Marco ZENG	(852) 2848 4430	marco.zeng@hk.daiwacm.com <i>Power, Utilities, Renewables and ESS (PURE) – US Power, PCB (China)</i>

CHINA		
Louis LUO	(86) 21 6841 3282	louis.luo@daiwacm.cn <i>Automobiles and Components, Industrials (China)</i>
Mavis MA	(86) 21 6841 3288	mavis.ma@daiwacm.cn <i>Battery Components; Industrials (China)</i>
Leo LU	(86) 21 6841 3286	leo.lu@daiwacm.cn <i>Power, Utilities, Renewables and ESS (PURE) – Solar, ESS, Lithium, AIDC Power (China)</i>
Skye LIANG	(86) 21 6841 3207	skye.liang@daiwacm.cn <i>Power, Utilities, Renewables and ESS (PURE) – Grid, Wind and Optical Fibre (China)</i>
William WU	(86) 21 6841 3200	william.wu@daiwacm.cn <i>Property (China/Hong Kong)</i>
Bintuo NI	(86) 21 6841 3228	bintuo.ni@daiwacm.cn <i>Tech Hardware and Semiconductors (China)</i>

SOUTH KOREA		
Yoonki BAE	(82) 2 787 9168	yoonki.bae@kr.daiwacm.com <i>Automobile Components, Healthcare</i>
Mike OH	(82) 2 787 9179	mike.oh@kr.daiwacm.com <i>Banking, Capital Goods – Defence, Construction and Shipbuilding, and Utilities</i>
Youngho JANG	(82) 2 787 9838	youngho.jang@kr.daiwacm.com <i>Consumer, Healthcare</i>
Henny JUNG	(82) 2 787 9182	henny.jung@kr.daiwacm.com <i>EV Batteries and Battery Materials, IT/Electronics (Small/Mid Cap), Automobiles and Components</i>
Daeho SON	(82) 2 787 9176	dh.son@kr.daiwacm.com <i>Industrials – Robotics</i>
Joon LEE	(82) 2 787 9151	hj.lee@kr.daiwacm.com <i>Media, Software and IT Services</i>
Thomas Y KWON	(82) 2 787 9181	yskwon@kr.daiwacm.com <i>Pan-Asia Head of Internet & Telecommunications; Online Games</i>
SK KIM	(82) 2 787 9173	sk.kim@kr.daiwacm.com <i>Head of Global Memory; IT/Electronics – Semiconductors/Displays and Tech Hardware</i>

TAIWAN		
Rick HSU	(886) 2 8758 6261	rick.hsu@daiwacm-cathay.com.tw <i>Head of Regional Technology; Head of Taiwan Research; Semiconductors (Regional)</i>
Sheng CHENG	(886) 2 8758 6253	sheng.cheng@daiwacm-cathay.com.tw <i>IT/Technology Hardware (Automation, Datacentre Components and PCB/CCL) (Greater China)</i>
Allan WANG	(886) 2 8758 6249	allan.wang@daiwacm-cathay.com.tw <i>IT/Technology Hardware (Automation, Datacentre Components and PCB/CCL) (Greater China)</i>
Leon CHEN	(886) 2 8758 6247	leon.chen@daiwacm-cathay.com.tw <i>IT/Technology – Semiconductors Supply Chain (Taiwan)</i>
Stacy LIN	(886) 2 8758 6252	stacy.lin@daiwacm-cathay.com.tw <i>IT/Technology – Semiconductors Supply Chain (Taiwan)</i>
Helen CHIEN	(886) 2 8758 6254	helen.chien@daiwacm-cathay.com.tw <i>Small/Mid Cap</i>

Daiwa's Offices

Office / Branch / Affiliate	Address	Tel	Fax
DAIWA SECURITIES GROUP INC			
HEAD OFFICE	Gran Tokyo North Tower, 1-9-1, Marunouchi, Chiyoda-ku, Tokyo, 100-6753	(81) 3 5555 3111	(81) 3 5555 0661
Daiwa Securities Trust Company	One Evertrust Plaza, Jersey City, NJ 07302, U.S.A.	(1) 201 333 7300	(1) 201 333 7726
Daiwa Securities Trust and Banking (Europe) PLC (Head Office)	5 King William Street, London EC4N 7JB, United Kingdom	(44) 207 320 8000	(44) 207 410 0129
Daiwa Europe Trustees (Ireland) Ltd	Level 3, Block 5, Harcourt Centre, Harcourt Road, Dublin 2, Ireland	(353) 1 603 9900	(353) 1 478 3469
Daiwa Capital Markets America Inc. New York Head Office	1251 Avenue of the Americas, 49th Floor, New York, NY 10020	(1) 212 612 7000	(1) 212 612 7100
Daiwa Capital Markets America Inc. San Francisco Branch	555 California Street, Suite 3360, San Francisco, CA 94104, U.S.A.	(1) 415 955 8100	(1) 415 956 1935
Daiwa Capital Markets Europe Limited, London Head Office	5 King William Street, London EC4N 7AX, United Kingdom	(44) 20 7597 8000	(44) 20 7597 8600
Daiwa Capital Markets Deutschland GmbH	Friedrich-Ebert-Anlage 35-37, 60327 Frankfurt am Main, Germany	(49) 69 27139 8100	(49) 69 27139 8190
Daiwa Capital Markets Europe Limited, Paris Representative Office	17, rue de Surène 75008 Paris, France	(33) 1 56 262 200	(33) 1 47 550 808
Daiwa Capital Markets Europe Limited, Moscow Representative Office	Midland Plaza 7th Floor, 10 Arbat Street, Moscow 119002, Russian Federation	(7) 495 641 3416	(7) 495 775 6238
Daiwa Capital Markets Europe Limited, Bahrain Branch	7th Floor, The Tower, Bahrain Commercial Complex, P.O. Box 30069, Manama, Bahrain	(973) 17 534 452	(973) 17 535 113
Daiwa Capital Markets Hong Kong Limited	Level 28, One Pacific Place, 88 Queensway, Hong Kong	(852) 2525 0121	(852) 2845 1621
Daiwa Capital Markets Singapore Limited	7 Straits View, Marina One East Tower, #16-05 & #16-06, Singapore 018936, Republic of Singapore	(65) 6387 8888	(65) 6282 8030
Daiwa Capital Markets Australia Limited	Level 34, Rialto North Tower, 525 Collins Street, Melbourne, Victoria 3000, Australia	(61) 3 9916 1300	(61) 3 9916 1330
DBP-Daiwa Capital Markets Philippines, Inc	18th Floor, Citibank Tower, 8741 Paseo de Roxas, Salcedo Village, Makati City, Republic of the Philippines	(632) 813 7344	(632) 848 0105
Daiwa-Cathay Capital Markets Co Ltd	14/F, 200, Keelung Road, Sec 1, Taipei, Taiwan, R.O.C.	(886) 2 2723 9698	(886) 2 2345 3638
Daiwa Securities Capital Markets Korea Co., Ltd.	20 Fl.& 21Fl. One IFC, 10 Gukjegeumyung-Ro, Yeongdeungpo-gu, Seoul, Korea	(82) 2 787 9100	(82) 2 787 9191
Daiwa Securities Co. Ltd., Beijing Representative Office	Room 301/302, Kerry Center, 1 Guanghua Road, Chaoyang District, Beijing 100020, People's Republic of China	(86) 10 6500 6688	(86) 10 6500 3594
Daiwa (Shanghai) Corporate Strategic Advisory Co. Ltd.	44/F, Hang Seng Bank Tower, 1000 Lujiazui Ring Road, Pudong, Shanghai China 200120, People's Republic of China	(86) 21 3858 2000	(86) 21 3858 2111
Daiwa Securities Co. Ltd., Bangkok Representative Office	18 th Floor, M Thai Tower, All Seasons Place, 87 Wireless Road, Lumpini, Pathumwan, Bangkok 10330, Thailand	(66) 2 252 5650	(66) 2 252 5665
Daiwa Securities Co. Ltd., Hanoi Representative Office	Suite 405, Pacific Palace Building, 83B, Ly Thuong Kiet Street, Hoan Kiem Dist. Hanoi, Vietnam	(84) 4 3946 0460	(84) 4 3946 0461
DAIWA INSTITUTE OF RESEARCH LTD			
HEAD OFFICE	15-6, Fuyuki, Koto-ku, Tokyo, 135-8460, Japan	(81) 3 5620 5100	(81) 3 5620 5603
MARUNOUCHI OFFICE	Gran Tokyo North Tower, 1-9-1, Marunouchi, Chiyoda-ku, Tokyo, 100-6756	(81) 3 5555 7011	(81) 3 5202 2021
New York Research Center	11th Floor, Financial Square, 32 Old Slip, NY, NY 10005-3504, U.S.A.	(1) 212 612 6100	(1) 212 612 8417
London Research Centre	3/F, 5 King William Street, London, EC4N 7AX, United Kingdom	(44) 207 597 8000	(44) 207 597 8550

Important Disclosures and Disclaimer

This publication is produced by Daiwa Securities Group Inc. and/or its non-U.S. affiliates, and distributed by Daiwa Securities Group Inc. and/or its non-U.S. affiliates, except to the extent expressly provided herein. This publication and the contents hereof are intended for information purposes only, and may be subject to change without further notice. Any use, disclosure, distribution, dissemination, copying, printing or reliance on this publication for any other purpose without our prior consent or approval is strictly prohibited. Neither Daiwa Securities Group Inc. nor any of its respective parent, holding, subsidiaries or affiliates, nor any of its respective directors, officers, servants and employees, represent nor warrant the accuracy or completeness of the information contained herein or as to the existence of other facts which might be significant, and will not accept any responsibility or liability whatsoever for any use of or reliance upon this publication or any of the contents hereof. Neither this publication, nor any content hereof, constitute, or are to be construed as, an offer or solicitation of an offer to buy or sell any of the securities or investments mentioned herein in any country or jurisdiction nor, unless expressly provided, any recommendation or investment opinion or advice. Any view, recommendation, opinion or advice expressed in this publication may not necessarily reflect those of Daiwa Securities Group Inc., and/or its affiliates nor any of its respective directors, officers, servants and employees except where the publication states otherwise. This research report is not to be relied upon by any person in making any investment decision or otherwise advising with respect to, or dealing in, the securities mentioned, as it does not take into account the specific investment objectives, financial situation and particular needs of any person.

Daiwa Securities Group Inc., its subsidiaries or affiliates, or its or their respective directors, officers and employees from time to time have trades as principals, or have positions in, or have other interests in the securities of the company under research including market making activities, derivatives in respect of such securities or may have also performed investment banking and other services for the issuer of such securities. Daiwa Securities Group Inc., its subsidiaries or affiliates do and seek to do business with the company(s) covered in this research report. Therefore, investors should be aware that a conflict of interest may exist. The following are additional disclosures.

Ownership of Securities

For "Ownership of Securities" information, please visit BlueMatrix disclosure link at <https://daiwa3.bluematrix.com/sellside/Disclosures.action>.

Investment Banking Relationship

For "Investment Banking Relationship", please visit BlueMatrix disclosure link at <https://daiwa3.bluematrix.com/sellside/Disclosures.action>.

Daiwa Securities Co. Ltd. and Daiwa Securities Group Inc.

Daiwa Securities Co. Ltd. is a subsidiary of Daiwa Securities Group Inc.

Investment Banking Relationship

Within the preceding 12 months, Daiwa Capital Markets Hong Kong Limited has lead-managed public offerings and/or secondary offerings of the equities or relevant securities of the following companies: WeRide Inc (800 HK/ WRD US), SEYOND (2665 HK) and SENASIC Electronics Technology Co Ltd (6675 HK).

*Subsidiaries of Daiwa Securities Group Inc. for the purposes of this section shall mean any one or more of: Daiwa Capital Markets Hong Kong Limited (大和資本市場香港有限公司), Daiwa Capital Markets Singapore Limited, Daiwa Capital Markets Australia Limited, Daiwa-Cathay Capital Markets Co., Ltd., Daiwa Securities Capital Markets Korea Co., Ltd.

Within the last 12 months, the subsidiaries and/or affiliates of Daiwa Securities Group Inc. * has received compensation for investment banking services from NH Investment & Securities Co Ltd (005940 KS), E-Mart Inc (139480 KS), Samsung Securities Co Ltd (016360 KS), KT Corporation (030200 KS) and Hyundai Department Store Co Ltd (069960 KS).

Hong Kong

This research is distributed in Hong Kong by Daiwa Capital Markets Hong Kong Limited (大和資本市場香港有限公司) ("DHK") which is regulated by the Hong Kong Securities and Futures Commission. Recipients of this research in Hong Kong may contact DHK in respect of any matter arising from or in connection with this research.

Investment Banking Relationship

Within the preceding 12 months, Daiwa Capital Markets Hong Kong Limited has lead-managed public offerings and/or secondary offerings of the equities or relevant securities of the following companies: WeRide Inc (800 HK/ WRD US), SEYOND (2665 HK) and SENASIC Electronics Technology Co Ltd (6675 HK).

Within the preceding 12 months, Daiwa Capital Markets Hong Kong Limited has received compensation for investment banking services from -.

Relevant Relationship (DHK)

DHK may from time to time have an individual employed by or associated with it serves as an officer of any of the companies under its research coverage.

Singapore

This research is distributed in Singapore by Daiwa Capital Markets Singapore Limited and it may only be distributed in Singapore to accredited investors, expert investors and institutional investors as defined in the Financial Advisers Regulations of Singapore and the Securities and Futures Act 2001 of Singapore, both as amended from time to time. By virtue of distribution to these category of investors, Daiwa Capital Markets Singapore Limited and its representatives are not required to comply with Section 45 of the Financial Advisers Act 2001 of Singapore (Section 45 relates to disclosure of Daiwa Capital Markets Singapore Limited's interest and/or any associated or connected person's interest in securities). Recipients of this research in Singapore may contact Daiwa Capital Markets Singapore Limited in respect of any matter arising from or in connection with this research.

Australia

This research is distributed in Australia by Daiwa Capital Markets Australia Limited and it may only be distributed in Australia to wholesale investors within the meaning of the Corporations Act. Recipients of this research in Australia may contact Daiwa Capital Markets Stockbroking Limited in respect of any matter arising from or in connection with the research.

Taiwan

This research is solely for reference and not intended to provide tailored investment recommendations. This research is distributed in Taiwan by Daiwa-Cathay Capital Markets Co., Ltd. and it may only be distributed in Taiwan to specific customers who have signed recommendation contracts with Daiwa-Cathay Capital Markets Co., Ltd. and non-customers including (i) professional institutional investors, (ii) TWSE or TPEX listed companies, upstream and downstream vendors, and specialists that offer or seek advice, and (iii) potential customers with an actual need for business development in accordance with the Operational Regulations Governing Securities Firms Recommending Trades in Securities to Customers. Recipients of this research including non-customer recipients of this research shall not provide it to others or engage in any activities in connection with this research which may involve conflicts of interests. Neither Daiwa-Cathay Capital Markets Co., Ltd. nor its personnel who writes or reviews the research report has any conflict of interest in this research. Since Daiwa-Cathay Capital Markets Co., Ltd. does not operate brokerage trading business in foreign markets, **this research is prepared on a "without recommendation" to any foreign securities basis** and Daiwa-Cathay Capital Markets Co., Ltd. does not accept orders from customers to trade in such foreign securities. Recipients of this research shall carefully judge their own investment risk and take full responsibility for the results of any resulting investments in the companies and/or sectors featured in this research. Without the prior written permission of Daiwa-Cathay Capital Markets Co., Ltd., recipients of this research are prohibited from disclosing the research to the media, reprinting the research, or quoting from the research to other parties. Recipients of this research in Taiwan may contact Daiwa-Cathay Capital Markets Co., Ltd. in respect of any matter arising from or in connection with the research.

Thailand

This research is distributed to only institutional investors in Thailand. This report is provided to you for informational purposes only. Information contained herein is not intended to be, and shall not be construed as, an offer or an invitation for subscription or purchase of securities in Thailand. This document has not been and will not be registered with, or approved by, the Office of the Securities and Exchange Commission of Thailand. As neither Daiwa Securities Group Inc. nor any of its subsidiaries or affiliates, nor any of the foregoing entities' respective directors, officers, servants and employees are licensed to carry on any securities business in Thailand, this document is intended to be read by the addressee who is an institutional investor only and may not be circulated or distributed, whether directly or indirectly, to the public or any members of the public in Thailand, unless to the extent permitted by applicable laws and regulations.

This report is prepared by analysts who are employed by Daiwa Securities Group Inc. and/or its non-U.S. affiliates. Neither Daiwa Securities Group Inc. nor any of its subsidiaries or affiliates, nor any of the foregoing entities' respective directors, officers, servants and employees accept any liability whatsoever for any direct or consequential loss arising from any use of this research or its contents.

The information and opinions contained herein have been compiled or arrived at from sources believed to be reliable. However, neither Daiwa Securities Group Inc. nor any of its subsidiaries or affiliates, nor any of the foregoing entities' respective directors, officers, servants and employees make any representation or warranty, express or implied, as to their accuracy or completeness. Expressions of opinion herein are subject to change without notice. The use of any information, forecasts and opinions contained in this report shall be at the sole discretion and risk of the user.

Daiwa Securities Group Inc., its subsidiaries and affiliates and their respective directors, officers, servants and employees may have positions and financial interest in securities mentioned in this research. Daiwa Securities Group Inc. and its subsidiaries and affiliates may from time to time perform investment banking or other services for, or solicit investment banking or other business from, any entity mentioned in this research. Therefore, investors should be aware of conflict of interest that may affect the objectivity of this research.

United Kingdom

These materials constitute independent research for the purposes of the Financial Conduct Authority ("FCA") Conduct of Business Rules. These materials have been produced by Daiwa Securities Co. Ltd ("DSCL") and/or its affiliates and are distributed by Daiwa Capital Markets Europe Limited ("DCME") to persons classified as (i) eligible counterparties; and (ii) professional clients for the purposes of the Conduct of Business Rules of the FCA. These materials should not be disseminated to any retail clients in the United Kingdom. DCME is authorised and regulated by the FCA and is a member of the London Stock Exchange. DCME's FCA registration number is 124490 and the FCA's registered address is 25 The North Colonnade, London, E14 5HS.

DCME makes no representation as to the completeness, accuracy or timeliness of any of the materials produced by DSCL, and does not accept any liability for any losses, costs, liabilities or expenses which may arise directly or indirectly from any use of, or reliance on the materials.

DCME and/or its affiliates may, from time to time, to the extent permitted by law, participate or invest in, or be mandated in respect of transactions with the issuer(s) referred to in the materials prepared by DSCL, perform services for or solicit business from such issuers, and/or have a position or effect transactions in a particular issuer's securities and/or may have acted as an underwriter during the past twelve months in respect of a particular issuer of its securities. DCME does not review DSCL's materials prior to publication on DCME's platforms and has no influence over its content. DCME has in place organisational arrangements for the prevention and avoidance of conflicts of interest, including information barriers to control the flow of information between the private and public sides of the firm. A summary of DCME's conflict management policies are available at <http://www.uk.daiwacm.com/about-us/corporate-governance-regulatory>.

DSCL's materials are distributed by DCME only as permitted by law. The materials are not directed to, or intended for distribution to or use by, any person or entity located in any jurisdiction where such distribution, publication, availability or use would be contrary to law or regulation or would subject DCME to any regulation or licensing requirement within such jurisdiction.

Germany

Research reports are produced by Daiwa Securities Co. Ltd. and/or its affiliates and are distributed by Daiwa Capital Markets Deutschland GmbH in the European Union. Daiwa Capital Markets Deutschland GmbH is authorised and regulated by Bundesanstalt für Finanzdienstleistungsaufsicht ("BaFin") under the reference number 149361. Daiwa Capital Markets Deutschland GmbH and its affiliates may, from time to time, to the extent permitted by law, participate or invest in, or be mandated in respect of, other financing transactions with the issuer(s) of the securities referred to herein (the "Securities"), perform services for or solicit business from such issuers, and/or have a position or effect transactions in the Securities or options thereof.

This publication is intended for investors who are Professional Clients or Eligible Counterparties in the European Union within the meaning of Directive 2014/65/EU ("MiFID II") and should therefore not be distributed to Retail Clients in the European Union.

Daiwa Capital Markets Deutschland GmbH has in place organisational arrangements for the identification, prevention and management of conflicts of interest. Information regarding our conflict management are available at: <https://www.de.daiwacm.com/policies/>

Additional disclosures regarding specific company names and related financial instruments are available at: <https://daiwa3.bluematrix.com/sellside/Disclosures.action>.

Bahrain

This research material is distributed in Bahrain by Daiwa Capital Markets Europe Limited, Bahrain Branch, regulated by The Central Bank of Bahrain and holds Investment Business Firm – Category 2 license and having its official place of business at the Bahrain World Trade Centre, South Tower, 7th floor, P.O. Box 30069, Manama, Kingdom of Bahrain. Tel No. +973 17534452 Fax No. +973 535113

United States

This research is distributed into the United States directly by Daiwa Capital Markets Hong Kong Limited (DCMHK) and in certain cases indirectly by Daiwa Capital Markets America Inc. (DCMA), a U.S. Securities and Exchange Commission registered broker-dealer and FINRA member firm, exclusively to "major U.S. institutional investors", as defined under Rule 15a-6 promulgated under the U.S. Securities Exchange Act of 1934, as amended, and as interpreted by the staff of the U.S. Securities and Exchange Commission (SEC). Where this is co-branded research published by Daiwa's partners in collaboration with other parties not affiliated with Daiwa, this research is distributed in the United States by DCMHK only. This report is not an offer to sell or the solicitation of any offer to buy securities. U.S. customers wishing to effect transactions in any designated investment discussed in this report should do so through a qualified salesperson of DCMA. Non-U.S. customers wishing to effect transactions in any designated investment discussed in this report should contact a Daiwa entity in their local jurisdiction. The securities or other investment products discussed in this report may not be eligible for sale in some jurisdictions.

Analysts employed outside the U.S., as specifically indicated elsewhere in this report, are not registered as research analysts with FINRA. These analysts may not be associated persons of DCMA, and therefore may not be subject to FINRA Rule 2241 restrictions on communications with a subject company, public appearances and trading securities held by a research analyst account.

ADDITIONAL IMPORTANT DISCLOSURES CAN BE FOUND AT:

<https://daiwa3.bluematrix.com/sellside/Disclosures.action>

Ownership of Securities

For "Ownership of Securities" information please visit BlueMatrix disclosure link at <https://daiwa3.bluematrix.com/sellside/Disclosures.action>.

Investment Banking Relationships

For "Investment Banking Relationships" please visit BlueMatrix disclosure link at <https://daiwa3.bluematrix.com/sellside/Disclosures.action>.

DCMA Market Making

For "DCMA Market Making" please visit BlueMatrix disclosure link at <https://daiwa3.bluematrix.com/sellside/Disclosures.action>.

Research Analyst Conflicts

For updates on "Research Analyst Conflicts" please visit BlueMatrix disclosure link at <https://daiwa3.bluematrix.com/sellside/Disclosures.action>. The principal research analysts who prepared this report have no financial interest in securities of the issuers covered in the report, are not (nor are any members of their household) an officer, director or advisory board member of the issuer(s) covered in the report, and are not aware of any material relevant conflict of interest involving the analyst or DCMA, and did not receive any compensation from the issuer during the past 12 months except as noted: no exceptions.

Research Analyst Certification

For updates on "Research Analyst Certification" and "Rating System" please visit BlueMatrix disclosure link at <https://daiwa3.bluematrix.com/sellside/Disclosures.action>. The views about any and all of the subject securities and issuers expressed in this Research Report accurately reflect the personal views of the research analyst(s) primarily responsible for this report (or the views of the firm producing the report if no individual analyst is named on the report); and no part of the compensation of such analyst (or no part of the compensation of the firm if no individual analyst is named on the report) was, is, or will be directly or indirectly related to the specific recommendations or views contained in this Research Report.

The following explains the rating system in the report as compared to relevant local indices, unless otherwise stated, based on the beliefs of the author of the report.

- "1": the security could outperform the local index by more than 15% over the next 12 months.
- "2": the security is expected to outperform the local index by 5-15% over the next 12 months.
- "3": the security is expected to perform within 5% of the local index (better or worse) over the next 12 months.
- "4": the security is expected to underperform the local index by 5-15% over the next 12 months.
- "5": the security could underperform the local index by more than 15% over the next 12 months.

Disclosure of investment ratings

Rating	Percentage of total
Buy*	72.63%
Hold**	20.53%
Sell***	6.84%

Source: Daiwa

Notes: data is for single-branded Daiwa research in Asia (ex Japan) and correct as of 30 June 2026.

* comprised of Daiwa's Buy and Outperform ratings.

** comprised of Daiwa's Hold ratings.

*** comprised of Daiwa's Underperform and Sell ratings.

Additional information may be available upon request.

Japan - Notification items pursuant to Article 37 of the Financial Instruments and Exchange Law

(This Notification is only applicable to where report is distributed by Daiwa Securities Co. Ltd.)

If you decide to enter into a business arrangement with our company based on the information described in this report, we ask you to pay close attention to the following items.

- In addition to the purchase price of a financial instrument, our company will collect a trading commission* for each transaction as agreed beforehand with you. Since commissions may be included in the purchase price or may not be charged for certain transactions, we recommend that you confirm the commission for each transaction. In some cases, our company also may charge a maximum of ¥2 million per year as a standing proxy fee for our deposit of your securities, if you are a non-resident.
- For derivative and margin transactions etc., our company may require collateral or margin requirements in accordance with an agreement made beforehand with you. Ordinarily in such cases, the amount of the transaction will be in excess of the required collateral or margin requirements**.
- There is a risk that you will incur losses on your transactions due to changes in the market price of financial instruments based on fluctuations in interest rates, exchange rates, stock prices, real estate prices, commodity prices, and others. In addition, depending on the content of the transaction, the loss could exceed the amount of the collateral or margin requirements.
- There may be a difference between bid price etc. and ask price etc. of OTC derivatives handled by our company.
- Before engaging in any trading, please thoroughly confirm accounting and tax treatments regarding your trading in financial instruments with such experts as certified public accountants.

* The amount of the trading commission cannot be stated here in advance because it will be determined between our company and you based on current market conditions and the content of each transaction etc.

** The ratio of margin requirements etc. to the amount of the transaction cannot be stated here in advance because it will be determined between our company and you based on current market conditions and the content of each transaction etc.

When making an actual transaction, please be sure to carefully read the materials presented to you prior to the execution of agreement, and to take responsibility for your own decisions regarding the signing of the agreement with our company.

Corporate Name:	Daiwa Securities Co. Ltd.
Registered:	Financial Instruments Business Operator, Chief of Kanto Local Finance Bureau (Kin-sho) No.108
Memberships:	Japan Securities Dealers Association, The Financial Futures Association of Japan, Investment Management Association of Japan, Type II Financial Instruments Firms Association, Japan Security Token Offering Association